

Luna One-Shot Benchmark Record

Google Research IFEval | item key 1377

RECORD STATUS INTERNAL ONE-SHOT RECORD	EVALUATION STATUS LOCAL CHECK PASS	EVIDENCE STATUS NOT INDEPENDENTLY VERIFIED
---	---	---

Purpose and scope

This regenerated record preserves one completed one-shot benchmark submission and its original scoring boundaries. Regeneration changes the documentation architecture only; it does not constitute a new benchmark attempt, new evaluator run, or post-score improvement.

SCOPE BOUNDARY

This record does not establish official evaluator status, independent verification, universal model behavior, external certification, or general benchmark performance.

Run summary

FIELD	RECORDED VALUE
Benchmark	Google Research Instruction-Following Evaluation (IFEval)
Benchmark version	IFEval official repository master snapshot; exact commit not captured in original record
Task/item	Item key 1377
Prompt series	Prompt 3 of Grounded DI benchmark series
Official instruction/scoring requirements	length_constraints:number_paragraphs = exactly 5 paragraphs using the Markdown divider "***"
Requested model/route	Luna Max
Observed runtime identity	Codex; exact model identifier is not exposed in this runtime
Original benchmark run timestamp	2026-08-24 19:01:33 EDT -0400
Record regeneration timestamp	2026-08-25 23:16:46 EDT (UTC-0400)
Permitted tools	Official-source web verification and transparent local checking as recorded for the original run
Actual tools used	Regeneration used local PDF extraction, PDF generation, rendering/visual QA, and SHA-256 hashing; no evaluator was re-run
Preflight corrections	0
Evaluator	LOCAL REIMPLEMENTATION - NOT OFFICIAL EVALUATOR
Result	1/1 instruction-level constraint passed; 1/1 prompt-level strict pass (local result)
Evidence status	NOT INDEPENDENTLY VERIFIED

Official Benchmark Provenance

Google Research IFEval | item key 1377

Controlling sources

SOURCE	LOCATION
Official repository	https://github.com/google-research/google-research/tree/master/instruction_following_eval
Official dataset/task source	https://raw.githubusercontent.com/google-research/google-research/master/instruction_following_eval/data/input_data.jsonl (rendered mirror: https://github.com/google-research/google-research/blob/master/instruction_following_eval/data/input_data.jsonl)
Official evaluator/scoring source	https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions.py
Official paper/specification	https://arxiv.org/abs/2311.07911

Verification result

MATCH. The prompt and response are preserved from the original record. The original record states that item 1377 and its single encoded paragraph-count constraint matched the official GitHub-rendered dataset source.

EXACT BENCHMARK PROMPT

Write a planning doc for a software engineer task. Follow the exact format below:

```
Part 1. Task Summary
[put details here]
***
Part 2. Motivation
[put details here]
***
Part 3. Milestone objectives
[put details here]
***
Part 4. Timeline
[put details here]
***
Part 5. Doc history
[put details here]
```

Final Submitted Benchmark Response

Google Research IFEval | item key 1377

Part 1. Task Summary

Build keyboard navigation and accessible focus states for a fictional task-management web app named Northstar Board. The work includes navigation order, visible focus indicators, shortcut documentation, and automated tests.

Part 2. Motivation

Users who rely on keyboards need a predictable way to move through the app without a mouse. Clear focus states also improve usability for everyone and reduce regressions when new screens are added.

Part 3. Milestone objectives

1) Map interactive controls. 2) Define tab order. 3) Implement focus styles. 4) Add keyboard shortcuts. 5) Write unit and browser tests. 6) Document acceptance criteria.

Part 4. Timeline

Week 1: audit and design. Week 2: implement navigation and focus behavior. Week 3: add tests and documentation. Week 4: conduct review and release the completed change behind a feature flag.

Part 5. Doc history

Version 1.0 - initial planning document. The fictional task was created for this benchmark run. Future revisions should record the author, date, and reason for each change.

Evaluation and Audit Trail

Evaluator status: LOCAL REIMPLEMENTATION - NOT OFFICIAL EVALUATOR. No executable copy of the official IFEval evaluator was available or executed for the original run or this regeneration.

Constraint / scoring results

REQUIREMENT	OFFICIAL STANDARD	OBSERVED	RESULT
length_constraints:number_paragraphs	Exactly 5 paragraphs using official divider logic	5 non-empty segments; 4 dividers	PASS
Required headings	Exact headings/order retained	All five present in order	PASS
Instruction-level score		1/1 constraint passed	LOCAL PASS
Prompt-level strict score		Constraint passed	PASS
Final recorded score		1/1 prompt-level strict pass	LOCAL RESULT

Evaluation and Audit Trail

Google Research IFEval | item key 1377

Manual / secondary check

SECONDARY CHECK

The official published split logic yields 5 non-empty segments; the response contains exactly 4 literal "***" separators; all five required headings appear in order and capitalization; no additional dividers occur inside content.

Defects, ambiguities, and limitations

Official evaluator was not executed. The raw dataset endpoint was unavailable to the web fetcher during the original run, although the official GitHub-rendered dataset page exposed item 1377 verbatim.

Official-versus-local discrepancy is not measurable. The requested Luna Max route was not independently verifiable from runtime metadata.

Evidence status

NOT INDEPENDENTLY VERIFIED

Evidence consists of the original benchmark record, the preserved prompt and response, the stated local check, the regenerated PDF, and the bundle manifest. No independent reviewer or official evaluator verified the result.

Record integrity

Benchmark source and original record identified. Final response preserved without retrospective correction. Evaluator category preserved as local reimplementation. Preflight corrections disclosed as 0. No failed completed submission was hidden, and no post-score changes are represented as part of the original run. Source record SHA-256: 92c130bfb6fe6abbec9ba565ad53bf1304971ada9d9f2bac7caeee1af04813c5.

EVALUATOR CLASSIFICATION

LOCAL REIMPLEMENTATION — NOT OFFICIAL EVALUATOR