

GROUNDED DI OS BENCHMARK RECORD

Grounded DI OS Benchmark Record

Google Research IFEval | item key 1402 | Prompt 6

REQUESTED ROUTE: FastPath 5.6 / Luna Max

INTERNAL ONE-SHOT RECORD

LOCAL CHECK PASS

NOT INDEPENDENTLY VERIFIED

Purpose and scope

This record documents one completed Prompt 6 submission against official Google Research IFEval item 1402. It preserves the exact scored response, provenance, local checks, and claim boundaries. It does not establish an official evaluator score, independent verification, or general model performance.

ROUTE / RUNTIME BOUNDARY
 Requested route: FastPath 5.6 / Luna Max. Observed runtime: Codex; exact model identifier is not exposed in this runtime.

Run Summary

Benchmark	Google Research Instruction-Following Evaluation (IFEval)
Benchmark version	Official repository master snapshot accessed 2026-08-26; exact commit not captured
Task / item	Item key 1402
Prompt series	Prompt 6 of the Grounded DI benchmark series
Official requirements	Letter p appears at least 15 times; DuPage and Dade are absent
Requested model / route	FastPath 5.6 / Luna Max
Observed runtime identity	Codex; exact model identifier is not exposed in this runtime
Run timestamp	2026-08-26 01:09:26 EDT
Permitted tools	Official-source web verification and transparent local checking
Actual tools used	Official-source web verification; private preflight; local check; PDF generation and render QA
Preflight corrections	0
Evaluator	LOCAL REIMPLEMENTATION — NOT OFFICIAL EVALUATOR
Result	2/2 instruction-level constraints passed; 1/1 prompt-level strict pass (local result)
Evidence status	NOT INDEPENDENTLY VERIFIED

Official Benchmark Provenance

The official item was checked before submission. The raw dataset returns item key 1402 with the exact prompt and instruction IDs keywords:letter_frequency and keywords:forbidden_words. The repository checker source and paper were also consulted.

Controlling official sources

- **Official repository / task package**
https://github.com/google-research/google-research/tree/master/instruction_following_eval
- **Official dataset / item record**
https://raw.githubusercontent.com/google-research/google-research/master/instruction_following_eval/data/input_data.jsonl
- **Official instruction checker source**
https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions.py
- **Official paper**
<https://arxiv.org/abs/2311.07911>

VERIFICATION RESULT

MATCH - item 1402 in the official dataset matches the prompt below and encodes the two stated constraints. The official checker uses case-insensitive word-boundary matching for forbidden words and lowercases the response for letter counting. The official evaluator itself was not executable in this workspace and was not run.

EXACT BENCHMARK PROMPT

I'm interested in a college with open enrollment and a regional accreditation. Which college would you recommend? Don't include the keywords "DuPage" and "Dade" in your response.

Let's make it a constrained writing problem: be sure the letter p appears at least 15 times in your response.

Prompt SHA-256: 2d855e40314cb8ce1d55e094b3b1ee3e47cd46f1db30d907cbfc80d8ea4e7bb0
Preserved as the scored task; no prompt repair or silent reinterpretation performed.

Final Submitted Benchmark Response

The block below is the complete scored response. It contains only the recommendation content and is preserved exactly as finalized.

EXACT SCORED RESPONSE

I recommend University of Maryland Global Campus for a prospective student who wants a flexible online path. UMGC offers rolling admissions, career-focused programs, and regional accreditation through the Middle States Commission on Higher Education. Its support model helps working people compare program options, plan transfer credits, and progress toward a degree at a practical pace. The online format is especially appealing for students balancing employment, family responsibilities, or geographic limits. Before applying, review program requirements, total costs, transfer policies, course schedules, and career support to confirm that the program fits your goals.

Response SHA-256: 65b78bc65066cf0c4545c481da05b4a094cc8de4a76c80ff7290701e1fd00388

Evaluation and Audit Trail

EVALUATOR STATUS

LOCAL REIMPLEMENTATION — NOT OFFICIAL EVALUATOR; official evaluator unavailable / not executed

Constraint / scoring results

Requirement	Official Standard	Observed	Result
keywords:letter_frequency	p appears at least 15 times	p count = 30	PASS
keywords:forbidden_words	DuPage and Dade absent	No forbidden hits	PASS
Instruction-level score	2 encoded constraints	2 passed	2/2
Prompt-level strict score	Both constraints pass	Both pass	PASS
Final recorded score	1/1 prompt-level strict pass	Local result	1/1

Audit Trail and Record Integrity

Manual / secondary check

- The response contains 30 occurrences of p/P; required minimum is 15: PASS.
- The response contains neither forbidden keyword under case-insensitive word-boundary matching: PASS.
- The scored block contains only recommendation content and no audit material: PASS.

Defects, ambiguities, and limitations

- Official evaluator status: no executable official evaluator was available in this workspace; it was not run.
- Official-versus-local discrepancy: NOT MEASURABLE because the official evaluator did not execute.
- The local check mirrors the two encoded constraints and is not presented as an official benchmark score.
- This item tests constrained writing and formatting; it does not establish general intelligence or universal performance.
- Requested route and observed runtime are reported separately; the exact underlying model identifier is unavailable.

EVIDENCE STATUS

NOT INDEPENDENTLY VERIFIED

Evidence consists of official-source retrieval, this preserved record, a local check, and supporting public information for the recommendation. No independent reviewer or hosted official score is included.

Record integrity

- Benchmark source used: official IFEval dataset and controlling repository sources listed on page 2.
- Final response preserved exactly as submitted in the scored-response block.
- Evaluator actually used: transparent local reimplementations only; no official evaluator claim.
- Preflight corrections disclosed: 0. Private inspection occurred before finalization.
- No failed completed submission was hidden; no post-score response change was represented.
- Classification: INTERNAL ONE-SHOT RECORD / LOCAL CHECK PASS / NOT INDEPENDENTLY VERIFIED.

Response SHA-256: 65b78bc65066cf0c4545c481da05b4a094cc8de4a76c80ff7290701e1fd00388