

GROUNDING DI OS BENCHMARK RECORD

Grounded DI OS Benchmark Record

Google Research IFEval | item key 1481 | Prompt 7

REQUESTED ROUTE: FastPath 5.6 / Luna Max

INTERNAL ONE-SHOT RECORD

LOCAL CHECK PASS

NOT INDEPENDENTLY VERIFIED

Purpose and scope

This record documents one completed Prompt 7 submission against official Google Research IFEval item 1481. It preserves the exact scored response, provenance, local checks, and claim boundaries. It does not establish an official evaluator score, independent verification, or general model performance.

ROUTE / RUNTIME BOUNDARY
 Requested route: FastPath 5.6 / Luna Max. Observed runtime: Codex; exact model identifier is not exposed in this runtime.

Run Summary

Benchmark	Google Research Instruction-Following Evaluation (IFEval)
Benchmark version	Official repository master snapshot accessed 2026-08-26; exact commit not captured
Task / item	Item key 1481
Prompt series	Prompt 7 of the Grounded DI benchmark series
Official requirements	Exactly 2 SECTION markers; exactly 2 Markdown bullet lines
Requested model / route	FastPath 5.6 / Luna Max
Observed runtime identity	Codex; exact model identifier is not exposed in this runtime
Run timestamp	2026-08-26 09:29:15 EDT
Permitted tools	Official-source web verification and transparent local checking
Actual tools used	Official-source web verification; private preflight; local check; PDF generation and render QA
Preflight corrections	0
Evaluator	LOCAL REIMPLEMENTATION — NOT OFFICIAL EVALUATOR
Result	2/2 instruction-level constraints passed; 1/1 prompt-level strict pass (local result)
Evidence status	NOT INDEPENDENTLY VERIFIED

Official Benchmark Provenance

The official item was checked before submission. The raw dataset returns item key 1481 with the exact prompt and instruction IDs detectable_format:multiple_sections and detectable_format:number_bullet_lists. The repository checker source and paper were also consulted.

Controlling official sources

- **Official repository / task package**
https://github.com/google-research/google-research/tree/master/instruction_following_eval
- **Official dataset / item record**
https://raw.githubusercontent.com/google-research/google-research/master/instruction_following_eval/data/input_data.jsonl
- **Official instruction checker source**
https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions.py
- **Official paper**
<https://arxiv.org/abs/2311.07911>

VERIFICATION RESULT

MATCH - item 1481 in the official dataset matches the prompt below and encodes two formatting constraints. The official evaluator itself was not executable in this workspace and was not run.

EXACT BENCHMARK PROMPT

Write a casual, interesting, and weird resume for Antonia Maj who is applying for a job at a coffee company. They have experience in marketing, customer service, and sales. They are educated but not majored in anything related to coffee.

Make sure to include at least two sections marking the beginning of each section with 'SECTION X'. In your entire response make sure to use exactly two bullet points in markdown format.

Please use the following bullet point format:

- * Text for bullet 1
- * Text for bullet 2

Prompt SHA-256: ab75fde6eac70021ff0720297ba736ffa38e11dd5854d2ffa07e0386e7578034

Preserved as the scored task; no prompt repair or silent reinterpretation performed.

Final Submitted Benchmark Response

The block below is the complete scored response. It contains only the resume and is preserved exactly as finalized.

EXACT SCORED RESPONSE

SECTION 1 - PROFILE

Antonia Maj brings a casual blend of marketing knowledge, customer service skill, and sales experience to a coffee company. An educated and curious professional, Antonia enjoys turning everyday customer questions into friendly conversations and memorable coffee moments.

* Coffee-minded communicator who can explain products clearly, welcome every guest, and build loyal customer relationships.

* Creative promoter who can plan playful campaigns, support a busy team, and track practical sales results.

SECTION 2 - EXPERIENCE AND STYLE

Antonia is prepared to learn quickly, organize a lively cafe workflow, and bring a slightly weird sense of humor to the workday. The goal is to make each customer feel noticed while helping the company grow.

Response SHA-256: a1943274a2ff9a8b3df490f2e31f2b00b04638c64c0b8c2c6a453ad131a95c78

Evaluation and Audit Trail

EVALUATOR STATUS

LOCAL REIMPLEMENTATION — NOT OFFICIAL EVALUATOR; official evaluator unavailable / not executed

Constraint / scoring results

Requirement	Official Standard	Observed	Result
detectable_format:multiple_sections	2 SECTION markers	SECTION count = 2	PASS
detectable_format:number_bullet_lists	2 Markdown bullets	Bullet count = 2	PASS
Instruction-level score	2 encoded constraints	2 passed	2/2
Prompt-level strict score	Both constraints pass	Both pass	PASS
Final recorded score	1/1 prompt-level strict pass	Local result	1/1

Audit Trail and Record Integrity

Manual / secondary check

- The response contains exactly 2 section markers beginning with SECTION: PASS.
- The response contains exactly 2 Markdown bullet lines beginning with '* ': PASS.
- The scored block contains only resume content and no audit material: PASS.

Defects, ambiguities, and limitations

- Official evaluator status: no executable official evaluator was available in this workspace; it was not run.
- Official-versus-local discrepancy: NOT MEASURABLE because the official evaluator did not execute.
- The local check mirrors the two encoded formatting constraints and is not presented as an official benchmark score.
- This item tests constrained resume formatting; it does not establish general intelligence or universal performance.
- Requested route and observed runtime are reported separately; the exact underlying model identifier is unavailable.

EVIDENCE STATUS

NOT INDEPENDENTLY VERIFIED

Evidence consists of official-source retrieval, this preserved record, and a transparent local check. No independent reviewer or hosted official score is included.

Record integrity

- Benchmark source used: official IFEval dataset and controlling repository sources listed on page 2.
- Final response preserved exactly as submitted in the scored-response block.
- Evaluator actually used: transparent local reimplementations only; no official evaluator claim.
- Preflight corrections disclosed: 0. Private inspection occurred before finalization.
- No failed completed submission was hidden; no post-score response change was represented.
- Classification: INTERNAL ONE-SHOT RECORD / LOCAL CHECK PASS / NOT INDEPENDENTLY VERIFIED.

Response SHA-256: a1943274a2ff9a8b3df490f2e31f2b00b04638c64c0b8c2c6a453ad131a95c78