

GROUNDING DI OS BENCHMARK RECORD

Grounding DI OS Benchmark Record

Google Research IFEval | item key 1508 | Prompt 8

REQUESTED ROUTE: FastPath 5.6 / Luna Max

INTERNAL ONE-SHOT RECORD

LOCAL CHECK PASS

NOT INDEPENDENTLY VERIFIED

Purpose and scope

This record documents one completed Prompt 8 submission against official Google Research IFEval item 1508. It preserves the exact scored response, provenance, local checks, and claim boundaries. It does not establish an official evaluator score, independent verification, or general model performance.

ROUTE / RUNTIME BOUNDARY

Requested route: FastPath 5.6 / Luna Max. Observed runtime: Codex; exact model identifier is not exposed in this runtime.

Run Summary

Benchmark	Google Research Instruction-Following Evaluation (IFEval)
Benchmark version	Official repository master snapshot accessed 2026-08-26; exact commit not captured
Task / item	Item key 1508
Prompt series	Prompt 8 of the Grounding DI benchmark series
Official requirements	Required keywords forests and riddle; zero comma characters
Requested model / route	FastPath 5.6 / Luna Max
Observed runtime identity	Codex; exact model identifier is not exposed in this runtime
Run timestamp	2026-08-26 09:43:10 EDT
Permitted tools	Official-source web verification and transparent local checking
Actual tools used	Official-source web verification; private preflight; local check; PDF generation and render QA
Preflight corrections	0
Evaluator	LOCAL REIMPLEMENTATION — NOT OFFICIAL EVALUATOR
Result	2/2 instruction-level constraints passed; 1/1 prompt-level strict pass (local result)
Evidence status	NOT INDEPENDENTLY VERIFIED

Official Benchmark Provenance

The official item was checked before submission. The raw dataset returns item key 1508 with the exact prompt and instruction IDs keywords:existence and punctuation:no_comma. The repository checker source and paper were also consulted.

Controlling official sources

- **Official repository / task package**
https://github.com/google-research/google-research/tree/master/instruction_following_eval
- **Official dataset / item record**
https://raw.githubusercontent.com/google-research/google-research/master/instruction_following_eval/data/input_data.jsonl
- **Official instruction checker source**
https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions.py
- **Official paper**
<https://arxiv.org/abs/2311.07911>

VERIFICATION RESULT

MATCH - item 1508 in the official dataset matches the prompt below and encodes the two stated constraints. The official evaluator itself was not executable in this workspace and was not run.

EXACT BENCHMARK PROMPT

Write a haiku about a lion that includes the keywords "forests" and "riddle". Refrain from using commas in your haiku.

Prompt SHA-256: dbb32d584afad9f0dc9375f0e741729980881f9c2e3694f2295dd4475d9ea5

Preserved as the scored task; no prompt repair or silent reinterpretation performed.

Final Submitted Benchmark Response

The block below is the complete scored response. It contains only the haiku and is preserved exactly as finalized.

EXACT SCORED RESPONSE

calm lion at dusk

forests guard a dark riddle

moonlight crowns the mane

Response SHA-256: e18cacaa8239b78f2e9ab2218c35240afdc2cee78a4aa5bbd2911b9bc02ba574

Evaluation and Audit Trail

EVALUATOR STATUS

LOCAL REIMPLEMENTATION — NOT OFFICIAL EVALUATOR; official evaluator unavailable / not executed

Constraint / scoring results

Requirement	Official Standard	Observed	Result
keywords:existence	forests and riddle	Both keywords present	PASS
punctuation:no_comma	Zero commas	Comma count = 0	PASS
Instruction-level score	2 encoded constraints	2 passed	2/2
Prompt-level strict score	Both constraints pass	Both pass	PASS
Final recorded score	1/1 prompt-level strict pass	Local result	1/1

Audit Trail and Record Integrity

Manual / secondary check

- The response contains both exact required keywords, forests and riddle: PASS.
- The response contains zero comma characters: PASS.
- The scored block contains only the haiku and no audit material: PASS.

Defects, ambiguities, and limitations

- Official evaluator status: no executable official evaluator was available in this workspace; it was not run.
- Official-versus-local discrepancy: NOT MEASURABLE because the official evaluator did not execute.
- The local check mirrors the two encoded constraints and is not presented as an official benchmark score.
- This item tests keyword and punctuation compliance; it does not establish general intelligence or universal performance.
- Requested route and observed runtime are reported separately; the exact underlying model identifier is unavailable.

EVIDENCE STATUS

NOT INDEPENDENTLY VERIFIED

Evidence consists of official-source retrieval, this preserved record, and a transparent local check. No independent reviewer or hosted official score is included.

Record integrity

- Benchmark source used: official IFEval dataset and controlling repository sources listed on page 2.
- Final response preserved exactly as submitted in the scored-response block.
- Evaluator actually used: transparent local reimplementations only; no official evaluator claim.
- Preflight corrections disclosed: 0. Private inspection occurred before finalization.
- No failed completed submission was hidden; no post-score response change was represented.
- Classification: INTERNAL ONE-SHOT RECORD / LOCAL CHECK PASS / NOT INDEPENDENTLY VERIFIED.

Response SHA-256: e18caca8239b78f2e9ab2218c35240afdc2cee78a4aa5bbd2911b9bc02ba574