

GROUNDING DI OS BENCHMARK RECORD

Grounding DI OS Benchmark Record

Google Research IFEval | item key 1418 | Prompt 9

REQUESTED ROUTE: FastPath 5.6 / Luna Max

INTERNAL ONE-SHOT RECORD

LOCAL CHECK PASS

NOT INDEPENDENTLY VERIFIED

Purpose and scope

This record documents one completed Prompt 9 submission against official Google Research IFEval item 1418. It preserves the exact scored response, provenance, local checks, and claim boundaries. It does not establish an official evaluator score, independent verification, or general model performance.

ROUTE / RUNTIME BOUNDARY

Requested route: FastPath 5.6 / Luna Max. Observed runtime: Codex; exact model identifier is not exposed in this runtime.

Run Summary

Benchmark	Google Research Instruction-Following Evaluation (IFEval)
Benchmark version	Official repository master snapshot accessed 2026-08-26; exact commit not captured
Task / item	Item key 1418
Prompt series	Prompt 9 of the Grounding DI benchmark series
Official requirements	No commas; sentence count at least 30 and less than 31
Requested model / route	FastPath 5.6 / Luna Max
Observed runtime identity	Codex; exact model identifier is not exposed in this runtime
Run timestamp	2026-08-26 09:50:11 EDT
Permitted tools	Official-source web verification and transparent local checking
Actual tools used	Official-source web verification; private preflight; local check; PDF generation and render QA
Preflight corrections	0
Evaluator	LOCAL REIMPLEMENTATION — NOT OFFICIAL EVALUATOR
Result	3/3 instruction-level constraints passed; 1/1 prompt-level strict pass (local result)
Evidence status	NOT INDEPENDENTLY VERIFIED

Official Benchmark Provenance

The official item was checked before submission. The raw dataset returns item key 1418 with the exact prompt and three instruction entries: punctuation:no_comma and two separate length_constraints:number_sentences thresholds. The checker and sentence-count utility sources were also consulted.

Controlling official sources

- **Official repository / task package**
https://github.com/google-research/google-research/tree/master/instruction_following_eval
- **Official dataset / item record**
https://raw.githubusercontent.com/google-research/google-research/master/instruction_following_eval/data/input_data.jsonl
- **Official instruction checker source**
https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions.py
- **Official sentence-count utility**
https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions_util.py
- **Official paper**
<https://arxiv.org/abs/2311.07911>

VERIFICATION RESULT

MATCH - item 1418 in the official dataset matches the prompt below and encodes the three stated constraints. The official evaluator was not executable in this workspace because its NLTK sentence-tokenizer dependency was unavailable.

EXACT BENCHMARK PROMPT

Write a 30-line poem with short sentences without any comma. Each line should contain exactly one sentence. Make sure that you put the right punctuation at the end of each line. Your entire response should contain the poem only.

Prompt SHA-256: 4d9439084968bd824aada27f4984788c8d2275a91c5f624bfa5c94d4e949a52a

Preserved as the scored task; no prompt repair or silent reinterpretation performed.

Final Submitted Benchmark Response

The block below is the complete scored response. It contains only the poem and is preserved exactly as finalized.

EXACT SCORED RESPONSE

Morning opens softly.
A blue bird sings.
Dew rests on leaves.
Tall grass holds the light.
A small path bends.
Footsteps wake the gravel.
Clouds drift above.
The river keeps moving.
Stones shine near shore.
A fox watches quietly.
Then it slips away.
Wind brushes the reeds.
The reeds bow low.
Sun warms the water.
Fish flash below.
A heron lifts off.
Its wings beat slowly.
Evening enters the valley.
Shadows cross the field.
A bell sounds far away.
The old gate swings.
Night gathers at last.
Stars appear one by one.
The moon climbs higher.
A lantern glows nearby.
The house grows still.
Dreams wait indoors.
The river hums softly.
Dawn will come again.
The world rests easy.

Response SHA-256: 9a813eca2f80af6470d35f461c85b4be5994a0799752455648c81f00b7928fc6

Evaluation and Audit Trail

EVALUATOR STATUS

LOCAL REIMPLEMENTATION — NOT OFFICIAL EVALUATOR; official evaluator unavailable / not executed

Constraint / scoring results

REQUIREMENT	OFFICIAL STANDARD	OBSERVED	RESULT
punctuation:no_comma	Zero commas	Comma count = 0	PASS
number_sentences >= 30	At least 30	Equivalent count = 30	PASS
number_sentences < 31	Fewer than 31	Equivalent count = 30	PASS
Instruction-level score	3 encoded constraints	3 passed	3/3
Prompt-level strict score	All three pass	All pass	PASS
Final recorded score	1/1 prompt-level strict pass	Local result	1/1

Audit Trail and Record Integrity

Manual / secondary check

- The response contains exactly 30 non-empty lines: PASS.
- Each line contains one short sentence ending in a period: PASS.
- The response contains zero comma characters: PASS.
- The scored block contains only poem content and no audit material: PASS.

Defects, ambiguities, and limitations

- Official evaluator status: the evaluator package was not executable because the required NLTK sentence-tokenizer dependency was unavailable; it was not run.
- Official-versus-local discrepancy: NOT MEASURABLE because the official evaluator did not execute.
- The local equivalent counts one terminal period per line for this deliberately simple response and mirrors the official threshold arithmetic; it is not presented as an official score.
- The 30-line and one-sentence-per-line checks are manual task checks unless separately encoded by the official evaluator.
- Requested route and observed runtime are reported separately; the exact underlying model identifier is unavailable.

EVIDENCE STATUS

NOT INDEPENDENTLY VERIFIED

Evidence consists of official-source retrieval, this preserved record, and a transparent local check. No independent reviewer or hosted official score is included.

Record integrity

- Benchmark source used: official IFEval dataset, checker, sentence utility, and paper listed on page 2.
- Final response preserved exactly as submitted in the scored-response block.
- Evaluator actually used: transparent local reimplementations only; no official evaluator claim.
- Preflight corrections disclosed: 0. Private inspection occurred before finalization.
- No failed completed submission was hidden; no post-score response change was represented.
- Classification: INTERNAL ONE-SHOT RECORD / LOCAL CHECK PASS / NOT INDEPENDENTLY VERIFIED.

Response SHA-256: 9a813eca2f80af6470d35f461c85b4be5994a0799752455648c81f00b7928fc6