

GROUNDING DI OS BENCHMARK RECORD

Grounded DI OS Benchmark Record

Google Research IFEval | item key 1512 | Prompt 10

REQUESTED ROUTE: FastPath 5.6 / Luna Max

INTERNAL ONE-SHOT RECORD

LOCAL CHECK PASS

NOT INDEPENDENTLY VERIFIED

Purpose and scope

This record documents one completed Prompt 10 submission against official Google Research IFEval item 1512. It preserves the exact scored response, provenance, local checks, and claim boundaries. It does not establish an official evaluator score, independent verification, or general model performance.

ROUTE / RUNTIME BOUNDARY

Requested route: FastPath 5.6 / Luna Max. Observed runtime: Codex; exact model identifier is not exposed in this runtime.

Run Summary

Benchmark	Google Research Instruction-Following Evaluation (IFEval)
Benchmark version	Official repository master snapshot accessed 2026-08-26; exact commit not captured
Task / item	Item key 1512
Prompt series	Prompt 10 of the Grounded DI benchmark series
Official requirements	At least 20 square-bracket placeholders; zero comma characters
Requested model / route	FastPath 5.6 / Luna Max
Observed runtime identity	Codex; exact model identifier is not exposed in this runtime
Run timestamp	2026-08-26 10:36:14 EDT
Permitted tools	Official-source web verification and transparent local checking
Actual tools used	Official-source web verification; private preflight; local check; PDF generation and render QA
Preflight corrections	0
Evaluator	LOCAL REIMPLEMENTATION — NOT OFFICIAL EVALUATOR
Result	2/2 instruction-level constraints passed; 1/1 prompt-level strict pass (local result)
Evidence status	NOT INDEPENDENTLY VERIFIED

Official Benchmark Provenance

The official item was checked before submission. The raw dataset returns item key 1512 with instruction IDs detectable_content:number_placeholders and punctuation:no_comma. The repository checker source and paper were also consulted.

Controlling official sources

- **Official repository / task package**
https://github.com/google-research/google-research/tree/master/instruction_following_eval
- **Official dataset / item record**
https://raw.githubusercontent.com/google-research/google-research/master/instruction_following_eval/data/input_data.jsonl
- **Official instruction checker source**
https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions.py
- **Official paper**
<https://arxiv.org/abs/2311.07911>

VERIFICATION RESULT

MATCH - item 1512 in the official dataset matches the prompt below and encodes the two stated constraints. The official evaluator itself was not executable in this workspace and was not run.

EXACT BENCHMARK PROMPT

Generate a forum thread about several people waiting to hear the latest local news. All sentences should be short. Refrain from using any commas. Use placeholders to represent different usernames. Use square brackets for placeholders, like [username1], [username2]. Please include at least 20 placeholders in the thread.

Prompt SHA-256: 53ce549b5375a88591be137b9aa95b39f55750b586e10ad068807d609e37175b

Preserved as the scored task; no prompt repair or silent reinterpretation performed.

Final Submitted Benchmark Response

The block below is the complete scored response. It contains only the forum thread and is preserved exactly as finalized.

EXACT SCORED RESPONSE

[alexal]: Has anyone heard the local news?
[bruno2]: Not yet.
[casey3]: I am checking the radio.
[drew4]: The station is quiet.
[ellie5]: My neighbor saw police cars.
[frank6]: That sounds serious.
[grace7]: Do we know the location?
[harper8]: The report mentioned Main Street.
[ivan9]: I am waiting for details.
[jordan10]: The town page says nothing.
[kail1]: The phones are buzzing.
[logan12]: I will check the county feed.
[morgan13]: Please share updates here.
[nora14]: The weather looks calm.
[owen15]: Traffic is moving slowly.
[piper16]: A news van just passed.
[quinn17]: The station posted a warning.
[river18]: The message is still vague.
[sage19]: I hope everyone is safe.
[taylor20]: The latest report should arrive soon.
[uma21]: I will keep watching.
[victor22]: Please stay near your phones.
[wren23]: I found a short update.
[xavier24]: It says officials are speaking.
[yara25]: We should wait for facts.
[zane26]: I will post the next update.

Response SHA-256: dcbf9d9d076a24d3801755b8175ca79284bb13626ef95dad118ad369dd2c493b

Evaluation and Audit Trail

EVALUATOR STATUS

LOCAL REIMPLEMENTATION — NOT OFFICIAL EVALUATOR; official evaluator unavailable / not executed

Constraint / scoring results

REQUIREMENT	OFFICIAL STANDARD	OBSERVED	RESULT
detectable_content:number_placeholders	At least 20	Placeholder count = 26	PASS
punctuation:no_comma	Zero commas	Comma count = 0	PASS
Instruction-level score	2 encoded constraints	2 passed	2/2
Prompt-level strict score	Both constraints pass	Both pass	PASS

Final recorded score: 1/1 prompt-level strict pass (local result)

Audit Trail and Record Integrity

Manual / secondary check

- The response contains 26 square-bracket placeholder matches and therefore at least 20: PASS.
- The response contains zero comma characters: PASS.
- The forum-thread lines remain short by manual check: PASS.
- The scored block contains only forum-thread content and no audit material: PASS.

Defects, ambiguities, and limitations

- Official evaluator status: no executable official evaluator was available in this workspace; it was not run.
- Official-versus-local discrepancy: NOT MEASURABLE because the official evaluator did not execute.
- The local check mirrors the official placeholder regex and comma rule and is not presented as an official benchmark score.
- The short-sentence requirement is a manual task check and is not separately encoded in the official instruction list.
- Requested route and observed runtime are reported separately; the exact underlying model identifier is unavailable.

EVIDENCE STATUS

NOT INDEPENDENTLY VERIFIED

Evidence consists of official-source retrieval, this preserved record, and a transparent local check. No independent reviewer or hosted official score is included.

Record integrity

- Benchmark source used: official IFEval dataset and controlling repository sources listed on page 2.
- Final response preserved exactly as submitted in the scored-response block.
- Evaluator actually used: transparent local reimplementations only; no official evaluator claim.
- Preflight corrections disclosed: 0. Private inspection occurred before finalization.
- No failed completed submission was hidden; no post-score response change was represented.
- Classification: INTERNAL ONE-SHOT RECORD / LOCAL CHECK PASS / NOT INDEPENDENTLY VERIFIED.

Response SHA-256: dcbf9d9d076a24d3801755b8175ca79284bb13626ef95dad118ad369dd2c493b