

GROUNDING DI OS BENCHMARK RECORD

Grounded DI OS Benchmark Record

Google Research IFEval | Prompt 13 | Item key 1609

RECORD STATUS

INTERNAL

EVALUATION

LOCAL CHECK PASS

EVIDENCE

UNVERIFIED

Purpose and scope

This record documents one attempted response to official Google Research IFEval item 1609 under the requested FastPath 5.6 / Luna Max route. It establishes the exact source item, preserved submitted response, evaluator status, local constraint results, and evidence boundary. It does not establish an official IFEval score, independent verification, or general model performance.

Run summary

Benchmark	Google Research IFEval
Benchmark version	Official master dataset retrieved 2026-08-26
Task / item	Prompt 13 / item key 1609
Official requirements	Four paragraphs; paragraph 1 starts with weekend; outer double quotation marks
Requested route	FastPath 5.6 / Luna Max
Observed runtime identity	Codex; exact model identifier not exposed by the runtime
Run timestamp	2026-08-26 12:58:31 EDT
Permitted tools	Official web sources; local evaluator if official evaluator unavailable
Actual tools used	Web source verification; local Python equivalent check; ReportLab; Poppler
Preflight corrections	0
Evaluator	LOCAL REIMPLEMENTATION — NOT OFFICIAL EVALUATOR
Result	2/2 encoded constraints passed; 1/1 prompt-level strict pass
Evidence status	NOT INDEPENDENTLY VERIFIED
Prompt SHA-256	0fb200ca0abfc8a9a492fd0500a7d97bfbfa25249fa9bb0944cb1ae1c5ca4fe3
Response SHA-256	87095362490046ca07568897e45a3873f81c35a7ea9c291f1390b2761e7fa0ed

CLAIM BOUNDARY

The local result is a documented reimplementation of the official encoded checks. It is not an official benchmark score. The exact model identifier was not exposed by the runtime.

Official Benchmark Provenance

SOURCE	URL	STATUS
Official repository	https://github.com/google-research/google-research/tree/master/instruction_following_eval	VERIFIED
Official dataset	https://raw.githubusercontent.com/google-research/google-research/master/instruction_following_eval/data/input_data.jsonl	ITEM 1609 MATCHED
Official checker	https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions.py	BEHAVIOR REVIEWED
Official paper	https://arxiv.org/abs/2311.07911	REFERENCE

Verification result

The live official dataset entry for key 1609 matched the supplied benchmark prompt exactly. Its instruction list contains length_constraints:nth_paragraph_first_word with first_word=weekend, num_paragraphs=4, nth_paragraph=1, plus startend:quotation. The checker source was reviewed for the paragraph split, first-word normalization, and quotation-boundary behavior. No source conflict was observed.

Exact benchmark prompt

What's a good way to ask Sonia out? Please reply with exactly 4 paragraphs and separate each paragraph with two new lines. Put double quotation marks around your entire response. The very first paragraph must start with the word "weekend".

The prompt above is reproduced as the scored task. Audit text appears only in this record and was not included in the scored response.

Final Submitted Benchmark Response

The following block is the complete scored response. It is preserved as submitted, including the outer quotation marks and the three blank-line paragraph separators.

"weekend, choose a relaxed moment and ask Sonia directly if she'd like to get coffee with you.
Keep the invitation simple and specific.

Suggest a day and a low-pressure plan. A short coffee date or walk makes the invitation easy to accept and easy to decline.

Be warm and confident without making the question feel like a test. Give her room to say yes, no, or think about it.

If she agrees, confirm the time and place. If she declines, respond kindly and keep your dignity."

Evaluation and Audit Trail

EVALUATOR STATUS

Official IFEval evaluator: NOT EXECUTED. The local runtime did not contain the official evaluator package or its dependencies (nltk, immutabledict, or an installed instruction_following_eval module).

LOCAL REIMPLEMENTATION — NOT OFFICIAL EVALUATOR

REQUIREMENT	OFFICIAL STANDARD	OBSERVED	RESULT
length_constraints:nth_paragraph_first_word	Exactly 4 paragraphs; paragraph 1 begins with weekend	4 paragraphs; first normalized word: weekend; separators: 3	PASS
startend:quotation	Stripped response begins and ends with straight double quotes	First=""; final=""	PASS

SCORE	RECORDED RESULT
Instruction-level score	2 / 2 encoded constraints passed locally
Prompt-level strict score	PASS
Final recorded score	1 / 1 prompt-level strict pass
Official-versus-local discrepancy	NOT MEASURABLE because the official evaluator was not executed

Manual / Secondary Check

CHECK	RESULT
Exactly four non-empty paragraphs	PASS
Exactly three paragraph separators	PASS
First content word is weekend	PASS
Outer quotation marks are straight double quotes	PASS
Response contains only the requested answer	PASS

Preflight and limitations

Private preflight corrections: 0. The response was drafted, checked, and finalized without a correction. A material evaluator nuance is disclosed: the outer opening quote is the literal first character of the complete response, while the official paragraph-first-word checker strips a leading quote before comparing the first word. The first paragraph's first content word is weekend.

The official evaluator was not executable in this workspace because its runtime dependencies were unavailable. The local check mirrors the relevant source behavior but cannot measure an official/local discrepancy. No independent reviewer or external evaluator verified this run.

Evidence status

NOT INDEPENDENTLY VERIFIED

Evidence consists of the official source retrieval, this preserved prompt and response, the local equivalent check, the runtime/tool record, and the generated PDF. No external certification or independent reproduction is claimed.

Record integrity

- Official benchmark source and item key were recorded.
- The final submitted response is preserved exactly in a separated block.
- The evaluator actually used is identified; the official evaluator is not represented as executed.
- Preflight corrections are disclosed.
- No post-score response change was made and no completed failed submission was hidden.
- The record does not claim certification, independence, universality, or representative model performance.

Record identity hashes: prompt 0fb200ca0abfc8a9a492fd0500a7d97bfbfa25249fa9bb0944cb1ae1c5ca4fe3; response 87095362490046ca07568897e45a3873f81c35a7ea9c291f1390b2761e7fa0ed. Run timestamp: 2026-08-26 12:58:31 EDT.

FINAL STATUS: INTERNAL ONE-SHOT RECORD | LOCAL CHECK PASS | LOCAL REIMPLEMENTATION — NOT OFFICIAL EVALUATOR | NOT INDEPENDENTLY VERIFIED