

GROUNDING DI OS BENCHMARK RECORD

Grounded DI OS Benchmark Record

Google Research IFEval | Prompt 14 | Item key 1705

RECORD STATUS

INTERNAL

EVALUATION

LOCAL CHECK PASS

EVIDENCE

UNVERIFIED

Purpose and scope

This record documents one attempted response to official Google Research IFEval item 1705 under the requested FastPath 5.6 / Luna Max route. It preserves the source item, exact submitted response, checker behavior, local result, evidence status, and claim boundaries. It does not establish an official IFEval score, independent verification, or general model performance.

Run summary

Benchmark	Google Research IFEval
Benchmark version	Official master dataset retrieved 2026-08-26
Task / item	Prompt 14 / item key 1705
Official requirements	Postscript beginning P.P.S.; exact ending phrase Hope you agree with me.
Model / route	FastPath 5.6 / Luna Max
Run timestamp	2026-08-26 13:33:45 EDT
Permitted tools	Official source verification; local equivalent check if needed
Actual tools used	GitHub source verification; public web verification; local Python check; ReportLab; Poppler
Preflight corrections	0
Evaluator	LOCAL REIMPLEMENTATION — NOT OFFICIAL EVALUATOR
Result	2 / 2 encoded constraints passed locally; 1 / 1 prompt-level strict pass
Evidence status	NOT INDEPENDENTLY VERIFIED
Prompt SHA-256	fdb2ba9342fc3d8bab4be8184dbde8bddc088914eaecd9ea5e25293f6a314324
Response SHA-256	fcee59fefeff0ce296a6a2566aeb301de6461fd87dee8e01e596a16f95e4c2c4

CLAIM BOUNDARY

The official evaluator was not executed in this workspace. The recorded pass is a local equivalent result and is not an official benchmark score. The route label is recorded directly as FastPath 5.6 / Luna Max per the run conditions.

Official Benchmark Provenance

SOURCE	URL	STATUS
Official repository	https://github.com/google-research/google-research/tree/master/instruction_following_eval	VERIFIED
Official dataset	https://raw.githubusercontent.com/google-research/google-research/master/instruction_following_eval/data/input_data.jsonl	ITEM 1705 MATCHED
Official checker	https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions.py	BEHAVIOR REVIEWED
Official paper	https://arxiv.org/abs/2311.07911	REFERENCE

Verification result

The live official dataset entry for key 1705 matched the supplied benchmark prompt exactly. It identifies the two encoded constraints as detectable_content:postscript with postscript_marker=P.P.S and startend:end_checker with end_phrase=Hope you agree with me. The official checker source was reviewed at commit 807d4a2f4120 for its P.P.S pattern and end-phrase normalization. No source conflict was observed.

Exact benchmark prompt

Write a professional email that you could send to ask your boss for a raise. At the end of your response, explicitly add a postscript starting with P.P.S

Finish your entire response with this exact phrase: Hope you agree with me.

The prompt above is the scored task. Audit text was excluded from the scored response.

Final Submitted Benchmark Response

The following block is the complete scored response. It is preserved exactly, including paragraph breaks and the terminal phrase.

Dear Manager,

I am writing to request a conversation about my compensation. I have valued the opportunity to contribute to the team and would appreciate discussing whether my current salary reflects my responsibilities and recent contributions.

I would be grateful for a convenient time to review my performance, responsibilities, and the possibility of a raise. I remain committed to producing strong work and supporting our team's goals.

P.P.S. Hope you agree with me.

Evaluation and Audit Trail

EVALUATOR STATUS
Official IFEval evaluator: NOT EXECUTED. The official package and its dependencies were not installed in this workspace.

LOCAL REIMPLEMENTATION — NOT OFFICIAL EVALUATOR

Requirement	Official Standard	Observed	Result
detectable_content:postscript	P.P.S marker detected by the official checker pattern	1 match; postscript at response end	PASS
startend:end_checker	After strip, quote strip, and lowercase normalization, response ends with the exact phrase	Normalized response ends with: hope you agree with me.	PASS

Score	Recorded Result
Instruction-level score	2 / 2 encoded constraints passed locally
Prompt-level strict score	PASS
Final recorded score	1 / 1 prompt-level strict pass
Official-versus-local discrepancy	NOT MEASURABLE because the official evaluator was not executed

Manual / Secondary Check

MANUAL TASK CHECK	RESULT
Professional email requesting a salary discussion or raise	PASS
Postscript is the final section and begins with P.P.S.	PASS
Final response ends exactly with Hope you agree with me.	PASS
Scored response contains email content only and no audit material	PASS

Official encoded constraints versus manual task checks

Official encoded constraints are the two rows in the Evaluation and Audit Trail table: the postscript detector and the end-phrase checker. The professional-email subject, email-only presentation, and postscript placement were recorded as manual task checks and were not added to the official IFEval score.

Preflight and blinding status

Private preflight corrections: 0. The response was drafted, inspected, and finalized without a correction.

INTERNAL RESPONSE-ONLY BLIND CHECK: YES

For the local check, the checker received only the response text and constraint parameters; the audit narrative and final score presentation were kept outside the scored response. This is an internal response-only blind boundary. It is not independent third-party blinding, and no external blind evaluator was used.

Defects, ambiguities, and limitations

The official evaluator was not executable because the required package/dependencies were unavailable. Consequently official-versus-local discrepancy is NOT MEASURABLE. No benchmark defect or source conflict was observed. No independent reviewer or external evaluator verified this run. The blinding statement is limited to the internal response-only check described above.

Evidence status

NOT INDEPENDENTLY VERIFIED

Evidence consists of the official source retrieval, the preserved prompt and response, the local equivalent check, the runtime/tool record, and this PDF. No certification, official score, universal claim, or representative-performance claim is made.

Record integrity

- Official benchmark source and item key were recorded.
- The final submitted response is preserved exactly in a separated block.
- The evaluator actually used is identified; the official evaluator is not represented as executed.
- Preflight corrections are disclosed.
- No post-score response change was made and no completed failed submission was hidden.
- The record does not claim certification, independence, universality, or representative model performance.

Hashes and timestamp are preserved in the Run summary table on page 1; the source commit is recorded in the provenance section.

FINAL STATUS: INTERNAL ONE-SHOT RECORD | LOCAL CHECK PASS | LOCAL REIMPLEMENTATION — NOT OFFICIAL EVALUATOR | NOT INDEPENDENTLY VERIFIED