

GROUNDING DI OS BENCHMARK RECORD

Grounded DI OS Benchmark Record

Google Research IFEval | Prompt 15 | Item key 1713

RECORD STATUS

INTERNAL

EVALUATION

LOCAL CHECK PASS

EVIDENCE

UNVERIFIED

Purpose and scope

This record documents one attempted response to official Google Research IFEval item 1713 under the requested FastPath 5.6 / Luna Max route. It preserves the source item, exact submitted response, checker behavior, local result, evidence status, and claim boundaries. It does not establish an official IFEval score, independent verification, or general model performance.

Run summary

Benchmark	Google Research IFEval
Benchmark version	Official master dataset retrieved 2026-08-26
Task / item	Prompt 15 / item key 1713
Official requirements	disappointed at least 2 times; at least 6 Markdown highlights
Model / route	FastPath 5.6 / Luna Max
Run timestamp	2026-08-26 14:05:14 EDT
Permitted tools	Official source verification; local equivalent check if needed
Actual tools used	GitHub source verification; public web verification; local Python check; ReportLab; Poppler
Preflight corrections	0
Evaluator	LOCAL REIMPLEMENTATION — NOT OFFICIAL EVALUATOR
Result	2 / 2 encoded constraints passed locally; 1 / 1 prompt-level strict pass
Evidence status	NOT INDEPENDENTLY VERIFIED
Prompt SHA-256	534963b09479799d142df3a83f19e11a3164d287181fcfd7f9ce7df90fdc8647
Response SHA-256	1081bade2f9fd18784f24c54c5eaa33aba05c4de73cad065c3ff16bca2c37bf9

CLAIM BOUNDARY

The official evaluator was not executed in this workspace. The recorded pass is a local equivalent result and is not an official benchmark score. The route label is recorded directly as FastPath 5.6 / Luna Max.

Official Benchmark Provenance

SOURCE	URL	STATUS
Official repository	https://github.com/google-research/google-research/tree/master/instruction_following_eval	VERIFIED
Official dataset	https://raw.githubusercontent.com/google-research/google-research/master/instruction_following_eval/data/input_data.jsonl	ITEM 1713 MATCHED
Official checker	https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions.py	BEHAVIOR REVIEWED
Official paper	https://arxiv.org/abs/2311.07911	REFERENCE

Verification result

The live official dataset entry for key 1713 matched the supplied benchmark prompt exactly. It identifies keywords:frequency with keyword disappointed, relation at least, and frequency 2, plus detectable_format:number_highlighted_sections with num_highlights 6. The official checker source was reviewed at commit 807d4a2f4120 for case-insensitive keyword regex matching and non-empty Markdown highlight counting. No source conflict was observed.

Exact benchmark prompt

Write an interesting and funny article about the biology of a banana peel. In your response, the word disappointed should appear at least 2 times, and at least six section should be highlighted with markdown, i.e *banana peel*.

The prompt above is the scored task. Audit text was excluded from the scored response.

Final Submitted Benchmark Response

The following block is the complete scored response. It is preserved exactly; no title, score, citation, or audit material was included in the submission.

A banana peel is the fruit's built-in raincoat, helmet, and emergency comedy prop. Its *waxy cuticle* slows water loss while its *cell walls* keep the skin flexible enough to bend without filing a complaint.

Inside, *vascular bundles* carry water and sugars through the peel. *Chloroplasts* give young peels their green color; as the fruit ripens, pigments fade and warmer yellow tones take over. The peel is basically running a tiny color-coded maturity dashboard.

Ethylene gas helps coordinate ripening, telling cells to soften the fruit and convert starches into sweeter sugars. Meanwhile *lignin* and other fibers provide structure, although eventually the peel becomes so soft that gravity starts acting like a personal enemy.

The *black spots* are not tiny bruises from a disappointed banana committee. They reflect localized tissue breakdown and pigment changes. The peel may look disappointed, but it is biologically busy: protecting the fruit, signaling ripeness, and preparing for its final career as compost.

Evaluation and Audit Trail

EVALUATOR STATUS

Official IFEval evaluator: NOT EXECUTED. The official package and its dependencies were not installed in this workspace.

LOCAL REIMPLEMENTATION — NOT OFFICIAL EVALUATOR

REQUIREMENT	OFFICIAL STANDARD	OBSERVED	RESULT
keywords:frequency	disappointed appears at least 2 times; case-insensitive regex	2 matches	PASS
detectable_format:number_highlighted_sections	At least 6 non-empty Markdown highlight spans	7 valid spans	PASS

SCORE	RECORDED RESULT
Instruction-level score	2 / 2 encoded constraints passed locally
Prompt-level strict score	PASS
Final recorded score	1 / 1 prompt-level strict pass
Official-versus-local discrepancy	NOT MEASURABLE because the official evaluator was not executed

Manual / Secondary Check

MANUAL TASK CHECK	RESULT
Interesting and funny article	PASS
Article concerns banana-peel biology	PASS
Response contains article content only and no audit material	PASS
At least six clear highlighted sections are present	PASS

Official encoded constraints versus manual task checks

The official score consists only of the keyword-frequency and highlighted-section constraints. Article quality, humor, banana-peel biology, and article-only presentation are recorded separately as manual checks and are not added to the IFEval score.

Preflight and blinding status

Private preflight corrections: 0. The response was drafted, inspected, and finalized without a correction.

INTERNAL RESPONSE-ONLY BLIND CHECK: YES

For the local check, the checker received only the response text and constraint parameters; the audit narrative and score presentation were kept outside the scored response. This is an internal response-only blind boundary, not independent third-party verification.

Defects, ambiguities, and limitations

The official evaluator was not executable because the required package and dependencies were unavailable. Official-versus-local discrepancy is therefore NOT MEASURABLE. No benchmark defect or source conflict was observed. No independent reviewer or external evaluator verified this run.

Evidence status

NOT INDEPENDENTLY VERIFIED

Evidence consists of the official source retrieval, preserved prompt and response, local equivalent check, tool record, and this PDF. No certification, universal claim, representative-performance claim, or official score is made.

Record integrity

- Official benchmark source and item key were recorded.
- The final submitted response is preserved exactly in a separated block.
- The evaluator actually used is identified; the official evaluator is not represented as executed.
- Preflight corrections are disclosed.
- No post-score response change was made and no completed failed submission was hidden.
- The record does not claim certification, independence, universality, or representative model performance.

Hashes and timestamp are preserved in the Run summary table on page 1; the source commit is recorded in the provenance section.

FINAL STATUS: INTERNAL ONE-SHOT RECORD | LOCAL CHECK PASS | LOCAL REIMPLEMENTATION — NOT OFFICIAL EVALUATOR | NOT INDEPENDENTLY VERIFIED