

GROUNDING DI OS BENCHMARK RECORD

Grounded DI OS Benchmark Record

Google Research IFEval | Prompt 17 | Item key 1936

RECORD STATUS

INTERNAL

EVALUATION

LOCAL CHECK PASS

EVIDENCE

UNVERIFIED

Purpose and scope

This record documents one attempted response to official Google Research IFEval item 1936 under the FastPath 5.6 / Luna Max route. It preserves the source item, exact submitted response, checker behavior, local result, evidence status, and claim boundaries. It does not establish an official IFEval score, independent verification, or general model performance.

Run summary

Benchmark	Google Research IFEval
Benchmark version	Official master dataset retrieved 2026-08-26
Task / item	Prompt 17 / item key 1936
Official requirements	Exactly 6 counted bullet lines; no forbidden whole words
Model / route	FastPath 5.6 / Luna Max
Run timestamp	2026-08-26 14:43:28 EDT
Permitted tools	Official source verification; local equivalent check if needed
Actual tools used	GitHub source verification; public web source check; local Python check; ReportLab; Poppler
Preflight corrections	0
Evaluator	LOCAL REIMPLEMENTATION — NOT OFFICIAL EVALUATOR
Result	2 / 2 encoded constraints passed locally; 1 / 1 prompt-level strict pass
Evidence status	NOT INDEPENDENTLY VERIFIED
Prompt SHA-256	89953cf6612debe5a1284ce5ce9771eb46e10137cff57f8429601c7638be8065
Response SHA-256	a187c162c42ebb1487df61cfb40161b54ba6062478bdcc1ac027e21b7c623aba

CLAIM BOUNDARY

The official evaluator was not executed in this workspace. The recorded pass is a local equivalent result and is not an official benchmark score. The model / route is recorded directly as FastPath 5.6 / Luna Max.

Official Benchmark Provenance

SOURCE	URL	STATUS
Official repository	https://github.com/google-research/google-research/tree/master/instruction_following_eval	VERIFIED
Official dataset	https://raw.githubusercontent.com/google-research/google-research/master/instruction_following_eval/data/input_data.jsonl	ITEM 1936 MATCHED
Official checker	https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions.py	BEHAVIOR REVIEWED
Official paper	https://arxiv.org/abs/2311.07911	REFERENCE

Verification result

The current official dataset entry for key 1936 matched the supplied benchmark prompt exactly. It identifies keywords:forbidden_words with bad and underperform, plus detectable_format:number_bullet_lists with an exact count of 6. The official checker source was reviewed at the current master file retrieval (dataset blob cbe52f6eecf3 and checker blob 0007359bd37c) for case-insensitive whole-word matching and the combined asterisk/hyphen bullet-count logic. No source conflict was observed.

Exact benchmark prompt

Write a rubric, in the form of a list of bullet points, for evaluating the performance of a customer service representative. Your answer must not include keywords ['bad', 'underperform'] and must contain exactly 6 bullet points in the following form:
 * Bullet point 1
 * Bullet point 2
 * Bullet point 3
 * Bullet point 4
 * Bullet point 5
 * Bullet point 6

The prompt above is the scored task. Audit text was excluded from the scored response.

Final Submitted Benchmark Response

The following block is the complete scored response. It is preserved exactly; no title, score, citation, or audit material was included in the submission.

- * Clarity: Explain each answer in plain language.
 * Courtesy: Maintain a calm and respectful tone.
 * Accuracy: Give reliable information without guessing.
 * Ownership: Follow each request through to completion.
 * Responsiveness: Acknowledge concerns and set clear expectations.
 * Documentation: Record the outcome and the next action.

Evaluation and Audit Trail

EVALUATOR STATUS

Official IFEval evaluator: NOT EXECUTED. The official package and its dependencies were not installed in this workspace.

LOCAL REIMPLEMENTATION — NOT OFFICIAL EVALUATOR

Requirement	Official Standard	Observed	Result
keywords:forbidden_words	No whole-word match for either forbidden term; case-insensitive	bad: absent; underperform: absent	PASS
detectable_format:number_bullet_lists	Asterisk and hyphen bullet matches total exactly 6	6 asterisk + 0 hyphen = 6	PASS

Score	Recorded Result
Instruction-level score	2 / 2 encoded constraints passed locally
Prompt-level strict score	PASS
Final recorded score	1 / 1 prompt-level strict pass
Official-versus-local discrepancy	NOT MEASURABLE because the official evaluator was not executed

Manual / Secondary Check

MANUAL TASK CHECK	RESULT
Exactly six substantive bullets are present	PASS
Every bullet begins with the exact '* ' marker	PASS
Rubric evaluates customer-service representative performance	PASS
Response contains only rubric content and no audit material	PASS

Official encoded constraints versus manual task checks

The official score consists only of the forbidden-word and exact bullet-count constraints. Bullet-marker uniformity, substantive rubric quality, customer-service relevance, and rubric-only presentation are recorded separately as manual checks and are not added to the IFEval score.

Preflight and blinding status

Private preflight corrections: 0. The response was drafted, inspected, and finalized without a correction.

INTERNAL RESPONSE-ONLY BLIND CHECK: YES

For the local check, the checker received only the response text and fixed constraint parameters; the audit narrative and score presentation were kept outside the scored response. This is an internal response-only blind boundary and not independent third-party verification.

Defects, ambiguities, and limitations

The official evaluator was not executable because the required package and dependencies were unavailable. Official-versus-local discrepancy is therefore NOT MEASURABLE. No benchmark defect or source conflict was observed. No independent reviewer or external evaluator verified this run.

Evidence status

NOT INDEPENDENTLY VERIFIED

Evidence consists of the official source retrieval, preserved prompt and response, local equivalent check, tool record, and this PDF. No certification, universal claim, representative-performance claim, or official score is made.

Record integrity

- Official benchmark source and item key were recorded.
- The final submitted response is preserved exactly in a separated block.
- The evaluator actually used is identified; the official evaluator is not represented as executed.
- Preflight corrections are disclosed.
- No post-score response change was made and no completed failed submission was hidden.
- The record does not claim certification, independence, universality, or representative model performance.

Hashes and timestamp are preserved in the Run summary table on page 1; source blob identifiers are recorded in the provenance section.

FINAL STATUS: INTERNAL ONE-SHOT RECORD | LOCAL CHECK PASS | LOCAL REIMPLEMENTATION — NOT OFFICIAL EVALUATOR | NOT INDEPENDENTLY VERIFIED