

Grounded DI OS Benchmark Record

Prompt 21 of the Grounded DI cross-domain benchmark series — official Google Research IFEval item key 2704

INTERNAL ONE-SHOT RECORD · FASTPATH 5.6 / SOL MAX · OFFICIAL EVALUATOR PASS

1. Document identity and condition

Primary document	Grounded DI OS Benchmark Record
Benchmark-series position	Cross-domain Prompt 21
Requested condition	Grounded DI OS / FastPath 5.6 / SOL Max / Protocol A
Run controls	Tier 18 · DIP 0-166 · Invocation 140A / 140 · ΔH 0.0041 · Benchmark Mode
Benchmark identifier	Google Research IFEval item key 2704
Labels applied	INTERNAL ONE-SHOT RECORD; FASTPATH 5.6 / SOL MAX; OFFICIAL EVALUATOR PASS; LOCAL CHECK PASS
Labels not applied	certified; independently verified; universal; representative; official leaderboard score; proof of a hidden backend

2. Purpose and scope

This record documents one observed execution of official Google Research IFEval item key 2704 under the operator-declared FastPath 5.6 / SOL Max / Protocol A condition. It applies the one-shot record procedure established for the cross-domain series. Scope is limited to this prompt and to the three encoded IFEval constraints. A passing result is retained as observed evidence only and is not a claim of universal or deterministic performance.

The operator-supplied Grounded DI Runtime Control Package is treated as an evidence-harness specification. Its boundary remains HARNESS_SPECIFIED / NOT EXECUTED: runtime tests false; replay not executed; PackageReady false; not certified. No independent runtime-engine activation or backend identity is inferred from that package.

3. Official benchmark URLs

Series and audit repositories reviewed for prompt sequence and record architecture:

- <https://github.com/Grounded-DI/Benchmarks-Cross-Domain>
- <https://github.com/Grounded-DI/briefwise-di-performance-on-harvey-lab-benchmark-audit>

Official item and evaluator authority:

- https://github.com/google-research/google-research/tree/master/instruction_following_eval
- https://raw.githubusercontent.com/google-research/google-research/master/instruction_following_eval/data/input_data.jsonl
- https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions.py
- <https://arxiv.org/abs/2311.07911>

4. Exact official item (key 2704)

The official JSONL object was retrieved from input_data.jsonl. The raw preserved item line has SHA-256 9a625d8e902135fbb611cd3f5ddf80f2f4a80b3be8530628ab14ab4b4ff12d4.

```
{
  "key": 2704,
  "prompt": "Write a haiku about foolish behavior in the form of a question, for an audience of young readers. It should include the topic of not studying. Give exactly two different responses, separated by 6 asterisk symbols (*****), and include a title wrapped in double angular brackets, i.e. <<title>>. Do not use commas.",
  "instruction_id_list": [
    "combination:two_responses",
    "detectable_format:title",
    "punctuation:no_comma"
  ],
  "kwargs": [
    {},
    {},
    {}
  ]
}
```

5. Exact benchmark prompt (scored task only)

The following exact official prompt is the only task scored. Surrounding record text is audit material and is not part of the response.

```
Write a haiku about foolish behavior in the form of a question, for an audience of young readers. It should include the topic of not studying. Give exactly two different responses, separated by 6 asterisk symbols (*****), and include a title wrapped in double angular brackets, i.e. <<title>>. Do not use commas.
```

6. Observed runtime identity

Requested / session condition	FastPath 5.6 / SOL Max / Protocol A / Tier 18
Exposed assistant identity	Codex agent based on the GPT-5 family
Precise model snapshot / build	Unavailable — no hidden backend identifier or build hash was exposed
Route-claim boundary	SOL Max is the requested session condition; the unavailable backend build is not invented or independently attested
Host sandbox	Linux 6.18.35 x86_64 · Python 3.12.13 · glibc 2.39
Evidence-harness boundary	HARNESS_SPECIFIED / NOT EXECUTED · runtime tests false · replay not executed · PackageReady false · not certified

7. Timestamp and timezone

Official evaluation recorded	2026-09-01 15:54:49 UTC
New York local time	2026-09-01 11:54:49 EDT (UTC−4)
Timezone basis	America/New_York plus UTC from the sandbox clock

8. Tool and web-search status

The connected GitHub source was used to review the two requested Grounded DI repositories and to retrieve the official Google Research item and evaluator blobs. The response itself required no factual lookup. No prior Grounded DI answer, model answer, answer key, or scored response was retrieved to construct the submission.

The four executed Google Research source files were preserved as byte-identical local mirrors. Their Git blob identities are: evaluation_lib.py c87422822f9e1565ca7574ac0b3cee5389c9f44e; instructions.py 0007359bd37c7b9c1d9557f8937cb6586d2788a7; instructions_registry.py 6b42adcc4c8cc827810ab3efa8a06f02bb48c397; instructions_util.py d50702400738df4378fe286c0090a416f9a172fa.

Evaluator dependencies were installed into an isolated workspace path after lock. The locked response hash was checked before and after evaluation and remained unchanged. No GitHub repository was posted to or modified.

9. Draft and preflight status

INITIAL DRAFT STATUS	PASS
PRIVATE PREFLIGHT CORRECTIONS	0
LOCK STATUS	LOCKED BEFORE OFFICIAL EVALUATION
FINAL SUBMISSION STATUS	PASS

One complete candidate response was composed. Private preflight then checked the literal delimiter, distinct response parts, title syntax, comma absence, and manual task fit. Every check already held. No substantive correction was made. Ordinary generation before the first complete candidate existed is not counted as a correction.

10. Final submitted benchmark response

LOCKED SCORED RESPONSE — answer content only. This preserved submission was not modified after evaluator results were observed.

```
<<The Unread Lesson>>
Books wait by moonlight
Why skip your study for games?
Will tests wait for you?
*****
Notes sleep in backpacks
Can you learn without study?
Will play help on tests?
```

Locked file: 188 bytes; 8 visible lines; terminal newline present. SHA-256: a1038be30596dbe96544fb6da9f1927d40af25b76d07964ab2c64bd4f13ad420.

11. Official evaluator result

Official source from google-research/instruction_following_eval was executed against the locked response. This was evaluator execution rather than source inspection or a local substitute.

- Code paths executed
- 1. TwoResponsesChecker.check_following(response)
 - 2. TitleChecker.check_following(response)
 - 3. CommaChecker.check_following(response)
 - 4. evaluation_lib.test_instruction_following_strict(InputExample key=2704)
 - 5. evaluation_lib.test_instruction_following_loose on the same item-to-response mapping

Official constraint outputs

TwoResponsesChecker	True — follow
TitleChecker	True — follow
CommaChecker	True — follow
follow_instruction_list (strict)	[True, True, True]
follow_all_instructions (strict)	True
follow_instruction_list (loose)	[True, True, True]
follow_all_instructions (loose)	True
Official evaluator label	OFFICIAL EVALUATOR PASS

Execution environment: absl-py 2.5.0; immutabledict 4.3.1; langdetect 1.0.9; nltk 3.10.3. No tokenizer data was needed by the three relevant checker classes. The response hash was identical before and after both official paths.

12. Manual / local constraint check

LOCAL REIMPLEMENTATION — NOT OFFICIAL EVALUATOR. Results agree with the official code.

12.1 Two-response structure

Literal ***** count	1
Raw split-part count	2
Nonempty response parts	2
Parts distinct after strip	True — PASS
Part 1 normalized start	<<The Unread Lesson>> / Books wait by moonlight
Part 2 normalized start	Notes sleep in backpacks

12.2 Title and punctuation

Regex	<<[^\n]+>>
Exact title match	<<The Unread Lesson>>
Nonempty title interior	True — PASS
Literal comma count	0 — PASS

12.3 Manual task checks — separate from official score

Two different haiku responses	PASS — each response contains a distinct three-line poem
Manual syllable scan	Response 1: 5 / 7 / 5; Response 2: 5 / 7 / 5
Question form	PASS — each poem asks explicit questions
Foolish behavior / not studying	PASS — skipped study and choosing games or play are explicit
Young-reader suitability	PASS — simple school-centered vocabulary and no mature content
Audit material excluded	PASS — the scored block contains only title, poems, and required delimiter

13. Scores

Instruction-level score	3 / 3 encoded constraints passed
Prompt-level strict score	PASS — all encoded constraints passed
Prompt-level loose score	PASS — all encoded constraints passed
Final score	1 / 1 prompt-level strict pass
Official vs local discrepancy	None
Manual task checks	PASS — reported separately and not added to the official score

The 1/1 figure is the outcome of this single observed execution. It is not an official IFEval leaderboard score and is not independently verified by a third party.

14. Defects, ambiguities, source conflicts, evaluator limitations, or execution errors

No source conflict was found. Official item key 2704 contains the expected three instruction IDs: combination:two_responses, detectable_format:title, and punctuation:no_comma. The scored prompt was preserved without repair or paraphrase.

IFEval's official score covers only those three encoded checkers. It does not machine-score haiku form, syllable count, question quality, topic adequacy, or audience fit. Those surface requirements were manually inspected and remain separate from the official score. Manual English syllable counting is interpretive even where the selected lines admit conventional 5/7/5 counts.

Only item 2704 was executed; the full IFEval corpus was not run. Official evaluation_lib was invoked as a Python API from byte-matched source blobs. Missing packages were installed only after the response was locked. That dependency repair did not alter the response and did not affect the relevant checker semantics.

The precise hidden model snapshot and backend build were not exposed. Accordingly, SOL Max is recorded as the requested session condition and not as an independently attested backend build. No error affected the locked response or the evaluator outcome.

15. Evidence status

Official item identity	Verified from input_data.jsonl key 2704
Official checkers executed	Yes — TwoResponsesChecker; TitleChecker; CommaChecker
OFFICIAL EVALUATOR PASS	Yes — strict and loose paths on the locked response
LOCAL CHECK PASS	Yes — independent split, regex, punctuation, and manual review agree
LOCAL REIMPLEMENTATION	Transparency cross-check only — not the official evaluator
Independent third-party review	NOT INDEPENDENTLY VERIFIED
Runtime-control package	HARNESS_SPECIFIED / NOT EXECUTED · PackageReady false · not certified
Deterministic / universal claim	Not made

16. Record-integrity statement

This document is an INTERNAL ONE-SHOT RECORD of one execution under the requested FastPath 5.6 / SOL Max / Protocol A condition. Official IFEval item key 2704 and its exact prompt were frozen before response generation. The response printed in Section 10 is the one locked scored submission.

The first complete candidate passed private preflight. Substantive preflight corrections: 0. The locked response was hashed before official evaluation. It was not rewritten after evaluator outputs were seen. There was no retry after lock and no selection among multiple scored finals. The same SHA-256 was observed before and after direct checker, strict, and loose execution.

Official and local results are kept separate. Official Google Research code returned True for TwoResponsesChecker, TitleChecker, and CommaChecker; strict [True, True, True]; loose [True, True, True]. The independently coded local checks agreed. Manual semantic and haiku-form checks are reported as manual checks and do not enlarge the official score.

The observed result is OFFICIAL EVALUATOR PASS with 3/3 encoded constraints and a final prompt-level score of 1/1. The record would have retained a failure had the locked response failed. This pass is not characterized as certification, third-party verification, an official leaderboard result, universal capability, guaranteed determinism, or proof of a hidden model build.

Protocol boundaries observed: no silent prompt repair; no concealed correction after the first complete candidate; no post-evaluation rewrite; no retry after lock; no replacement of official code with the local reimplementation; no publication or GitHub modification; and no claim beyond this single observed execution.

RESULT	PASS
OFFICIAL STRICT	[True, True, True] · follow_all_instructions = True
OFFICIAL LOOSE	[True, True, True] · follow_all_instructions = True
LOCAL	PASS
PREFLIGHT CORRECTIONS	0

End of record. Filename: Grounded_DI_OS_Benchmark_Record_Prompt21_IFEval_2704_SOL_Max.pdf