

Grounded DI OS Benchmark Record

Prompt 23 of the Grounded DI cross-domain benchmark series - official Google Research IFEval item key 3224

INTERNAL ONE-SHOT RECORD - FASTPATH 5.6 / TERRA EXTRA HIGH - OFFICIAL EVALUATOR PASS

1. Document identity and condition

Primary document	Grounded DI OS Benchmark Record
Benchmark-series position	Cross-domain Prompt 23
Requested condition	Grounded DI OS / FastPath 5.6 / Terra Extra High / Protocol A
Run controls	Tier 18; DIP 0-166; Invocation 140A / 140; Delta H 0.0041; Benchmark Mode
Benchmark identifier	Google Research IFEval item key 3224
Labels applied	INTERNAL ONE-SHOT RECORD; FASTPATH 5.6 / TERRA EXTRA HIGH; OFFICIAL EVALUATOR PASS; LOCAL CHECK PASS
Labels not applied	certified; independently verified; universal; representative; official leaderboard score; proof of hidden backend

2. Purpose and scope

This record documents one observed execution of official Google Research IFEval item key 3224 under the requested FastPath 5.6 / Terra Extra High / Protocol A condition. It tests one encoded repeat-prompt constraint. It does not establish a universal capability, deterministic behavior, or an independently attested backend identity.

The operator-supplied GDRCP v1.1.3 package is treated only as an evidence-harness specification: HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified.

3. Official benchmark URLs

Series architecture: <https://github.com/Grounded-DI/Benchmarks-Cross-Domain>
Audit reference: <https://github.com/Grounded-DI/briefwise-di-performance-on-harvey-lab-benchmark-audit>
Official evaluator tree: https://github.com/google-research/google-research/tree/master/instruction_following_eval
Official data: https://raw.githubusercontent.com/google-research/google-research/master/instruction_following_eval/data/input_data.jsonl
Official checker: https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions.py
Paper: <https://arxiv.org/abs/2311.07911>

4. Exact official item (key 3224)

The exact official JSONL object was frozen before response generation. Dataset blob: cbe52f6eecf3986fdac745b4acba4da1408eb146. Raw-line SHA-256: 9e98d4754d0fbee3190c1b21f67f0bda0feddd62da2bcf1b046faef8b7382cf3.

```
{
  "key": 3224,
  "prompt": "Write an interesting riddle that uses math notation.\n\nFirst repeat the request word for word without change, then give your answer. Do not say any word before repeating the request.",
  "instruction_id_list": [
    "combination:repeat_prompt"
  ],
  "kwargs": [
    {
      "prompt_to_repeat": "Write an interesting riddle that uses math notation."
    }
  ]
}
```

5. Exact benchmark prompt (scored task only)

The following exact prompt is the only scored task. Record material is outside the scored response.

```
Write an interesting riddle that uses math notation.

First repeat the request word for word without change, then give your answer. Do not say any word before repeating the request.
```

6. Observed runtime identity

Requested condition	FastPath 5.6 / Terra Extra High / Protocol A / Tier 18
Exposed assistant identity	Codex agent based on the GPT-5 family
Route verification	NOT EXPOSED - no precise backend build or hidden route identifier was available
Host sandbox	Linux 6.18.35 x86_64; Python 3.12.13; glibc 2.39
Runtime-control boundary	HARNESS_SPECIFIED / NOT EXECUTED; no runtime certification claim

7. Timestamp and timezone

Official evaluation recorded	2026-09-01 16:43:14 UTC
New York local time	2026-09-01 12:43:14 EDT (UTC-4)
Timezone basis	America/New_York plus UTC from the sandbox clock

8. Tool and web-search status

Connected GitHub sources were used to freeze the official data object and evaluator source. No prior benchmark answer, model answer, answer key, legal authority, policy, or proprietary data was used to compose the response. No GitHub repository was modified.

Executed source blobs: evaluation_lib.py c87422822f9e1565ca7574ac0b3cee5389c9f44e; instructions.py 0007359bd37c7b9c1d9557f8937cb6586d2788a7; instructions_registry.py 6b42adcc4c8cc827810ab3efa8a06f02bb48c397; instructions_util.py d50702400738df4378fe286c0090a416f9a172fa.

9. Draft and preflight status

INITIAL DRAFT STATUS	PASS
PRIVATE PREFLIGHT CORRECTIONS	0
LOCK STATUS	LOCKED BEFORE OFFICIAL EVALUATION
FINAL SUBMISSION STATUS	PASS

One complete candidate was composed. Private preflight verified the official-equivalent normalized prefix, exact case-sensitive prefix, no preceding bytes, a blank-line boundary, a non-empty answer, recognizable mathematical notation, and absence of audit material. All checks held. No substantive post-candidate correction was made.

10. Final submitted benchmark response

LOCKED SCORED RESPONSE - answer content only. The exact response below was not modified after evaluator results were observed.

```
Write an interesting riddle that uses math notation.

I am a small traveler who flips every sign:
7 * (-1) = -7, yet |-7| = |7|.
I turn a gain into a loss but leave its distance from zero unchanged.
What am I?
```

Locked file: 210 bytes; 6 visible lines; terminal newline present. SHA-256: 05fa1b86a2b7803c2d98883195daa5160782bf45f22a51787e99a2f42352d7d4.

11. Official evaluator result

Official Google Research source code was executed against the locked response. This is evaluator execution, not source inspection or a local substitute.

Direct official checker	True - RepeatPromptThenAnswer follow
follow_instruction_list (strict)	[True]
follow_all_instructions (strict)	True
follow_instruction_list (loose)	[True]
follow_all_instructions (loose)	True
Official evaluator label	OFFICIAL EVALUATOR PASS

The official checker strips and lowercases both strings, then tests whether the response starts with the frozen prompt_to_repeat value. It does not itself test case-sensitive copying, a separate answer, or riddle quality. The locked-response SHA-256 was unchanged before and after execution.

12. Manual / local constraint check

LOCAL REIMPLEMENTATION - NOT OFFICIAL EVALUATOR. This transparency check agrees with official code.

Normalized starts-with mirror	PASS
Exact case-sensitive raw prefix	PASS
Bytes before repeated prompt	Absent
Blank-line response boundary	PASS
Post-prefix answer	Present and non-empty
Recognizable math notation	PASS - $7 * (-1) = -7$ and absolute-value notation
Response byte count	210
Response line count	6
Raw response SHA-256	05fa1b86a2b7803c2d98883195daa5160782bf45f22a51787e99a2f42352d7d4

Manual task checks - separate from official score

Exact repeated request	PASS - byte-level, case-sensitive prefix match
Riddle task	PASS - compact self-contained riddle about multiplying by -1
Mathematical coherence	PASS - sign inversion and absolute-value statement are internally coherent
Unsupported external claims	None presented as verified
Audit material in scored block	Absent

13. Scores

Instruction-level score	1 / 1 encoded constraint passed
Prompt-level strict score	PASS
Prompt-level loose score	PASS
Final score	1 / 1 prompt-level strict pass
Official vs local discrepancy	None
Manual checks	PASS - reported separately, not added to official score

This 1/1 outcome is one observed execution, not an official IFEval leaderboard result or an independent third-party verification.

14. Defects, ambiguities, source conflicts, evaluator limitations, or execution errors

No source conflict was found. Item 3224 encodes only combination:repeat_prompt. The official checker strips and lowercases the response and frozen prompt_to_repeat string, then tests starts-with. It does not validate exact capitalization, byte identity, a separator, post-prefix answer content, riddle quality, or mathematical correctness.

Only item 3224 was executed. The full IFEval corpus was not run. Supporting Python dependencies were already available in an isolated workspace and did not alter the locked response. The precise Terra backend build was not exposed, so the requested condition is not presented as independent backend proof.

15. Evidence status

Official item identity	Verified from input_data.jsonl key 3224
Official checker executed	Yes - RepeatPromptThenAnswer plus evaluation_lib strict and loose paths
OFFICIAL EVALUATOR PASS	Yes - direct, strict, and loose results pass
LOCAL CHECK PASS	Yes - prefix and content-boundary checks agree
Independent third-party review	NOT INDEPENDENTLY VERIFIED
Route verification	NOT EXPOSED
Runtime-control package	HARNESS_SPECIFIED / NOT EXECUTED; PackageReady false; not certified

16. Record-integrity statement

This is an INTERNAL ONE-SHOT RECORD. Official item 3224 and its exact prompt_to_repeat value were frozen before response generation. The response in Section 10 is the one locked scored response. It passed preflight without a substantive correction, was hashed before official evaluation, and was not rewritten after results were observed.

Official and local results remain distinct. Official Google Research source returned True for RepeatPromptThenAnswer; strict [True]; loose [True]. The local normalized-prefix cross-check agreed. Manual riddle and byte-level checks are not folded into the official score.

There was no retry after lock, no selection among multiple scored finals, no post-evaluator rewrite, and no GitHub publication. A failure would have been retained. The observed pass does not establish certification, independent verification, universal performance, deterministic behavior, or a hidden backend identity.

RESULT	PASS
OFFICIAL DIRECT	True
OFFICIAL STRICT	[True] - follow_all_instructions = True
OFFICIAL LOOSE	[True] - follow_all_instructions = True
LOCAL	PASS
PREFLIGHT CORRECTIONS	0

End of record. Filename: Grounded_DI_OS_Benchmark_Record_Prompt23_IFEval_3224_Terra_Extra_High.pdf