

Grounded DI OS Benchmark Record

Prompt 25 of the Grounded DI cross-domain benchmark series - official Google Research IFEval item key 2225

INTERNAL ONE-SHOT RECORD - FASTPATH 5.6 / SOL EXTRA HIGH - OFFICIAL EVALUATOR PASS

1. Document identity and condition

Primary document	Grounded DI OS Benchmark Record
Benchmark-series position	Cross-domain Prompt 25
Requested condition	Grounded DI OS / FastPath 5.6 / SOL Extra High / Protocol A
Run controls	Tier 18; DIP 0-166; Invocation 140A / 140; Delta H 0.0041; Benchmark Mode
Benchmark identifier	Google Research IFEval item key 2225
Labels applied	INTERNAL ONE-SHOT RECORD; FASTPATH 5.6 / SOL EXTRA HIGH; OFFICIAL EVALUATOR PASS; LOCAL CHECK PASS
Labels not applied	certified; independently verified; universal; representative; official leaderboard score; proof of hidden backend; linguistic certification

2. Purpose and scope

This record documents one observed execution of official Google Research IFEval item key 2225 under the requested FastPath 5.6 / SOL Extra High / Protocol A condition. It tests one encoded response-language control and separately checks whether the Korean answer preserves a stable civil-engineering distinction. It does not establish universal capability, deterministic behavior, independent linguistic certification, or an independently attested backend identity.

The operator-supplied GDRCP v1.1.3 package is treated only as an evidence-harness specification: HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified.

3. Official benchmark URLs

Series architecture: <https://github.com/Grounded-DI/Benchmarks-Cross-Domain>
Audit reference: <https://github.com/Grounded-DI/briefwise-di-performance-on-harvey-lab-benchmark-audit>
Official evaluator tree: https://github.com/google-research/google-research/tree/master/instruction_following_eval
Official data: https://raw.githubusercontent.com/google-research/google-research/master/instruction_following_eval/data/input_data.jsonl
Official checker: https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions.py
Paper: <https://arxiv.org/abs/2311.07911>

4. Exact official item (key 2225)

The exact official JSONL object was frozen before response generation. Dataset blob: cbe52f6eecf3986fdac745b4acba4da1408eb146. Raw-line SHA-256: f7fa86564414a8dbf94901b756e8430633736233589c7cbfea934a942dbf190f.

```
{
  "key": 2225,
  "prompt": "what is the difference between a levee and an embankment? Please respond to me only in Korean.",
  "instruction_id_list": [
    {
      "language": "response_language"
    }
  ],
  "kwargs": [
    {
      "language": "ko"
    }
  ]
}
```

5. Exact benchmark prompt (scored task only)

The following exact prompt is the only scored task. Record material is outside the scored response.

```
what is the difference between a levee and an embankment? Please respond to me only in Korean.
```

6. Observed runtime identity

Requested condition	FastPath 5.6 / SOL Extra High / Protocol A / Tier 18
Exposed assistant identity	Codex agent based on the GPT-5 family
Route verification	NOT EXPOSED - no precise backend build or hidden route identifier was available
Host sandbox	Linux 6.18.35 x86_64; Python 3.12.13; glibc 2.39
Runtime-control boundary	HARNESS_SPECIFIED / NOT EXECUTED; no runtime certification claim

7. Timestamp and timezone

Official evaluation recorded	2026-09-01 21:14:28 UTC
New York local time	2026-09-01 17:14:28 EDT (UTC-4)
Timezone basis	America/New_York plus UTC from the sandbox clock

8. Tool and web-search status

Official GitHub sources were used to freeze the data object and evaluator source. No prior benchmark answer, answer key, legal authority, policy, current fact, or proprietary data was used to compose the response. No GitHub repository was modified.

Frozen source identities: dataset blob cbe52f6eecf3986fdac745b4acba4da1408eb146; instructions.py blob 0007359bd37c7b9c1d9557f8937cb6586d2788a7. The official evaluator was executed from the frozen local source after response lock.

9. Draft and preflight status

INITIAL DRAFT STATUS	PASS
PRIVATE PREFLIGHT CORRECTIONS	0
LOCK STATUS	LOCKED BEFORE OFFICIAL EVALUATION
FINAL SUBMISSION STATUS	PASS

One complete candidate was composed. Private preflight verified Korean-only presentation, absence of Latin letters and audit material, a meaningful distinction between the two structures, and preservation of the broad-category relationship. All checks held. No substantive post-candidate correction was made.

10. Final submitted benchmark response

LOCKED SCORED RESPONSE - answer content only. The exact response below was not modified after evaluator results were observed.

제방은 강이나 바다의 물이 넘쳐 주변 지역이 침수되는 것을 막기 위해 물가를 따라 쌓는 구조물이다. 성토 구조물은 흙이나 돌 같은 재료를 쌓아 지면을 높이거나 도로와 철도 등을 지지하도록 만든 더 넓은 개념이다. 따라서 제방은 주로 홍수 방지와 물의 흐름 통제를 목적으로 하는 성토 구조물의 한 종류이다.

Locked file: 415 bytes; 1 visible line; terminal newline present. SHA-256: 5d64c20ddbe8b19f4238dff5e4a555d571bcd1108b24b64240a65908710f875.

11. Official evaluator result

Official Google Research source code was executed against the locked response. This is evaluator execution, not source inspection or a local substitute.

Direct official checker	True - ResponseLanguageChecker returned Korean match
follow_instruction_list (strict)	[True]
follow_all_instructions (strict)	True
follow_instruction_list (loose)	[True]
follow_all_instructions (loose)	True
Official evaluator label	OFFICIAL EVALUATOR PASS

ResponseLanguageChecker called langdetect.detect on the locked response and returned True because the detected code was ko. No language-detection exception occurred. The official score does not itself establish exclusive Korean usage, technical accuracy, or useful content.

12. Manual / local constraint check

LOCAL REIMPLEMENTATION - NOT OFFICIAL EVALUATOR. This transparency check agrees with the official result.

Local detected language	ko
Official-equivalent language match	PASS
Language-detection exception	None
Latin alphabet count	0
Hangul characters	122 of 125 non-whitespace characters
Hangul ratio	0.976; remaining 3 characters are ordinary punctuation
Other script characters	None
Response byte / line count	415 bytes / 1 visible line
Raw response SHA-256	5d64c20ddbe8b19f4238dff5e4a555d571bc1108b24b64240a65908710f875

Manual task checks - separate from official score

Korean-only presentation	PASS - no Latin letters or romanization
Levee function	PASS - water containment, flood prevention, and flow control
Embankment category	PASS - broader raised earth or fill structure
Relationship	PASS - levee identified as a purpose-specific embankment
Unsupported external claims	None presented as verified
Audit material in scored block	Absent

Plain-English back-translation for manual review only: A levee is built along water to prevent surrounding areas from flooding when river or sea water overflows. An embankment is the broader concept of a structure made by piling soil, stone, or similar material to raise ground or support roads and railways. Therefore, a levee is a type of embankment whose principal purposes are flood prevention and control of water flow. This back-translation is not scored and is not independent linguistic certification.

13. Scores

Instruction-level score	1 / 1 encoded language constraint passed
Prompt-level strict score	PASS
Prompt-level loose score	PASS
Final score	1 / 1 prompt-level strict pass
Official vs local discrepancy	None
Manual checks	PASS - reported separately, not added to official score

This 1/1 outcome is one observed execution, not an official IFEval leaderboard result, independent third-party review, or linguistic certification.

14. Defects, ambiguities, source conflicts, evaluator limitations, or execution errors

No source conflict or execution error was found. Item 2225 encodes only language:response_language. The official checker identifies a likely language by calling langdetect.detect; it does not prove that every token is Korean, validate the civil-engineering distinction, or measure usefulness.

The checker also returns True if langdetect raises LangDetectException. That exception path was not triggered here: direct, strict, loose, and local executions all detected ko. Script-purity and technical-content checks remain manual and local. The precise SOL backend build was not exposed.

15. Evidence status

Official item identity	Verified from input_data.jsonl key 2225
Official checker executed	Yes - ResponseLanguageChecker plus evaluation_lib strict and loose paths
OFFICIAL EVALUATOR PASS	Yes - direct, strict, and loose results pass
LOCAL CHECK PASS	Yes - language, script, and technical checks agree
Independent linguistic review	NOT INDEPENDENTLY VERIFIED
Route verification	NOT EXPOSED
Runtime-control package	HARNESS_SPECIFIED / NOT EXECUTED; PackageReady false; not certified

16. Record-integrity statement

This is an INTERNAL ONE-SHOT RECORD. Official item 2225 and its exact prompt were frozen before response generation. The response in Section 10 is the one locked scored response. It passed preflight without a substantive correction, was hashed before official evaluation, and was not rewritten after results were observed.

Official and local results remain distinct. Official Google Research source returned True for ResponseLanguageChecker; strict [True]; loose [True]. Local detection returned ko without exception. Manual script-purity and civil-engineering checks agreed but are not folded into the official score.

There was no retry after lock, no selection among multiple scored finals, no post-evaluator rewrite, and no GitHub publication. A failure would have been retained. The observed pass does not establish certification, independent verification, universal performance, deterministic behavior, independent linguistic certification, or a hidden backend identity.

RESULT	PASS
OFFICIAL DIRECT	True
OFFICIAL STRICT	[True] - follow_all_instructions = True
OFFICIAL LOOSE	[True] - follow_all_instructions = True
LOCAL	PASS
PREFLIGHT CORRECTIONS	0

End of record. Filename: Grounded_DI_OS_Benchmark_Record_Prompt25_IFEval_2225_SOL_Extra_High.pdf