

Grounded DI OS Benchmark Record

Prompt 28 of the Grounded DI cross-domain benchmark series - official Google Research IFEval item key 3386

INTERNAL ONE-SHOT RECORD - FASTPATH 5.6 / TERRA EXTRA HIGH - OFFICIAL EVALUATOR PASS

1. Document identity and condition

Primary document	Grounded DI OS Benchmark Record
Benchmark-series position	Cross-domain Prompt 28
Requested condition	Grounded DI OS / FastPath 5.6 / Terra Extra High / Protocol A
Run controls	Tier 18; DIP 0-166; Invocation 140A / 140; Delta H 0.0041; Benchmark Mode
Benchmark identifier	Google Research IFEval item key 3386
Encoded controls	keywords:forbidden_words; keywords:existence
Benchmark phase	Composite interaction test - not a new official checker family
Labels applied	INTERNAL ONE-SHOT RECORD; FASTPATH 5.6 / TERRA EXTRA HIGH; OFFICIAL EVALUATOR PASS; LOCAL CHECK PASS
Labels not applied	certified; independently verified; universal; representative; official leaderboard score; proof of hidden backend identity; retail-pricing certification

2. Purpose and scope

This record documents one observed composite execution of official Google Research IFEval item key 3386 under the requested FastPath 5.6 / Terra Extra High / Protocol A condition. The prompt combines a forbidden-word restriction with a required-keyword restriction. It also permits a transparent manual arithmetic check for a whole-gallon purchase: three gallons cost \$9 and leave \$1 from \$10.

Prompts 1-26 already exercised every registered IFEval checker family. Prompt 28 is therefore identified as a new composite item and interaction test rather than an unused checker-family test.

The operator-supplied GDRCP v1.1.3 package is treated only as an evidence-harness specification: HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified.

3. Official benchmark URLs

Series architecture: <https://github.com/Grounded-DI/Benchmarks-Cross-Domain>
Audit reference: <https://github.com/Grounded-DI/briefwise-di-performance-on-harvey-lab-benchmark-audit>
Official evaluator tree: https://github.com/google-research/google-research/tree/master/instruction_following_eval
Official data: https://raw.githubusercontent.com/google-research/google-research/master/instruction_following_eval/data/input_data.jsonl
Official checker: https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions.py
Paper: <https://arxiv.org/abs/2311.07911>

4. Exact official item (key 3386)

The exact official JSONL object was frozen before response generation. Dataset blob: cbe52f6eecf3986fdac745b4acba4da1408eb146. instructions.py blob: 0007359bd37c7b9c1d9557f8937cb6586d2788a7. Raw-line SHA-256: e2fa9a3d16de5d1fa277868a4b72e4a93c227329cfbe9ac45620fc6d7486b12c.

```
{
  "key": 3386,
  "prompt": "Jennifer goes to the store to buy milk. She has 10 dollars in her pocket
and milk costs 3 dollars per gallon. How many gallons of milk can she buy? Explain
your thinking. Avoid the keywords: 'divide', 'answer'. Include the keyword
'remainder'.",
  "instruction_id_list": [
    "keywords:forbidden_words",
    "keywords:existence"
  ],
  "kwargs": [
    {
      "forbidden_words": [
        "divide",
        "answer"
      ]
    },
    {
      "keywords": [
        "remainder"
      ]
    }
  ]
}
```

5. Exact benchmark prompt (scored task only)

The following exact prompt is the only scored task. Record material is outside the scored response.

Jennifer goes to the store to buy milk. She has 10 dollars in her pocket and milk costs 3 dollars per gallon. How many gallons of milk can she buy? Explain your thinking. Avoid the keywords: 'divide', 'answer'. Include the keyword 'remainder'.

6. Observed runtime identity

Requested condition	FastPath 5.6 / Terra Extra High / Protocol A / Tier 18
Exposed assistant identity	Codex agent based on the GPT-5 family
Route verification	ROUTE VERIFICATION NOT EXPOSED - no precise backend build or hidden route identifier was available
Host sandbox	Linux 6.18.35 x86_64; Python 3.12.13; glibc 2.39
Runtime-control boundary	HARNESS_SPECIFIED / NOT EXECUTED; no runtime certification claim

7. Timestamp and timezone

Official evaluation recorded	2026-09-01 23:50:48 UTC
Observed local system clock	2026-09-02 02:50:48 UTC+03:00
Timezone basis	UTC and the executing sandbox clock

8. Tool and web-search status

Official Google Research sources were used to freeze the data object and evaluator source. No answer key, prior benchmark answer, current fact research, legal authority, policy, proprietary data, or retail-pricing source was used to compose the scored response. No GitHub repository was modified.

The official source executed from the frozen local checkout with its established local Python dependencies. The response lock preceded all evaluation. No dependency repair was required after lock and the locked response bytes were not altered.

9. Draft and preflight status

INITIAL DRAFT STATUS	PASS
PRIVATE PREFLIGHT CORRECTIONS	0
LOCK STATUS	LOCKED BEFORE OFFICIAL EVALUATION
FINAL SUBMISSION STATUS	PASS

The supplied complete candidate was preserved as the single scored response. Private preflight confirmed absence of both prohibited strings, presence of remainder, whole-gallon arithmetic, and absence of audit material. No substantive post-candidate correction was made.

10. Final submitted benchmark response

LOCKED SCORED RESPONSE - answer content only. The exact response below was not modified after evaluator results were observed.

Four gallons would cost \$12 which is more than she has. Jennifer can buy 3 whole gallons because three gallons cost \$9. She has a remainder of \$1.

Locked file: 147 bytes; 1 visible line; terminal newline present. SHA-256: a9f50d69d063355f5c9dac2ab2965b3477577703b1755aa0ac54fd2c20ed34f7.

11. Official evaluator result

Official Google Research source code was executed against the locked response. This is evaluator execution, not source inspection or a local substitute.

Direct ForbiddenWords checker	True
Direct KeywordChecker	True
follow_instruction_list (strict)	[True, True]
follow_all_instructions (strict)	True
follow_instruction_list (loose)	[True, True]
follow_all_instructions (loose)	True
Official evaluator label	OFFICIAL EVALUATOR PASS

Official ForbiddenWords applies case-insensitive whole-word searches. Official KeywordChecker applies a case-insensitive substring search. Neither checker scores the arithmetic or purchase assumption.

12. Manual / local constraint check

LOCAL REIMPLEMENTATION - NOT OFFICIAL EVALUATOR. The final local check agrees with the official result.

Response byte / visible-line count	147 bytes / 1 line; terminal newline present
Raw response SHA-256	a9f50d69d063355f5c9dac2ab2965b3477577703b1755aa0ac54fd2c20ed34f7
Forbidden whole-word matches	divide: 0; answer: 0
Forbidden substring matches	divide: 0; answer: 0
Official-equivalent forbidden result	PASS
Official-equivalent keyword result	PASS - remainder found at characters 130-139
Standalone remainder result	PASS - one whole-word match at characters 130-139
Arithmetic: 10 / 3	3.333333...
Arithmetic: 3 x 3 / 10 - 9	\$9 / \$1
Four-gallon cost	\$12
Audit material in scored block	Absent

Manual task checks - separate from official score

Whole-gallon result	PASS - 3 whole gallons stated
Remainder amount	PASS - \$1 stated
Fractional-sales claim	Absent - no independent retail-practice claim
Interaction boundary	PASS - required keyword used without prohibited terms
Retail-pricing certification	NOT CLAIMED
Audit material in scored block	Absent

Local diagnostic note: the first independent standalone-word test contained an over-escaped regular-expression boundary and produced a false negative. That diagnostic was corrected after response lock and rerun locally. The locked response bytes and all official direct, strict, and loose results were unchanged. This was not a substantive response correction.

13. Scores

Instruction-level score	2 / 2 encoded constraints passed
Prompt-level strict score	PASS
Prompt-level loose score	PASS
Final score	1 / 1 prompt-level strict pass
Official vs local discrepancy	None after local diagnostic correction
Manual checks	PASS - reported separately, not added to official score

This 1/1 outcome is one observed composite execution, not an official IFEval leaderboard result, independent third-party review, retail-pricing certification, proof of universal capability, or a claim of a verified hidden backend.

14. Defects, ambiguities, source conflicts, evaluator limitations, or execution errors

No source conflict or official evaluator execution error was found. A local diagnostic implementation defect was found after official evaluation: escaped regular-expression boundaries prevented the first standalone-word check from matching the visibly present word remainder. The local diagnostic was corrected and rerun without changing the response or rerunning the official evaluator. The repair is disclosed in Section 12 and does not affect the recorded official results.

Item 3386 uses two existing official checker families. ForbiddenWords uses whole-word matching and therefore does not itself ban lexical variants such as answers or division. KeywordChecker uses substring matching and therefore does not itself require remainder to appear as a standalone word. Neither evaluator verifies arithmetic, whole-gallon purchasing, or actual retail sales practices. Route verification was not exposed.

15. Evidence status

Official item identity	Verified from input_data.jsonl key 3386
Official checkers executed	Yes - ForbiddenWords, KeywordChecker, evaluation_lib strict, and evaluation_lib loose
OFFICIAL EVALUATOR PASS	Yes - direct, strict, and loose results pass
LOCAL CHECK PASS	Yes - final local checker and manual arithmetic checks agree
Route verification	ROUTE VERIFICATION NOT EXPOSED
Runtime-control package	HARNESS_SPECIFIED / NOT EXECUTED; PackageReady false; not certified

16. Record-integrity statement

This is an INTERNAL ONE-SHOT RECORD. Official item 3386 and its exact prompt were frozen before response generation. The response in Section 10 is the one locked scored response. It passed preflight with zero substantive corrections, was hashed before official evaluation, and was not rewritten after results were observed.

Official and local results remain distinct. Official Google Research source returned True for both direct checkers; strict [True, True]; loose [True, True]. The corrected local diagnostic and the manual arithmetic checks agreed but are not folded into the official score.

There was no selection among multiple scored finals, no post-evaluator rewrite, and no GitHub publication. A failed locked response would have been retained. The observed pass does not establish certification, independent verification, universal performance, deterministic behavior, retail-pricing certification, or a hidden backend identity.

RESULT	PASS
OFFICIAL DIRECT	ForbiddenWords True; KeywordChecker True
OFFICIAL STRICT	[True, True] - follow_all_instructions = True
OFFICIAL LOOSE	[True, True] - follow_all_instructions = True
LOCAL	PASS
PREFLIGHT CORRECTIONS	0

End of record. Filename: Grounded_DI_OS_Benchmark_Record_Prompt28_IFEval_3386_Terra_Extra_High.pdf