

Grounded DI OS Benchmark Record

Prompt 29 of the Grounded DI cross-domain benchmark series - official Google Research IFEval item key 143

INTERNAL ONE-SHOT RECORD - FASTPATH 5.6 / TERRA EXTRA HIGH - OFFICIAL EVALUATOR PASS

1. Document identity and condition

Primary document	Grounded DI OS Benchmark Record
Benchmark-series position	Cross-domain Prompt 29
Requested condition	Grounded DI OS / FastPath 5.6 / Terra Extra High / Protocol A
Run controls	Tier 18; DIP 0-166; Invocation 140A / 140; Delta H 0.0041; Benchmark Mode
Benchmark identifier	Google Research IFEval item key 143
Encoded controls	length_constraints:number_paragraphs; detectable_content:postscript
Benchmark phase	Composite interaction test - not a new official checker family
Labels applied	INTERNAL ONE-SHOT RECORD; FASTPATH 5.6 / TERRA EXTRA HIGH; OFFICIAL EVALUATOR PASS; LOCAL CHECK PASS
Labels not applied	certified; independently verified; universal; representative; official leaderboard score; proof of hidden backend identity; veterinary certification

2. Purpose and scope

This record documents one observed composite execution of official Google Research IFEval item key 143 under the requested FastPath 5.6 / Terra Extra High / Protocol A condition. The prompt combines exact markdown-divider segmentation with a P.P.S postscript marker and contains a transparent arithmetic check.

Prompt 29 is a new composite item and interaction test, not a newly covered IFEval checker family. The required calculation is self-contained: 3 years is 36 months; 36 times 6 inches is 216 inches; 216 inches is 18 feet; and 5 plus 18 is 23 feet.

The operator-supplied GDRCP v1.1.3 package is treated only as an evidence-harness specification: HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified.

3. Official benchmark URLs

Series architecture: <https://github.com/Grounded-DI/Benchmarks-Cross-Domain>
Audit reference: <https://github.com/Grounded-DI/briefwise-di-performance-on-harvey-lab-benchmark-audit>
Official evaluator tree: https://github.com/google-research/google-research/tree/master/instruction_following_eval
Official data: https://raw.githubusercontent.com/google-research/google-research/master/instruction_following_eval/data/input_data.jsonl
Official checker: https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions.py
Paper: <https://arxiv.org/abs/2311.07911>

4. Exact official item (key 143)

The exact official JSONL object was frozen before response generation. Dataset blob: cbe52f6eecf3986fdac745b4acba4da1408eb146. instructions.py blob: 0007359bd37c7b9c1d9557f8937cb6586d2788a7. Raw-line SHA-256: 2c506938ce245222b863287f76787d1d34832d25a019046e7131f4ddc6c418e2.

```
{
  "key": 143,
  "prompt": "A colt is 5 feet tall. It will grow 6 inches every month. How tall will it be in 3 years? Think step-by-step, then give your answer. Separate your thinking and the final answer by a line with just three \"**\" symbols: ***\nAt the end of your response, please explicitly add a postscript starting with P.P.S",
  "instruction_id_list": [
    {
      "length_constraints": "number_paragraphs",
      "detectable_content": "postscript"
    }
  ],
  "kwargs": [
    {
      "num_paragraphs": 2
    },
    {
      "postscript_marker": "P.P.S"
    }
  ]
}
```

5. Exact benchmark prompt (scored task only)

The following exact prompt is the only scored task. Record material is outside the scored-response boundary.

```
A colt is 5 feet tall. It will grow 6 inches every month. How tall will it be in 3 years? Think step-by-step, then give your answer. Separate your thinking and the final answer by a line with just three "*" symbols: ***
At the end of your response, please explicitly add a postscript starting with P.P.S
```

6. Observed runtime identity

Requested condition	Grounded DI OS / FastPath 5.6 / Terra Extra High / Protocol A
Exposed runtime identity	Codex agent environment; no precise backend build identifier exposed
Route verification	ROUTE VERIFICATION NOT EXPOSED - requested condition is not independent proof of a hidden route
Host sandbox	Linux 6.18.35 x86_64; Python 3.12.13
Runtime-control boundary	HARNESS_SPECIFIED / NOT EXECUTED; no runtime-certification claim

7. Timestamp and timezone

Official evaluation recorded	2026-09-02 02:05 UTC
Timezone basis	UTC
Document generation	After evaluator result was known

8. Tool and web-search status

Official Google Research data and checker source were retrieved and matched to the listed freeze identities before response generation. No answer key, prior benchmark answer, current-fact research, legal authority, policy, proprietary data, or veterinary source was used to compose the scored response. No GitHub repository was modified.

The official evaluator required post-lock dependency repair for absl-py, langdetect, nltk, and immutabledict in a temporary evaluator dependency location. The locked response file was not changed.

9. Draft and preflight status

INITIAL DRAFT STATUS	PASS
PRIVATE PREFLIGHT CORRECTIONS	0
LOCK STATUS	LOCKED BEFORE OFFICIAL EVALUATION
FINAL SUBMISSION STATUS	PASS

The supplied complete candidate was preserved as the single scored response. Private preflight confirmed exactly two non-empty blocks, one separator line containing only *******, a terminal P.P.S postscript, correct unit conversion, and absence of audit material. No substantive post-candidate correction was made.

10. Final submitted benchmark response

LOCKED SCORED RESPONSE - answer content only. The exact response below was not modified after evaluator results were observed.

Three years is 36 months. At 6 inches per month the colt grows 216 inches. Since 216 inches equals 18 feet, adding the original 5 feet gives 23 feet. *** The colt will be 23 feet tall after 3 years. P.P.S. This calculation uses the constant monthly growth rate stated in the problem.

Locked file: 284 bytes; 4 visible lines; terminal newline present. SHA-256: 00b854b0859697e13c189709fad7dff74e07114286a0c3970a293ca75ebd4ac.

11. Official evaluator result

Official Google Research source code was executed against the locked response. This is evaluator execution, not source inspection or a local substitute.

Direct ParagraphChecker	True
Direct PostscriptChecker	True
follow_instruction_list (strict)	[True, True]
follow_all_instructions (strict)	True
follow_instruction_list (loose)	[True, True]
follow_all_instructions (loose)	True
Official evaluator label	OFFICIAL EVALUATOR PASS

Official ParagraphChecker splits on evaluator-recognized ******* dividers and counts non-empty segments. Official PostscriptChecker searches for a recognized P.P.S form. Neither scores terminal placement or arithmetic.

12. Manual / local constraint check

LOCAL REIMPLEMENTATION - NOT OFFICIAL EVALUATOR. The final local check agrees with the official result.

Response byte / visible-line count	284 bytes / 4 lines; terminal newline present
Raw response SHA-256	00b854b0859697e13c189709fad7dfc74e07114286a0c3970a293ca75ebd4ac
*** count / byte offset	1 / 150
Separator presentation	PASS - one line containing only *** (line 2)
Non-empty blocks after exact line split	2
Official-equivalent ParagraphChecker	PASS
Official-equivalent PostscriptChecker	PASS
P.P.S placement	PASS - begins final non-whitespace line; no content follows
Arithmetic: 3 x 12 months	36 months
Arithmetic: 36 x 6 inches	216 inches
Arithmetic: 216 / 12	18 feet
Final height	5 + 18 = 23 feet
Audit material in scored block	Absent

Manual task checks - separate from official score

Two-block ordered response	PASS - reasoning then final answer
Exact separator-line boundary	PASS
Terminal P.P.S presentation	PASS
Unit conversion	PASS - 23 feet under the prompt's stated rate
Growth realism	Not claimed - constant monthly growth is treated only as the stated word-problem premise
Audit material in scored block	Absent

13. Scores

Instruction-level score	2 / 2 encoded constraints passed
Prompt-level strict score	PASS
Prompt-level loose score	PASS
Final score	1 / 1 prompt-level strict pass
Official vs local discrepancy	None
Manual checks	PASS - reported separately, not added to official score

This 1/1 outcome is one observed composite execution, not an official IFEval leaderboard result, independent third-party review, proof of universal capability, or proof of a verified hidden backend.

14. Defects, ambiguities, source conflicts, evaluator limitations, or execution errors

No source conflict or official evaluator execution error was found. Before official execution the environment lacked evaluator imports. After lock, absl-py, langdetect, nltk, and immutabledict were installed solely to enable the official source. The locked response hash remained 00b854b0859697e13c189709fad7dff74e07114286a0c3970a293ca75ebd4ac.

ParagraphChecker counts ***-separated non-empty segments and does not itself require *** to stand alone on a line. PostscriptChecker detects a recognized P.P.S form, but does not itself prove the marker is the final non-whitespace content. Neither checker assesses the arithmetic, unit conversion, ordered reasoning, substantive quality, or real-world growth assumptions. Route verification was not exposed.

15. Evidence status

Official item identity	Verified from input_data.jsonl key 143
Official checkers executed	Yes - ParagraphChecker, PostscriptChecker, evaluation_lib strict, and evaluation_lib loose
OFFICIAL EVALUATOR PASS	Yes - direct, strict, and loose results pass
LOCAL CHECK PASS	Yes - final local structure and arithmetic checks agree
Route verification	ROUTE VERIFICATION NOT EXPOSED
Runtime-control package	HARNESS_SPECIFIED / NOT EXECUTED; PackageReady false; not certified

16. Record-integrity statement

This is an INTERNAL ONE-SHOT RECORD. Official item 143 and its exact prompt were frozen before response generation. The response in Section 10 is the one locked scored response. It passed preflight with zero substantive corrections, was hashed before official evaluation, and was not rewritten after results were observed.

Official and local results remain distinct. Official Google Research source returned True for both direct checkers; strict [True, True]; loose [True, True]. Local structural and arithmetic checks agreed but are not folded into the official score.

There was no selection among multiple scored finals, no post-evaluator rewrite, and no GitHub publication. A failed locked response would have been retained. The observed pass does not establish certification, independent verification, universal performance, deterministic behavior, veterinary validity, or a hidden backend identity.

RESULT	PASS
OFFICIAL DIRECT	ParagraphChecker True; PostscriptChecker True
OFFICIAL STRICT	[True, True] - follow_all_instructions = True
OFFICIAL LOOSE	[True, True] - follow_all_instructions = True
LOCAL	PASS
PREFLIGHT CORRECTIONS	0

End of record. Filename: Grounded_DI_OS_Benchmark_Record_Prompt29_IFEval_143_Terra_Extra_High.pdf