

# Grounded DI OS Benchmark Record

Prompt 31 of the Grounded DI cross-domain benchmark series - official Google Research IFEval item key 1375

INTERNAL ONE-SHOT RECORD - FASTPATH 5.6 / TERRA EXTRA HIGH - OFFICIAL EVALUATOR PASS

## 1. Document identity and condition

|                           |                                                                                                                                                                                   |
|---------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Primary document          | Grounded DI OS Benchmark Record                                                                                                                                                   |
| Benchmark-series position | Cross-domain Prompt 31                                                                                                                                                            |
| Requested condition       | Grounded DI OS / FastPath 5.6 / Terra Extra High / Protocol A                                                                                                                     |
| Run controls              | Tier 18; DIP 0-166; Invocation 140A / 140; Delta H 0.0041; Benchmark Mode                                                                                                         |
| Benchmark identifier      | Google Research IFEval item key 1375                                                                                                                                              |
| Encoded controls          | length_constraints:number_paragraphs; keywords:existence                                                                                                                          |
| Benchmark phase           | Composite interaction test - not a new official checker family                                                                                                                    |
| Labels applied            | INTERNAL ONE-SHOT RECORD; FASTPATH 5.6 / TERRA EXTRA HIGH; OFFICIAL EVALUATOR PASS; LOCAL CHECK PASS                                                                              |
| Labels not applied        | certified; independently verified; universal; representative; official leaderboard score; proof of hidden backend identity; security certification; production-security assurance |

## 2. Purpose and scope

This record documents one observed composite execution of official Google Research IFEval item key 1375 under the requested FastPath 5.6 / Terra Extra High / Protocol A condition. The item combines exact two-block segmentation with a required keyword and is evaluated only against one locked response.

The task is a limited defensive Java illustration: an unsafe dynamic SQL construction pattern followed by parameter binding. The record neither supplies an exploit payload nor claims the snippets are complete, compilable, production-ready, or a complete security design.

The operator-supplied GDRCP v1.1.3 package is treated only as an evidence-harness specification: HARNESS\_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified.

## 3. Official benchmark URLs

Series architecture: <https://github.com/Grounded-DI/Benchmarks-Cross-Domain>  
Audit reference: <https://github.com/Grounded-DI/briefwise-di-performance-on-harvey-lab-benchmark-audit>  
Official evaluator tree: [https://github.com/google-research/google-research/tree/master/instruction\\_following\\_eval](https://github.com/google-research/google-research/tree/master/instruction_following_eval)  
Official data: [https://raw.githubusercontent.com/google-research/google-research/master/instruction\\_following\\_eval/data/input\\_data.jsonl](https://raw.githubusercontent.com/google-research/google-research/master/instruction_following_eval/data/input_data.jsonl)  
Official checker: [https://github.com/google-research/google-research/blob/master/instruction\\_following\\_eval/instructions.py](https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions.py)  
Paper: <https://arxiv.org/abs/2311.07911>

## 4. Exact official item (key 1375)

The exact official JSONL line was frozen before response generation. Dataset blob: cbe52f6eecf3986fdac745b4acba4da1408eb146. instructions.py blob: 0007359bd37c7b9c1d9557f8937cb6586d2788a7. Raw-line SHA-256: 0a0a00e9e45b77ce1d2a5d48ca6307fb4fae50e1cbe38f47f795f0472bd1d22b.

```
{
  "key": 1375,
  "prompt": "Improper use of the Java API can lead to vulnerabilities in web applications. Write a code snippet to demonstrate this. Then demonstrate how to fix the code. Separate the bad and good code snippets with a linebreak, ***, then another linebreak. Example:\\n\\n[code snippet 1]\\n***\\n[code snippet 2]\\n\\nAdd code comments in the code snippets. One of the comments should contain the keyword \\\"vulnerable\\\".",
  "instruction_id_list": [
    {
      "length_constraints": {
        "number_paragraphs": 2
      },
      "keywords": {
        "keywords": [
          "vulnerable"
        ]
      }
    }
  ]
}
```

### 5. Exact benchmark prompt (scored task only)

The following exact prompt is the only scored task. Record material is outside the scored-response boundary.

```
Improper use of the Java API can lead to vulnerabilities in web applications. Write a code snippet to demonstrate this. Then demonstrate how to fix the code. Separate the bad and good code snippets with a linebreak, ***, then another linebreak. Example:

[code snippet 1]
***
[code snippet 2]

Add code comments in the code snippets. One of the comments should contain the keyword "vulnerable".
```

### 6. Observed runtime identity

|                          |                                                                                                 |
|--------------------------|-------------------------------------------------------------------------------------------------|
| Requested condition      | Grounded DI OS / FastPath 5.6 / Terra Extra High / Protocol A                                   |
| Exposed runtime identity | Codex agent environment; no precise backend build identifier exposed                            |
| Route verification       | ROUTE VERIFICATION NOT EXPOSED - requested condition is not independent proof of a hidden route |
| Runtime-control boundary | HARNESS_SPECIFIED / NOT EXECUTED; no runtime-certification claim                                |

### 7. Timestamp and timezone

|                              |                                  |
|------------------------------|----------------------------------|
| Official source freeze       | 2026-09-02 14:59:07 UTC          |
| Official evaluation recorded | 2026-09-02 15:00 UTC             |
| Document generation          | After evaluator result was known |
| Timezone basis               | UTC                              |

### 8. Tool and web-search status

Official Google Research data and checker source were retrieved and matched to the listed freeze identities before response generation. No answer key, prior benchmark answer, current-fact research, proprietary code, credential, external service, legal authority, or application target was used. No GitHub repository was modified.

Official evaluator code executed with existing temporary evaluation dependencies. No dependency repair was required after lock. The locked response file was not changed.

### 9. Draft and preflight status

|                               |                                   |
|-------------------------------|-----------------------------------|
| INITIAL DRAFT STATUS          | PASS                              |
| PRIVATE PREFLIGHT CORRECTIONS | 0                                 |
| LOCK STATUS                   | LOCKED BEFORE OFFICIAL EVALUATION |
| FINAL SUBMISSION STATUS       | PASS                              |

The supplied complete candidate was preserved as the single scored response. Private preflight confirmed one separator line containing only `***`, two non-empty code blocks, a case-insensitive vulnerable match inside a code comment, parameter binding in the second block, and no audit or disallowed material. No substantive post-candidate correction was made.

### 10. Final submitted benchmark response

LOCKED SCORED RESPONSE - answer content only. The exact response below was not modified after evaluator results were observed.

```
String username = request.getParameter("username");
String sql = "SELECT id FROM users WHERE name = '" + username + "'";
try (Statement stmt = connection.createStatement()) {
    ResultSet rs = stmt.executeQuery(sql); // vulnerable: untrusted input is concatenated
}
***
String username = request.getParameter("username");
String sql = "SELECT id FROM users WHERE name = ?";
try (PreparedStatement stmt = connection.prepareStatement(sql)) {
    stmt.setString(1, username); // bind untrusted input as data
    ResultSet rs = stmt.executeQuery();
}
```

Locked file: 542 bytes; 12 visible lines; terminal newline present. SHA-256: 9909c1e1773759782cc5401300cf982995c5616732cb42a8d33ace0e9fcfb58.

### 11. Official evaluator result

Official Google Research source code was executed against the locked response. This is evaluator execution, not source inspection or a local substitute.

|                                  |                         |
|----------------------------------|-------------------------|
| Direct ParagraphChecker          | True                    |
| Direct KeywordChecker            | True                    |
| follow_instruction_list (strict) | [True, True]            |
| follow_all_instructions (strict) | True                    |
| follow_instruction_list (loose)  | [True, True]            |
| follow_all_instructions (loose)  | True                    |
| Official evaluator label         | OFFICIAL EVALUATOR PASS |

Official ParagraphChecker only splits on the evaluator-recognized `***` pattern. Official KeywordChecker uses a case-insensitive substring search. Neither determines whether code is safe, compilable, Java, correctly ordered, or free of exploit material.

## 12. Manual / local constraint check

**LOCAL REIMPLEMENTATION - NOT OFFICIAL EVALUATOR. The final local check agrees with the official result.**

|                                             |                                                                 |
|---------------------------------------------|-----------------------------------------------------------------|
| <b>Response byte / visible-line count</b>   | 542 bytes / 12 lines; terminal newline present                  |
| <b>Raw response SHA-256</b>                 | 9909c1e1773759782cc5401300cf982995c5616732cb42a8d33ace0e9fcfb58 |
| <b>*** count / byte offset</b>              | 1 / 265                                                         |
| <b>Separator presentation</b>               | PASS - one line containing only ***                             |
| <b>Non-empty blocks after split</b>         | 2                                                               |
| <b>Official-equivalent ParagraphChecker</b> | PASS                                                            |
| <b>vulnerable matches</b>                   | 1; code-comment match PASS                                      |
| <b>Official-equivalent KeywordChecker</b>   | PASS                                                            |
| <b>Unsafe first-block pattern</b>           | PASS - dynamic query construction and createStatement present   |
| <b>Defensive second-block pattern</b>       | PASS - placeholder, prepareStatement, and setString present     |
| <b>Disallowed attack or target material</b> | Absent                                                          |
| <b>Audit material in scored block</b>       | Absent                                                          |

**Manual task checks - separate from official score**

|                                                 |                                                                                                             |
|-------------------------------------------------|-------------------------------------------------------------------------------------------------------------|
| <b>Two-block ordered code response</b>          | PASS - unsafe pattern first; parameterized correction second                                                |
| <b>Required comment keyword</b>                 | PASS - vulnerable appears in the first code comment                                                         |
| <b>Defensive binding</b>                        | PASS - user-controlled value bound with setString                                                           |
| <b>Payload, credential, target, or endpoint</b> | Absent                                                                                                      |
| <b>Scope boundary</b>                           | No claim of compilation, production readiness, complete authorization, or comprehensive security validation |
| <b>Audit material in scored block</b>           | Absent                                                                                                      |

## 13. Scores

|                                      |                                                         |
|--------------------------------------|---------------------------------------------------------|
| <b>Instruction-level score</b>       | 2 / 2 encoded constraints passed                        |
| <b>Prompt-level strict score</b>     | PASS                                                    |
| <b>Prompt-level loose score</b>      | PASS                                                    |
| <b>Final score</b>                   | 1 / 1 prompt-level strict pass                          |
| <b>Official vs local discrepancy</b> | None                                                    |
| <b>Manual checks</b>                 | Reported separately; not added to official IFEval score |

## 14. Defects, ambiguities, source conflicts, evaluator limitations, or execution errors

No source conflict or evaluator execution error was observed. The official ParagraphChecker splits on the \*\*\* pattern and the official KeywordChecker searches case-insensitively for a substring. Neither evaluator validates code behavior, compilation, safety, the ordered bad/fixed distinction, or absence of exploit material. The Java snippets are intentionally narrow illustrations and not an application-security assessment.

ROUTE VERIFICATION NOT EXPOSED remains a material limitation. The requested condition label is recorded as supplied and is not proof of a hidden backend build or runtime engine.

15. Evidence status

|                             |                                                                                                                                   |
|-----------------------------|-----------------------------------------------------------------------------------------------------------------------------------|
| Official item evidence      | FROZEN - exact raw JSONL line preserved; raw-line SHA-256 recorded                                                                |
| Official checker evidence   | FROZEN - instructions.py blob identity recorded                                                                                   |
| Scored response evidence    | LOCKED - SHA-256 recorded before official evaluation                                                                              |
| Official evaluator evidence | EXECUTED - direct, strict, and loose results recorded                                                                             |
| Local-check evidence        | EXECUTED - transparent reimplementation and manual checks recorded                                                                |
| Runtime-control package     | HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified                     |
| PDF visual QA               | PASS - final rendered pages inspected; no clipping, overlap, malformed code display, missing glyphs, or footer collision observed |

16. Record-integrity statement

This is an internal one-shot benchmark record. It binds the named official source freeze, the single locked response, direct official-checker results, strict and loose official evaluator results, and separately labeled local and manual checks. The PDF and package checksum sidecars provide file-integrity evidence for the delivered artifacts.

This record does not claim certification, independent verification, an official leaderboard score, universal or representative behavior, deterministic guarantees, security certification, production-security assurance, or proof of a particular hidden backend identity. The requested FastPath condition is recorded as an operator-supplied label subject to the stated route-verification limitation.

Source freeze summary

```
PROMPT 31 OFFICIAL SOURCE FREEZE
Frozen at: 2026-09-02T14:59:07Z
Repository: google-research/google-research
Dataset path: instruction_following_eval/data/input_data.jsonl
Dataset blob: cbe52f6eecf3986fdac745b4acba4da1408eb146
Checker path: instruction_following_eval/instructions.py
Checker blob: 0007359bd37c7b9c1d9557f8937cb6586d2788a7
Item key: 1375
Encoded instruction IDs, ordered: length_constraints:number_paragraphs; keywords:existence
Raw official item-line SHA-256:
0a0a00e9e45b77ce1d2a5d48ca6307fb4fae50e1cbe38f47f795f0472bd1d22b
Exact official JSONL line is preserved in item_1375_raw.jsonl.
```