

Grounded DI OS Benchmark Record

Prompt 32 of the Grounded DI cross-domain benchmark series - official Google Research IFEval item key 1936

INTERNAL ONE-SHOT RECORD - FASTPATH 5.6 / TERRA EXTRA HIGH - OFFICIAL EVALUATOR PASS

1. Document identity and condition

Primary document	Grounded DI OS Benchmark Record
Benchmark-series position	Cross-domain Prompt 32
Requested condition	Grounded DI OS / FastPath 5.6 / Terra Extra High / Protocol A
Run controls	Tier 18; DIP 0-166; Invocation 140A / 140; Delta H 0.0041; Benchmark Mode
Benchmark identifier	Google Research IFEval item key 1936
Encoded controls	keywords:forbidden_words; detectable_format:number_bullet_lists
Benchmark phase	Composite interaction test - not a new official checker family
Labels applied	INTERNAL ONE-SHOT RECORD; FASTPATH 5.6 / TERRA EXTRA HIGH; OFFICIAL EVALUATOR PASS; LOCAL CHECK PASS
Labels not applied	certified; independently verified; universal; representative; official leaderboard score; proof of hidden backend identity

2. Purpose and scope

This record documents one observed composite execution of official Google Research IFEval item key 1936 under the requested FastPath 5.6 / Terra Extra High / Protocol A condition. The item combines an exact bullet-count constraint with a forbidden-word constraint and is evaluated only against one locked response.

The scored task is a concise customer-service performance rubric. Manual review considers whether the six bullets are understandable and evaluable, but those content-quality considerations are distinct from the official IFEval score.

The operator-supplied GDRCP v1.1.3 package is treated only as an evidence-harness specification: HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified.

3. Official benchmark URLs

Series architecture: <https://github.com/Grounded-DI/Benchmarks-Cross-Domain>
Audit reference: <https://github.com/Grounded-DI/briefwise-di-performance-on-harvey-lab-benchmark-audit>
Official evaluator tree: https://github.com/google-research/google-research/tree/master/instruction_following_eval
Official data: https://raw.githubusercontent.com/google-research/google-research/master/instruction_following_eval/data/input_data.jsonl
Official checker: https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions.py
Paper: <https://arxiv.org/abs/2311.07911>

4. Exact official item (key 1936)

The exact official JSONL line was frozen before response generation. Dataset blob: cbe52f6eef3986fdac745b4acba4da1408eb146. instructions.py blob: 0007359bd37c7b9c1d9557f8937cb6586d2788a7. Raw-line SHA-256: e5c85222806feeac47dbc5645d7de4651278905d4d799941772074f8e9b46e23.

```
{
  "key": 1936,
  "prompt": "Write a rubric, in the form of a list of bullet points, for evaluating the performance of a customer service representative. Your answer must not include keywords ['bad', 'underperform'] and must contain exactly 6 bullet points in the following form:\n* Bullet point 1\n* Bullet point 2\n* Bullet point 3\n* Bullet point 4\n* Bullet point 5\n* Bullet point 6",
  "instruction_id_list": [
    "keywords:forbidden_words",
    "detectable_format:number_bullet_lists"
  ],
  "kwargs": {
    "forbidden_words": ["bad", "underperform"],
    "num_bullets": 6
  }
}
```

5. Exact benchmark prompt (scored task only)

The following exact prompt is the only scored task. Record material is outside the scored-response boundary.

```
Write a rubric, in the form of a list of bullet points, for evaluating the performance of a customer service representative. Your answer must not include keywords ['bad', 'underperform'] and must contain exactly 6 bullet points in the following form:
* Bullet point 1
* Bullet point 2
* Bullet point 3
* Bullet point 4
* Bullet point 5
* Bullet point 6
```

6. Observed runtime identity

Requested condition	Grounded DI OS / FastPath 5.6 / Terra Extra High / Protocol A
Exposed runtime identity	Codex agent environment; no precise backend build identifier exposed
Route verification	ROUTE VERIFICATION NOT EXPOSED - requested condition is not independent proof of a hidden route
Runtime-control boundary	HARNESS_SPECIFIED / NOT EXECUTED; no runtime-certification claim

7. Timestamp and timezone

Official source freeze	2026-09-02 15:19:39 UTC
Official evaluation recorded	2026-09-02 UTC, after response lock
Document generation	After evaluator result was known
Timezone basis	UTC

8. Tool and web-search status

Official Google Research data and checker source were retrieved and matched to the listed freeze identities before response generation. No answer key, prior benchmark answer, current-fact research, proprietary material, credential, external service, legal authority, or repository modification was used.

Official evaluator code executed with existing temporary evaluation dependencies. The locked response file was not changed.

9. Draft and preflight status

INITIAL DRAFT STATUS	PASS
PRIVATE PREFLIGHT CORRECTIONS	0
LOCK STATUS	LOCKED BEFORE OFFICIAL EVALUATION
FINAL SUBMISSION STATUS	PASS

The supplied complete candidate was preserved as the single scored response. Private preflight confirmed six conventional asterisk bullet lines, no hyphen bullets, no forbidden-word match, no extra prose, and no audit material. No substantive post-candidate correction was made.

10. Final submitted benchmark response

LOCKED SCORED RESPONSE - answer content only. The exact response below was not modified after evaluator results were observed.

- * Demonstrates active listening and accurately identifies the customer's issue.
- * Uses clear respectful language and maintains a professional tone.
- * Provides timely next steps that match the customer's stated need.
- * Documents the interaction accurately and completely.
- * Escalates matters through the appropriate channel when needed.
- * Follows through to confirm that the issue has been resolved.

Locked file: 400 bytes; 7 visible lines; terminal newline present. SHA-256: 551aea32e4b49d4bf0e1fb18fe1f79270a3e6db9f7f276391d90b3e74ad50a62.

11. Official evaluator result

Official Google Research source code was executed against the locked response. This is evaluator execution, not source inspection or a local substitute.

Direct ForbiddenWords	True
Direct BulletListChecker	True
follow_instruction_list (strict)	[True, True]
follow_all_instructions (strict)	True
follow_instruction_list (loose)	[True, True]
follow_all_instructions (loose)	True
Official evaluator label	OFFICIAL EVALUATOR PASS

Official ForbiddenWords uses a case-insensitive whole-word search. BulletListChecker counts recognized asterisk and hyphen bullet patterns. Neither evaluator establishes substantive rubric quality or absence of extra prose.

12. Manual / local constraint check

LOCAL REIMPLEMENTATION - NOT OFFICIAL EVALUATOR. The final local check agrees with the official result.

Response byte / visible-line count	400 bytes / 7 lines; terminal newline present
Raw response SHA-256	551aea32e4b49d4bf0e1fb18fe1f79270a3e6db9f7f276391d90b3e74ad50a62
Raw official item SHA-256	e5c85222806feeac47dbc5645d7de4651278905d4d799941772074f8e9b46e23
Forbidden whole-word matches	0 for each forbidden word
Forbidden substring matches	0 for each forbidden string
Official-equivalent ForbiddenWords	PASS
Asterisk / hyphen bullet counts	6 / 0
Official-equivalent total bullet count	6
Official-equivalent BulletListChecker	PASS
Conventional asterisk bullets	6
Non-bullet content	Absent
Audit material in scored block	Absent

Manual task checks - separate from official score

Six meaningful evaluation criteria	PASS
Listening, professionalism, and clarity	PASS
Timeliness, documentation, and escalation	PASS
Follow-through	PASS
Unsupported external claim	Absent
Audit material in scored block	Absent

13. Scores

Instruction-level score	2 / 2 encoded constraints passed
Prompt-level strict score	PASS
Prompt-level loose score	PASS
Final score	1 / 1 prompt-level strict pass
Official vs local discrepancy	None after documented local-mirror repair
Manual checks	Reported separately; not added to official IFEval score

14. Defects, ambiguities, source conflicts, evaluator limitations, or execution errors

No source conflict or official evaluator error was observed. The first local mirror contained an escaped-regex defect that incorrectly counted zero bullet lines. That local mirror was corrected after lock; the locked response bytes were unchanged, official direct and strict/loose evaluator outputs were already preserved as True, and the corrected local mirror returned the expected pass. This was a local-check implementation repair, not a response correction or a second candidate.

ROUTE VERIFICATION NOT EXPOSED remains a material limitation. The requested condition label is recorded as supplied and is not proof of a hidden backend build or runtime engine.

15. Evidence status

Official item evidence	FROZEN - exact raw JSONL line preserved; raw-line SHA-256 recorded
Official checker evidence	FROZEN - instructions.py blob identity recorded
Scored response evidence	LOCKED - SHA-256 recorded before official evaluation
Official evaluator evidence	EXECUTED - direct, strict, and loose results recorded
Local-check evidence	EXECUTED - corrected transparent mirror and manual checks recorded
Runtime-control package	HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified
PDF visual QA	PASS - final rendered pages inspected; no clipping, overlap, missing glyphs, or footer collision observed

16. Record-integrity statement

This is an internal one-shot benchmark record. It binds the named official source freeze, the single locked response, direct official-checker results, strict and loose official evaluator results, and separately labeled local and manual checks. The PDF and package checksum sidecars provide file-integrity evidence for the delivered artifacts.

This record does not claim certification, independent verification, an official leaderboard score, universal or representative behavior, deterministic guarantees, or proof of a particular hidden backend identity. The requested FastPath condition is recorded as an operator-supplied label subject to the stated route-verification limitation.

Source freeze summary

```
PROMPT 32 OFFICIAL SOURCE FREEZE
Frozen at: 2026-09-02T15:19:39Z
Repository: google-research/google-research
Dataset path: instruction_following_eval/data/input_data.jsonl
Dataset blob: cbe52f6eecf3986fdac745b4acba4da1408eb146
Checker path: instruction_following_eval/instructions.py
Checker blob: 0007359bd37c7b9c1d9557f8937cb6586d2788a7
Item key: 1936
Encoded instruction IDs, in order: keywords:forbidden_words; detectable_format:
number_bullet_lists
Exact official JSONL line preserved in item_1936_raw.jsonl.
```