

Grounded DI OS Benchmark Record

Prompt 33 of the Grounded DI cross-domain benchmark series - official Google Research IFEval item key 1262

INTERNAL ONE-SHOT RECORD - FASTPATH 5.6 / LUNA MAX - OFFICIAL EVALUATOR PASS

1. Document identity and condition

Primary document	Grounded DI OS Benchmark Record
Benchmark-series position	Cross-domain Prompt 33
Requested condition	Grounded DI OS / FastPath 5.6 / Luna Max / Protocol A
Run controls	Tier 18; DIP 0-166; Invocation 140A / 140; Delta H 0.0041; Benchmark Mode
Benchmark identifier	Google Research IFEval item key 1262
Encoded controls	detectable_format:title; length_constraints:number_sentences
Benchmark phase	Composite interaction test - not a new official checker family
Labels applied	INTERNAL ONE-SHOT RECORD; FASTPATH 5.6 / LUNA MAX; OFFICIAL EVALUATOR PASS; LOCAL CHECK PASS
Labels not applied	certified; independently verified; universal; representative; official leaderboard score; proof of hidden backend identity

2. Purpose and scope

This record documents one observed composite execution of official Google Research IFEval item key 1262 under the requested FastPath 5.6 / Luna Max / Protocol A condition. The item combines a double-angular-bracket title control with a strict sentence-budget control and is evaluated against one locked response.

The scored task is a concise explanation of how increased axle weight affects road damage. Manual review considers whether the explanation is technically coherent and useful, but those content-quality considerations are distinct from the official IFEval score.

The operator-supplied GDRCP v1.1.3 package is treated only as an evidence-harness specification: HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified.

3. Official benchmark URLs

Series architecture: <https://github.com/Grounded-DI/Benchmarks-Cross-Domain>

Audit reference: <https://github.com/Grounded-DI/briefwise-di-performance-on-harvey-lab-benchmark-audit>

Official evaluator tree: https://github.com/google-research/google-research/tree/master/instruction_following_eval

Official data:
https://raw.githubusercontent.com/google-research/google-research/master/instruction_following_eval/data/input_data.jsonl

Official checker: https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions.py

Official utility: https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions_util.py

Paper: <https://arxiv.org/abs/2311.07911>

4. Exact official item (key 1262)

The exact official JSONL line was fetched from the default branch and matched before response generation. Dataset blob: cbe52f6eefc3986fdac745b4acba4da1408eb146. instructions.py blob: 0007359bd37c7b9c1d9557f8937cb6586d2788a7. instructions_util.py blob: d50702400738df4378fe286c0090a416f9a172fa. The preserved line includes its terminal LF; its SHA-256 is e02fe8e418c57ced929be2d7ae1d76472e494fd0262ee5f19566936096c3164a. The same UTF-8 line without the terminal LF hashes to 5b6e7a9557247e5a4beab39a3dc4da96fcd8d0247de1baf2e5b6905131bdfa3.

```
{
  "key": 1262,
  "prompt": "Can you elaborate on the following text: \"Increased axle weight increased the damage to the road\"? Your response must contain a title wrapped in double angular brackets, i.e. <<title>>. Use less than 5 sentences.",
  "instruction_id_list": [
    "detectable_format:title",
    "length_constraints:number_sentences"
  ],
  "kwargs": {
    "relation": "less than",
    "num_sentences": 5
  }
}
```

5. Exact benchmark prompt (scored task only)

The following exact prompt is the only scored task. Record material is outside the scored-response boundary.

Can you elaborate on the following text: "Increased axle weight increased the damage to the road"? Your response must contain a title wrapped in double angular brackets, i.e. <<title>>. Use less than 5 sentences.

6. Observed runtime identity

Requested condition	Grounded DI OS / FastPath 5.6 / Luna Max / Protocol A
Exposed runtime identity	Codex agent environment; no precise backend or model build identifier exposed
Route verification	ROUTE VERIFICATION NOT EXPOSED - requested condition is not independent proof of a hidden route
Runtime-control boundary	HARNESS_SPECIFIED / NOT EXECUTED; no runtime-certification claim

7. Timestamp and timezone

Official source freeze	2026-09-02 17:18:13 UTC
Official evaluation recorded	2026-09-02 UTC, after response lock
Document generation	After evaluator result was known
Timezone basis	UTC

8. Tool and web-search status

Official Google Research data, instructions.py, instructions_util.py, instructions_registry.py, and evaluation_lib.py were retrieved from the official repository and matched to the recorded source identities before response generation. No answer key, current-fact research, proprietary material, credential, external service, legal authority, or repository modification was used.

The exact official checker code executed after lock with temporary evaluation dependencies and Punkt data. Dependency repair did not alter the locked response. No general web search was used.

9. Draft and preflight status

Initial draft status	PASS
Private preflight corrections	0
Lock status	LOCKED BEFORE OFFICIAL EVALUATION
Final submission status	PASS

The complete candidate was preserved as the single scored response. Private preflight confirmed one non-empty double-angular-bracket title, fewer than five sentences, a substantive explanation, and no audit material. No substantive post-candidate correction was made.

10. Final submitted benchmark response

LOCKED SCORED RESPONSE - answer content only. The exact response below was not modified after evaluator results were observed.

<<Axle Weight and Road Damage>>
Heavier axle loads place greater stress on pavement and its supporting layers.
Repeated loading accelerates fatigue cracking and permanent deformation so damage increases as axle weight rises.
This relationship explains why axle limits help protect road life.

Locked file: 292 bytes; 4 visible lines; terminal newline present. SHA-256:
 bbd1baa7d04b91fd3fc058e13c24b805764e62c17534d7b2b98df713ca191109.

11. Official evaluator result

Official Google Research source code was executed against the locked response. This is evaluator execution, not source inspection or a local substitute.

Direct TitleChecker	True
Direct NumberOfSentences	True
follow_instruction_list (strict)	[True, True]
follow_all_instructions (strict)	True
follow_instruction_list (loose)	[True, True]
follow_all_instructions (loose)	True
Official evaluator label	OFFICIAL EVALUATOR PASS

TitleChecker accepts a non-empty single-line double-angular-bracket span. NumberOfSentences uses the official Punkt sentence tokenizer and applies the strict less-than threshold. Neither evaluator establishes engineering accuracy or title quality.

12. Manual / local constraint check

LOCAL REIMPLEMENTATION - NOT OFFICIAL EVALUATOR. The independent local checks agree with the official result.

Response byte / visible-line count	292 bytes / 4 lines; terminal newline present
Raw response SHA-256	bbd1baa7d04b91fd3fc058e13c24b805764e62c17534d7b2b98df713ca191109
Raw official item SHA-256	e02fe8e418c57ced929be2d7ae1d76472e494fd0262ee5f19566936096c3164a (line including terminal LF)
Title regex matches	1; non-empty interior: PASS
Title at first non-whitespace line	PASS
Official-equivalent sentence count	3; threshold < 5: PASS
Independent punctuation sentence count	3; threshold < 5: PASS
Non-empty explanation	PASS
Engineering terms and causal explanation	PASS
Audit material in scored block	Absent

13. Scores

Instruction-level score	2 / 2 encoded constraints passed
Prompt-level strict score	PASS
Prompt-level loose score	PASS
Final score	1 / 1
Official versus local discrepancy	None
Manual checks	Reported separately; not added to official IFEval score

14. Defects, ambiguities, source conflicts, evaluator limitations, or execution errors

No source conflict or official evaluator execution error was observed. The TitleChecker recognizes any non-empty content on one line between double angular brackets and does not assess title quality. The NumberOfSentences checker depends on the official Punkt tokenizer and only measures sentence segmentation against the threshold; it does not assess the road-engineering explanation, causal precision, or usefulness.

The loose evaluator path is an upper-bound variant that tests transformations of the response; it is not an independent quality assessment. Route verification remains unavailable: the requested Luna Max label is recorded as supplied and is not proof of a hidden backend identity.

15. Evidence status

Official item evidence	FROZEN - exact official JSONL line preserved; source blob and line hashes recorded
Official checker evidence	FROZEN - instructions.py, instructions_util.py, registry, and evaluator identities recorded
Scored response evidence	LOCKED - SHA-256 recorded before evaluator execution
Official evaluator evidence	EXECUTED - direct, strict, and loose results recorded
Local-check evidence	EXECUTED - independent sentence, title, boundary, and content checks recorded
Runtime-control package	HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified
PDF visual QA	PASS - rendered pages inspected; no clipping, overlap, missing glyphs, or footer collision observed

16. Record-integrity statement

This is an internal one-shot benchmark record. It binds the named official source freeze, one locked response, direct official-checker results, strict and loose evaluator results, and separately labeled local and manual checks. The PDF and package checksum sidecars provide file-integrity evidence for the delivered artifacts.

This record does not claim certification, independent verification, an official leaderboard score, universal or representative behavior, deterministic guarantees, engineering certification, or proof of a particular hidden backend identity. The requested Luna Max condition is an operator-supplied label subject to the route-verification limitation stated above.

Source freeze summary

```
PROMPT 33 OFFICIAL SOURCE FREEZE
Frozen item: key 1262
Dataset blob: cbe52f6eecf3986fdac745b4acba4da1408eb146
Checker: instructions.py
Checker blob: 0007359bd37c7b9c1d9557f8937cb6586d2788a7
Utility: instructions_util.py
Utility blob: d50702400738df4378fe286c0090a416f9a172fa
Encoded instruction IDs in order: detectable_format:title; length_constraints:number_sentences
Raw JSONL line SHA-256 including LF: e02fe8e418c57ced929be2d7ae1d76472e494fd0262ee5f19566936096c3164a
Locked response SHA-256: bbd1baa7d04b91fd3fc058e13c24b805764e62c17534d7b2b98df713ca191109
Exact official JSONL line preserved in item_1262_raw.jsonl.
```