

Grounded DI OS Benchmark Record

Prompt 36 of the Grounded DI cross-domain benchmark series - official Google Research IFEval item key 1705

INTERNAL ONE-SHOT RECORD - FASTPATH 5.6 / LUNA MAX - OFFICIAL EVALUATOR PASS

1. Document identity and condition

Primary document	Grounded DI OS Benchmark Record
Benchmark-series position	Cross-domain Prompt 36
Requested condition	Grounded DI OS / FastPath 5.6 / Luna Max / Protocol A
Run controls	Tier 18; DIP 0-166; Invocation 140A / 140; Delta H 0.0041; Benchmark Mode
Benchmark identifier	Google Research IFEval item key 1705
Encoded controls	detectable_content:postscript; startend:end_checker
Benchmark phase	Composite interaction test - previously covered checker families
Labels applied	INTERNAL ONE-SHOT RECORD; FASTPATH 5.6 / LUNA MAX; OFFICIAL EVALUATOR PASS; LOCAL CHECK PASS
Labels not applied	certified; independently verified; universal; representative; official leaderboard score; proof of hidden backend identity

2. Purpose and scope

This record documents one observed composite execution of official Google Research IFEval item key 1705 under the requested FastPath 5.6 / Luna Max / Protocol A condition. The item combines a postscript-marker condition with a terminal exact-phrase condition and is evaluated against one locked response.

The scored task asks for a professional email requesting a raise. Manual review considers whether the email is professional, clearly requests compensation review, contains a distinct postscript, and ends with the required phrase. Those content and presentation checks remain distinct from the official IFEval score.

The operator-supplied GDRCP v1.1.3 package is treated only as an evidence-harness specification: HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified.

3. Official benchmark URLs

Series architecture: <https://github.com/Grounded-DI/Benchmarks-Cross-Domain>

Audit reference: <https://github.com/Grounded-DI/briefwise-di-performance-on-harvey-lab-benchmark-audit>

Official evaluator tree: https://github.com/google-research/google-research/tree/master/instruction_following_eval

Official data: https://raw.githubusercontent.com/google-research/google-research/master/instruction_following_eval/data/input_data.jsonl

Official checker: https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions.py

Paper: <https://arxiv.org/abs/2311.07911>

4. Exact official item (key 1705)

The exact official JSONL line was frozen before response generation. Dataset blob: cbe52f6eecf3986fdac745b4acba4da1408eb146. instructions.py blob: 0007359bd37c7b9c1d9557f8937cb6586d2788a7. Raw-line SHA-256: 787a3deca4963e87b0088f3f63edbab45676d43c98eca22745c855f04e320da3.

```
{
  "key": 1705,
  "prompt": "Write a professional email that you could send to ask your boss for a raise. At the end of your response, explicitly add a postscript starting with P.P.S\n\nFinish your entire response with this exact phrase: Hope you agree with me.",
  "instruction_id_list": [
    "detectable_content:postscript",
    "startend:end_checker"
  ],
  "kwargs": [
    [{"postscript_marker": "P.P.S"}, {"end_phrase": "Hope you agree with me."}]
  ]
}
```

5. Exact benchmark prompt (scored task only)

The following exact prompt is the only scored task. Record material is outside the scored-response boundary.

Write a professional email that you could send to ask your boss for a raise. At the end of your response, explicitly add a postscript starting with P.P.S

Finish your entire response with this exact phrase: Hope you agree with me.

6. Observed runtime identity

Requested condition	Grounded DI OS / FastPath 5.6 / Luna Max / Protocol A
Exposed runtime identity	Codex agent environment; precise backend build or model snapshot not exposed
Route verification	ROUTE VERIFICATION NOT EXPOSED - requested condition is not independent proof of a hidden route
Runtime-control boundary	HARNESS_SPECIFIED / NOT EXECUTED; no runtime-certification claim

7. Timestamp and timezone

Official source freeze	2026-09-02 UTC before response generation
Initial candidate	2026-09-02 UTC before private preflight
Response lock	2026-09-02 UTC after preflight and before official evaluation
Official evaluation recorded	2026-09-02 UTC after response lock
Document generation	After evaluator result was known
Timezone basis	UTC

8. Tool and web-search status

Official Google Research data and checker source were retrieved and matched to the listed freeze identities before response generation. No web search was used for the scored answer. No prior benchmark answer, current-fact research, proprietary material, credential, external service, legal authority, or repository modification was used.

Official evaluator code executed against the locked response using the preserved Google Research source. The locked response file was not changed after evaluation began.

9. Draft and preflight status

INITIAL DRAFT STATUS	PASS
PRIVATE PREFLIGHT CORRECTIONS	1
Correction	Inserted missing commas and corrected the phrase 'supported team priorities' to a grammatical coordinated list
LOCK STATUS	LOCKED BEFORE OFFICIAL EVALUATION
FINAL SUBMISSION STATUS	PASS

One complete candidate was generated. During private preflight, one substantive grammar correction was made before the response was locked. No retry, selection among multiple finals, or post-lock rewrite occurred.

10. Final submitted benchmark response

LOCKED SCORED RESPONSE - answer content only. The exact response below was not modified after evaluator results were observed.

```
Subject: Compensation review

Dear [Boss],

I would appreciate an opportunity to discuss my compensation. Over the past year I have consistently delivered high-quality work, supported team priorities, and taken on additional responsibilities. I would welcome a conversation about whether my contributions and current scope support an appropriate salary adjustment. I am committed to continuing to grow with the team and to identifying goals for the coming year.

Could we schedule a time to discuss this? Thank you for considering my request.

Best,
[Your Name]

P.P.S. Hope you agree with me.
```

Locked file: 594 bytes; 12 visible lines; terminal newline present. SHA-256: b865b1e1cfd88e5d6c7ed06a73b82f7e91452abad1ee047b0f563be007486f3e.

11. Official evaluator result

Official Google Research source code was executed against the locked response. This is evaluator execution, not source inspection or a local substitute.

Direct PostscriptChecker	True
Direct EndChecker	True
follow_instruction_list (strict)	[True, True]
follow_all_instructions (strict)	True
follow_instruction_list (loose)	[True, True]
follow_all_instructions (loose)	True
Official evaluator label	OFFICIAL EVALUATOR PASS

PostscriptChecker uses a case-insensitive line pattern for the configured P.P.S marker. EndChecker strips surrounding whitespace and double quotes before checking the configured ending phrase. Neither evaluator establishes professional quality or substantive compensation grounds.

12. Manual / local constraint check

LOCAL REIMPLEMENTATION - NOT OFFICIAL EVALUATOR. The transparent local check agrees with the official result.

Response byte / visible-line count	594 bytes / 12 lines; terminal newline present
Raw response SHA-256	b865b1e1cfd88e5d6c7ed06a73b82f7e91452abad1ee047b0f563be007486f3e
Raw official item SHA-256	787a3deca4963e87b0088f3f63edbab45676d43c98eca22745c855f04e320da3
Leading non-whitespace bytes	0
Trailing non-whitespace bytes	1
Official-equivalent P.P.S result	PASS
P.P.S match count	1
P.P.S line starts with marker	PASS
Official-equivalent ending result	PASS
Exact final phrase	PASS
Final non-whitespace line	P.P.S. Hope you agree with me.
Professional email check	PASS
Raise-request check	PASS
Audit material in scored block	Absent

Manual task checks - separate from official score

Professional email structure	PASS
Clear request to discuss compensation	PASS
Distinct P.P.S line	PASS
P.P.S is terminal non-whitespace line	PASS
Required exact phrase is final	PASS
No invented authority or current fact	PASS
No audit material in scored block	PASS

13. Scores

Instruction-level score	2 / 2 encoded constraints passed
Prompt-level strict score	PASS
Prompt-level loose score	PASS
Final score	1 / 1 prompt-level strict pass
Official vs local discrepancy	None
Manual checks	Reported separately; not added to official IFEval score

14. Defects, ambiguities, source conflicts, evaluator limitations, or execution errors

No source conflict or official evaluator execution error was observed. The initial draft had one grammar defect in a coordinated list; it was corrected during private preflight before lock and is counted above. No content was changed after lock.

The official PostscriptChecker detects a configured P.P.S-like marker anywhere in the response and does not prove that the postscript is terminal. The official EndChecker checks only a normalized suffix and does not prove that the response is an email, that the request is professional, or that the exact phrase is isolated on its own line.

The local manual checks apply a stricter presentation boundary: one visible P.P.S line at the end, exact terminal phrase, professional email structure, and no audit material. These checks are not additional official IFEval criteria and are not professional or employment-law certification.

ROUTE VERIFICATION NOT EXPOSED remains a material limitation. The requested Luna Max condition is recorded as supplied and is not proof of a hidden backend build or route.

15. Evidence status

Official item evidence	FROZEN - exact raw JSONL line preserved; raw-line SHA-256 recorded
Official checker evidence	FROZEN - instructions.py blob identity recorded
Scored response evidence	LOCKED - SHA-256 recorded before official evaluation
Official evaluator evidence	EXECUTED - direct, strict, and loose results recorded
Local-check evidence	EXECUTED - transparent mirror and manual checks recorded
Runtime-control package	HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified
PDF visual QA	PASS - final rendered pages inspected; no clipping, overlap, missing glyphs, or footer collision observed

16. Record-integrity statement

This is an INTERNAL ONE-SHOT RECORD for the requested FASTPATH 5.6 / LUNA MAX condition. The frozen item, initial candidate, one disclosed preflight correction, locked response, evaluator outputs, local checks, and PDF were recorded in order. Locked response SHA-256: b865b1e1cfd88e5d6c7ed06a73b82f7e91452abad1ee047b0f563be007486f3e. Official item-line SHA-256: 787a3deca4963e87b0088f3f63edbab45676d43c98eca22745c855f04e320da3. The record preserves the observed result and does not claim certification, independent verification, leaderboard standing, guaranteed determinism, or proof of hidden backend identity.