

Grounded DI OS Benchmark Record

Prompt 37 of the Grounded DI cross-domain benchmark series - official Google Research IFEval item key 1730

INTERNAL ONE-SHOT RECORD - FASTPATH 5.6 / LUNA MAX - OFFICIAL EVALUATOR PASS

1. Document identity and condition

Primary document	Grounded DI OS Benchmark Record
Benchmark-series position	Cross-domain Prompt 37
Requested condition	Grounded DI OS / FastPath 5.6 / Luna Max / Protocol A
Run controls	Tier 18; DIP 0-166; Invocation 140A / 140; Delta H 0.0041; Benchmark Mode
Benchmark identifier	Google Research IFEval item key 1730
Encoded controls	startend:quotation; detectable_format:multiple_sections
Benchmark phase	Composite interaction test - previously covered checker families
Labels applied	INTERNAL ONE-SHOT RECORD; FASTPATH 5.6 / LUNA MAX; OFFICIAL EVALUATOR PASS; LOCAL CHECK PASS
Labels not applied	certified; independently verified; universal; representative; official leaderboard score; proof of hidden backend identity

2. Purpose and scope

This record documents one observed composite execution of official Google Research IFEval item key 1730 under the requested FastPath 5.6 / Luna Max / Protocol A condition. The item combines whole-response double quotation marks with a five-section structure and is evaluated against one locked response.

The scored task asks for a blog post about raising awareness for a cause. Manual review considers whether the five sections form a useful, respectful awareness plan and whether the response remains entirely inside the outer quotation marks. Those content and presentation checks remain distinct from the official IFEval score.

The operator-supplied GDRCP v1.1.3 package is treated only as an evidence-harness specification: HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified.

3. Official benchmark URLs

Series architecture: <https://github.com/Grounded-DI/Benchmarks-Cross-Domain>

Audit reference: <https://github.com/Grounded-DI/briefwise-di-performance-on-harvey-lab-benchmark-audit>

Official evaluator tree: https://github.com/google-research/google-research/tree/master/instruction_following_eval

Official data: https://raw.githubusercontent.com/google-research/google-research/master/instruction_following_eval/data/input_data.jsonl

Official checker: https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions.py

Paper: <https://arxiv.org/abs/2311.07911>

4. Exact official item (key 1730)

The exact official JSONL line was frozen before response generation. Dataset blob: cbe52f6eecf3986fdac745b4acba4da1408eb146. instructions.py blob: 0007359bd37c7b9c1d9557f8937cb6586d2788a7. Raw-line SHA-256: 354ecceb360f22e34bef684822915524d44fa464aa05171f60f91f160545631c.

```
{
  "key": 1730,
  "prompt": "Write a blog post about how to raise awareness for a cause. Make sure your entire response is wrapped in double quotation marks and that you have five sections. Mark the beginning of each section with Section X.",
  "instruction_id_list": [
    "startend:quotation",
    "detectable_format:multiple_sections"
  ],
  "kwargs": {
    "section_splitter": "Section",
    "num_sections": 5
  }
}
```

5. Exact benchmark prompt (scored task only)

The following exact prompt is the only scored task. Record material is outside the scored-response boundary.

Write a blog post about how to raise awareness for a cause. Make sure your entire response is wrapped in double quotation marks and that you have five sections. Mark the beginning of each section with Section X.

6. Observed runtime identity

Requested condition	Grounded DI OS / FastPath 5.6 / Luna Max / Protocol A
Exposed runtime identity	Codex agent environment; precise backend build or model snapshot not exposed
Route verification	ROUTE VERIFICATION NOT EXPOSED - requested condition is not independent proof of a hidden route
Runtime-control boundary	HARNESS_SPECIFIED / NOT EXECUTED; no runtime-certification claim

7. Timestamp and timezone

Official source freeze	2026-09-02 UTC before response generation
Initial candidate	2026-09-02 UTC before private preflight
Response lock	2026-09-02 UTC after preflight and before official evaluation
Official evaluation recorded	2026-09-02 UTC after response lock
Document generation	After evaluator result was known
Timezone basis	UTC

8. Tool and web-search status

Official Google Research data and checker source were retrieved and matched to the listed freeze identities before response generation. No web search was used for the scored answer. No current-fact research, proprietary material, credential, external service, or repository modification was used.

Official evaluator code executed against the locked response using the preserved Google Research source. The locked response file was not changed after evaluation began.

9. Draft and preflight status

INITIAL DRAFT STATUS	PASS
PRIVATE PREFLIGHT CORRECTIONS	0
Correction	None. The complete candidate already satisfied the requested outer-quote and five-section presentation.
LOCK STATUS	LOCKED BEFORE OFFICIAL EVALUATION
FINAL SUBMISSION STATUS	PASS

One complete candidate was generated. Private preflight found no substantive defect, so no correction was made. No retry, selection among multiple finals, or post-lock rewrite occurred.

10. Final submitted benchmark response

LOCKED SCORED RESPONSE - answer content only. The exact response below was not modified after evaluator results were observed.

```
"Section 1
A cause becomes easier to support when people understand its human stakes. Begin with a clear
description of the problem, the community members affected, and the change you want to see.

Section 2
Tell a short respectful story that makes the issue real without exposing private details. Pair it
with plain facts from trustworthy sources and explain why the information matters locally.

Section 3
Invite people to participate through schools, neighborhood groups, libraries, and online
communities. Offer specific roles such as sharing a resource, hosting a conversation, or
volunteering an hour.

Section 4
Make action simple and concrete. Give readers one immediate step, a way to learn more, and a
clear explanation of how their contribution supports the cause.

Section 5
Keep awareness growing by listening to questions, reporting progress, and thanking contributors.
Review what worked, correct mistakes openly, and adapt the campaign so attention becomes
sustained community action."
```

Locked file: 1002 bytes; 14 visible lines; terminal newline present. SHA-256: 6fe8cf338d4beb883c9e426bd2dc4c483a4ff6630250d6eec4cb0ee2b7b1161b.

11. Official evaluator result

Official Google Research source code was executed against the locked response. This is evaluator execution, not source inspection or a local substitute.

Direct QuotationChecker	True
Direct SectionChecker	True
follow_instruction_list (strict)	[True, True]
follow_all_instructions (strict)	True
follow_instruction_list (loose)	[True, True]
follow_all_instructions (loose)	True
Official evaluator label	OFFICIAL EVALUATOR PASS

QuotationChecker strips whitespace and checks only the first and last characters. SectionChecker splits on the configured Section-plus-number pattern and checks that at least five sections are recognized. Neither evaluator establishes blog quality, factual accuracy, or absence of audit material.

12. Manual / local constraint check

LOCAL REIMPLEMENTATION - NOT OFFICIAL EVALUATOR. The transparent local check agrees with the official result.

Response byte / visible-line count	1002 bytes / 14 lines; terminal newline present
Raw response SHA-256	6fe8cf338d4beb883c9e426bd2dc4c483a4ff6630250d6eec4cb0ee2b7b1161b
Raw official item SHA-256	354ecceb360f22e34bef684822915524d44fa464aa05171f60f91f160545631c
Leading non-whitespace bytes	0
Trailing non-whitespace bytes	1
Official-equivalent quotation	PASS
Outer double-quote boundary	PASS
Internal double-quote count	0
Official-equivalent section count	5
Section marker match count	5
Conventional section-marker lines	5
Non-empty section bodies	5
No material outside outer quotes	PASS
Audit material in scored block	Absent

Manual task checks - separate from official score

Five meaningful section blocks	PASS
Awareness strategy is substantive	PASS
Respectful privacy-aware communication	PASS
Concrete participation and next steps	PASS
No invented authority or current fact	PASS
No audit material in scored block	PASS

13. Scores

Instruction-level score	2 / 2 encoded constraints passed
Prompt-level strict score	PASS
Prompt-level loose score	PASS
Final score	1 / 1 prompt-level strict pass
Official vs local discrepancy	None
Manual checks	Reported separately; not added to official IFEval score

14. Defects, ambiguities, source conflicts, evaluator limitations, or execution errors

No source conflict or official evaluator execution error was observed. The initial candidate passed private preflight with zero substantive corrections. No content was changed after lock.

The official QuotationChecker checks only an outer pair of double-quote characters after trimming whitespace; it does not require a particular quotation-mark style inside the response or validate a blog post. The official SectionChecker recognizes occurrences matching the configured Section-plus-number pattern and only requires at least five recognized sections; it does not prove that the sections are distinct, ordered, meaningful, or free of extra prose.

The local manual checks apply a stricter presentation boundary: exactly five non-empty section bodies, five conventional Section-number marker lines, no internal double quotes, useful cause-awareness content, and no audit material. These checks are not additional official IFEval criteria.

ROUTE VERIFICATION NOT EXPOSED remains a material limitation. The requested Luna Max condition is recorded as supplied and is not proof of a hidden backend build or route.

15. Evidence status

Official item evidence	FROZEN - exact raw JSONL line preserved; raw-line SHA-256 recorded
Official checker evidence	FROZEN - instructions.py blob identity recorded
Scored response evidence	LOCKED - SHA-256 recorded before official evaluation
Official evaluator evidence	EXECUTED - direct, strict, and loose results recorded
Local-check evidence	EXECUTED - transparent mirror and manual checks recorded
Runtime-control package	HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified
PDF visual QA	PASS - final rendered pages inspected; no clipping, overlap, missing glyphs, or footer collision observed

16. Record-integrity statement

This is an INTERNAL ONE-SHOT RECORD for the requested FASTPATH 5.6 / LUNA MAX condition. The frozen item, initial candidate, zero disclosed preflight corrections, locked response, evaluator outputs, local checks, and PDF were recorded in order. Locked response SHA-256: 6fe8cf338d4beeb883c9e426bd2dc4c483a4ff6630250d6eec4cb0ee2b7b1161b. Official item-line SHA-256: 354ecceb360f22e34bef684822915524d44fa464aa05171f60f91f160545631c. The record preserves the observed result and does not claim certification, independent verification, leaderboard standing, guaranteed determinism, or proof of hidden backend identity.