

Grounded DI OS Benchmark Record

Prompt 38 of the Grounded DI cross-domain benchmark series - official Google Research IFEval item key 3518

INTERNAL ONE-SHOT RECORD - FASTPATH 5.6 / LUNA MAX - OFFICIAL EVALUATOR PASS

1. Document identity and condition

Primary document	Grounded DI OS Benchmark Record
Benchmark-series position	Cross-domain Prompt 38
Requested condition	Grounded DI OS / FastPath 5.6 / Luna Max / Protocol A
Run controls	Tier 18; DIP 0-166; Invocation 140A / 140; Delta H 0.0041; Benchmark Mode
Benchmark identifier	Google Research IFEval item key 3518
Encoded controls	detectable_format:json_format; keywords:existence
Benchmark phase	Composite interaction test - previously covered checker families
Labels applied	INTERNAL ONE-SHOT RECORD; FASTPATH 5.6 / LUNA MAX; OFFICIAL EVALUATOR PASS; LOCAL CHECK PASS
Labels not applied	certified; independently verified; universal; representative; official leaderboard score; proof of hidden backend identity

2. Purpose and scope

This record documents one observed composite execution of official Google Research IFEval item key 3518 under the requested FastPath 5.6 / Luna Max / Protocol A condition. The item combines JSON serialization with required-keyword presence and is evaluated against one locked response.

The scored task asks for pros and cons of working abroad. Manual review considers whether the JSON is well formed, contains meaningful pros and cons, uses the required terms, and avoids unsupported current or organizational claims. Those content checks remain distinct from the official IFEval score.

The operator-supplied GDRCP v1.1.3 package is treated only as an evidence-harness specification: HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified.

3. Official benchmark URLs

Series architecture: <https://github.com/Grounded-DI/Benchmarks-Cross-Domain>

Audit reference: <https://github.com/Grounded-DI/briefwise-di-performance-on-harvey-lab-benchmark-audit>

Official evaluator tree: https://github.com/google-research/google-research/tree/master/instruction_following_eval

Official data: https://raw.githubusercontent.com/google-research/google-research/master/instruction_following_eval/data/input_data.jsonl

Official checker: https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions.py

Paper: <https://arxiv.org/abs/2311.07911>

4. Exact official item (key 3518)

The exact official JSONL line was frozen before response generation. Dataset blob: cbe52f6eecf3986fdac745b4acba4da1408eb146. instructions.py blob: 0007359bd37c7b9c1d9557f8937cb6586d2788a7. Raw-line SHA-256: fb152545716be53a8cbcbd0c56c98ea60a40bc76526ed324a1c725763acf6eca.

```
{
  "key": 3518,
  "prompt": "Could you please give me the pros and cons of working abroad wrapped in JSON format. Please make sure that your response only contains a JSON block. Please also make sure to include keywords 'compensated' and 'immigrants' in the response.",
  "instruction_id_list": [
    "detectable_format:json_format",
    "keywords:existence"
  ],
  "kwargs": {
    {},
    {
      "keywords": [
        "compensated",
        "immigrants"
      ]
    }
  ]
}
```

5. Exact benchmark prompt (scored task only)

The following exact prompt is the only scored task. Record material is outside the scored-response boundary.

```
Could you please give me the pros and cons of working abroad wrapped in JSON format. Please make sure that your response only contains a JSON block. Please also make sure to include keywords 'compensated' and 'immigrants' in the response.
```

6. Observed runtime identity

Requested condition	Grounded DI OS / FastPath 5.6 / Luna Max / Protocol A
Exposed runtime identity	Codex agent environment; precise backend build or model snapshot not exposed
Route verification	ROUTE VERIFICATION NOT EXPOSED - requested condition is not independent proof of a hidden route
Runtime-control boundary	HARNESS_SPECIFIED / NOT EXECUTED; no runtime-certification claim

7. Timestamp and timezone

Official source freeze	2026-09-02 UTC before response generation
Initial candidate	2026-09-02 UTC before private preflight
Response lock	2026-09-02 UTC after preflight and before official evaluation
Official evaluation recorded	2026-09-02 UTC after response lock
Document generation	After evaluator result was known
Timezone basis	UTC

8. Tool and web-search status

Official Google Research data and checker source were retrieved and matched to the listed freeze identities before response generation. No web search was used for the scored answer. No current-fact research, proprietary material, credential, external service, or repository modification was used.

Official evaluator code executed against the locked response using the preserved Google Research source. The locked response file was not changed after evaluation began.

9. Draft and preflight status

INITIAL DRAFT STATUS	PASS
PRIVATE PREFLIGHT CORRECTIONS	0
Correction	None. The complete candidate already satisfied JSON-only presentation and both required keywords.
LOCK STATUS	LOCKED BEFORE OFFICIAL EVALUATION
FINAL SUBMISSION STATUS	PASS

One complete candidate was generated. Private preflight found no substantive defect, so no correction was made. No retry, selection among multiple finals, or post-lock rewrite occurred.

10. Final submitted benchmark response

LOCKED SCORED RESPONSE - answer content only. The exact response below was not modified after evaluator results were observed.

```
{
  "topic": "Working abroad",
  "pros": [
    "A new country can broaden cultural understanding and professional experience.",
    "Some roles may offer stronger pay or benefits so a worker may feel fairly compensated.",
    "Daily life in another place can build language skills confidence and an international network."
  ],
  "cons": [
    "Relocation can be expensive and unfamiliar systems may take time to learn.",
    "Immigrants may face language barriers loneliness or difficulty transferring credentials.",
    "Distance from family and uncertainty about housing work or healthcare can add stress."
  ],
  "decision_guide": "Compare pay living costs visa requirements support networks career goals and personal readiness before deciding."
}
```

Locked file: 752 bytes; 14 visible lines; terminal newline present. SHA-256: 477456f283d39ce53bc848d511b872b86c40e122202c25b8f03a263c05394ec1.

11. Official evaluator result

Official Google Research source code was executed against the locked response. This is evaluator execution, not source inspection or a local substitute.

Direct JsonFormat	True
Direct KeywordChecker	True
follow_instruction_list (strict)	[True, True]
follow_all_instructions (strict)	True
follow_instruction_list (loose)	[True, True]
follow_all_instructions (loose)	True
Official evaluator label	OFFICIAL EVALUATOR PASS

JsonFormat strips whitespace, removes one recognized opening or closing Markdown fence when present, and applies json.loads. KeywordChecker sorts the configured keywords and uses case-insensitive substring search. Neither evaluator establishes a particular schema, complete topical coverage, or substantive accuracy.

12. Manual / local constraint check

LOCAL REIMPLEMENTATION - NOT OFFICIAL EVALUATOR. The transparent local check agrees with the official result.

Response byte / visible-line count	752 bytes / 14 lines; terminal newline present
Raw response SHA-256	477456f283d39ce53bc848d511b872b86c40e122202c25b8f03a263c05394ec1
Raw official item SHA-256	fb152545716be53a8cbcbd0c56c98ea60a40bc76526ed324a1c725763acf6eca
Leading non-whitespace bytes	0
Trailing non-whitespace bytes	1
Official-equivalent normalized parse	PASS
Parser exception	None
Parsed top-level JSON type	dict
Raw JSON parse	PASS
Recognized Markdown fence present	No
Exact JSON-only boundary	PASS
Keyword compensated matches	1
Keyword immigrants matches	1
Official-equivalent keyword result	PASS
Audit material in scored block	Absent

Manual task checks - separate from official score

Pros and cons are separately represented	PASS
Meaningful advantages and drawbacks	PASS
Decision guidance is present	PASS
Required terms are used in context	PASS
No invented authority or current fact	PASS
No audit material in scored block	PASS

13. Scores

Instruction-level score	2 / 2 encoded constraints passed
Prompt-level strict score	PASS
Prompt-level loose score	PASS
Final score	1 / 1 prompt-level strict pass
Official vs local discrepancy	None
Manual checks	Reported separately; not added to official IFEval score

14. Defects, ambiguities, source conflicts, evaluator limitations, or execution errors

No source conflict or official evaluator execution error was observed. The initial candidate passed private preflight with zero substantive corrections. No content was changed after lock.

The official JsonFormat checker tests parseability after trimming and removing one recognized Markdown fence at each boundary. It does not enforce a schema, forbid extra JSON fields, or assess whether the object answers the working-abroad task. The official KeywordChecker performs case-insensitive substring searches and does not require standalone words or semantic relevance.

The local manual checks apply a stricter JSON-only boundary, require a dictionary with separate substantive pros and cons plus decision guidance, and review unsupported-claim and audit-material absence. These checks are not additional official IFEval criteria and are not immigration, employment, or financial advice.

ROUTE VERIFICATION NOT EXPOSED remains a material limitation. The requested Luna Max condition is recorded as supplied and is not proof of a hidden backend build or route.

15. Evidence status

Official item evidence	FROZEN - exact raw JSONL line preserved; raw-line SHA-256 recorded
Official checker evidence	FROZEN - instructions.py blob identity recorded
Scored response evidence	LOCKED - SHA-256 recorded before official evaluation
Official evaluator evidence	EXECUTED - direct, strict, and loose results recorded
Local-check evidence	EXECUTED - transparent mirror and manual checks recorded
Runtime-control package	HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified
PDF visual QA	PASS - final rendered pages inspected; no clipping, overlap, missing glyphs, or footer collision observed

16. Record-integrity statement

This is an INTERNAL ONE-SHOT RECORD for the requested FASTPATH 5.6 / LUNA MAX condition. The frozen item, initial candidate, zero disclosed preflight corrections, locked response, evaluator outputs, local checks, and PDF were recorded in order. Locked response SHA-256: 477456f283d39ce53bc848d511b872b86c40e122202c25b8f03a263c05394ec1. Official item-line SHA-256: fb152545716be53a8cbcbd0c56c98ea60a40bc76526ed324a1c725763acf6eca. The record preserves the observed result and does not claim certification, independent verification, leaderboard standing, guaranteed determinism, or proof of hidden backend identity.