

Grounded DI OS Benchmark Record

Prompt 41 of the Grounded DI cross-domain benchmark series - official Google Research IFEval item key 2571

INTERNAL ONE-SHOT RECORD - FASTPATH 5.6 / LUNA MAX - OFFICIAL EVALUATOR PASS

1. Document identity and condition

Primary document	Grounded DI OS Benchmark Record
Benchmark-series position	Cross-domain Prompt 41
Requested condition	Grounded DI OS / FastPath 5.6 / Luna Max / Protocol A
Run controls	Tier 18; DIP 0-166; Invocation 140A / 140; Delta H 0.0041; Benchmark Mode
Benchmark identifier	Google Research IFEval item key 2571
Encoded controls	change_case:english_capital; punctuation:no_comma; length_constraints:number_sentences
Benchmark phase	Composite interaction test - previously covered checker families
Labels applied	INTERNAL ONE-SHOT RECORD; FASTPATH 5.6 / LUNA MAX; OFFICIAL EVALUATOR PASS; LOCAL CHECK PASS
Labels not applied	certified; independently verified; universal; representative; official leaderboard score; proof of hidden backend identity

2. Purpose and scope

This record documents one observed composite execution of official Google Research IFEval item key 2571 under the requested FastPath 5.6 / Luna Max / Protocol A condition. The item combines an English-all-caps control with an ASCII-comma prohibition and a sentence-count threshold, evaluated against one locked response.

The scored task asks which option better describes support for ecological landscapes. Manual review considers whether the response selects the sustainability-oriented option, explains ecological function clearly, and treats aesthetic value as secondary rather than unsupported fact. These content checks remain distinct from the official IFEval score.

The operator-supplied GDRCP v1.1.3 package is treated only as an evidence-harness specification: HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified.

3. Official benchmark URLs

Series architecture: <https://github.com/Grounded-DI/Benchmarks-Cross-Domain>

Audit reference: <https://github.com/Grounded-DI/briefwise-di-performance-on-harvey-lab-benchmark-audit>

Official evaluator tree: https://github.com/google-research/google-research/tree/master/instruction_following_eval

Official data: https://raw.githubusercontent.com/google-research/google-research/master/instruction_following_eval/data/input_data.jsonl

Official checker: https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions.py

Word and sentence utility: https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions_util.py

Paper: <https://arxiv.org/abs/2311.07911>

4. Exact official item (key 2571)

The exact official JSONL line was frozen before response generation. Dataset blob: cbe52f6eecf3986fdac745b4acba4da1408eb146. instructions.py blob: 0007359bd37c7b9c1d9557f8937cb6586d2788a7. instructions_util.py blob: d50702400738df4378fe286c0090a416f9a172fa. Raw-line SHA-256: f7d61dbdc55a72d2dd8c999d9d1e9dd85ff0c5f24f793983bd893dd02c952d51.

```
{"key": 2571, "prompt": "Which of the following is a better way to describe supporting ecological landscapes: (A) the practice of using the natural features of a landscape to create a more sustainable environment, or (B) the practice of using the natural features of a landscape to create a more aesthetically pleasing environment? Your response should be in English, all capital letters, contain no commas, and be fewer than 16 sentences.", "instruction_id_list": [{"change_case": "english_capital", "punctuation": "no_comma", "length_constraints": "number_sentences"}], "kwargs": [{"}, {}], {"relation": "less than", "num_sentences": 16}}}
```

5. Exact benchmark prompt (scored task only)

The following exact prompt is the only scored task. Record material is outside the scored-response boundary.

Which of the following is a better way to describe supporting ecological landscapes: (A) the practice of using the natural features of a landscape to create a more sustainable environment, or (B) the practice of using the natural features of a landscape to create a more aesthetically pleasing environment? Your response should be in English, all capital letters, contain no commas, and be fewer than 16 sentences.

6. Observed runtime identity

Requested condition	Grounded DI OS / FastPath 5.6 / Luna Max / Protocol A
Exposed runtime identity	Codex agent environment; precise backend build or model snapshot not exposed
Route verification	ROUTE VERIFICATION NOT EXPOSED - requested condition is not independent proof of a hidden route
Runtime-control boundary	HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified

7. Timestamp and timezone

Source freeze	2026-09-03 01:16:02 UTC report time; source and item freeze occurred before candidate lock
Initial candidate	One complete candidate preserved before private preflight
Response lock	Locked before official evaluation; no post-lock response mutation
Official evaluation	Recorded after response lock; direct, strict, and loose paths executed
Timezone basis	UTC

8. Tool and web-search status

Official Google Research data and checker sources were retrieved and matched to the listed freeze identities before response generation. No web search was used for the scored answer. No current-fact research, proprietary material, credential, external service, or repository modification was used.

Dependency repair occurred only after response lock: the local evaluator environment made NLTK Punkt and Punkt_tab resources available for the official sentence-count path. The locked response bytes were unchanged.

9. Draft and preflight status

INITIAL DRAFT STATUS	PASS
PRIVATE PREFLIGHT CORRECTIONS	0
Correction	None. The complete candidate already satisfied all three encoded controls and the manual content boundary.
LOCK STATUS	LOCKED BEFORE OFFICIAL EVALUATION
FINAL SUBMISSION STATUS	PASS

One complete candidate was generated. Private preflight found no substantive defect, so no correction was made. No retry, selection among multiple finals, or post-lock rewrite occurred.

10. Final submitted benchmark response

LOCKED SCORED RESPONSE - answer content only. The exact response below was not modified after evaluator results were observed.

```
OPTION A IS THE BETTER DESCRIPTION.
SUPPORTING ECOLOGICAL LANDSCAPES SHOULD WORK WITH NATIVE SOIL WATER AND HABITAT PATTERNS.
THIS APPROACH CAN REDUCE EROSION SUPPORT BIODIVERSITY AND IMPROVE RESILIENCE WHILE STILL ALLOWING BEAUTIFUL DESIGN.
AESTHETIC APPEAL CAN MATTER BUT IT SHOULD FOLLOW ECOLOGICAL FUNCTION RATHER THAN REPLACE IT.
```

Locked file: 335 bytes; 4 visible lines; terminal newline present. SHA-256: 25287b4bd0b229d28425f4cc6815e98a9c638c108b2e7cb3c81483b2dd6d7245.

11. Official evaluator result

Official Google Research source code was executed against the locked response. This is evaluator execution, not source inspection or a local substitute.

Direct CapitalLettersEnglishChecker	True
Direct CommaChecker	True
Direct NumberOfSentences	True
follow_instruction_list (strict)	[True, True, True]
follow_all_instructions (strict)	True
follow_instruction_list (loose)	[True, True, True]
follow_all_instructions (loose)	True
Official evaluator label	OFFICIAL EVALUATOR PASS

CapitalLettersEnglishChecker requires str.isupper() and langdetect.detect(value) == 'en' and returns True on its LangDetectException path. CommaChecker searches only for ASCII U+002C. NumberOfSentences delegates to the Punkt tokenizer. None of these checkers evaluates the ecological reasoning or response usefulness.

12. Manual / local constraint check

LOCAL REIMPLEMENTATION - NOT OFFICIAL EVALUATOR. The transparent local check agrees with the official result.

Response byte / visible-line count	335 bytes / 4 lines; terminal newline present
Raw response SHA-256	25287b4bd0b229d28425f4cc6815e98a9c638c108b2e7cb3c81483b2dd6d7245
Raw official item SHA-256	f7d61dbdc55a72d2dd8c999d9d1e9dd85ff0c5f24f793983bd893dd02c952d51
Leading non-whitespace bytes	0
Trailing non-whitespace bytes	1
str.isupper()	True
Local language detector	en
Language-detector exception	None
Official-equivalent capital result	PASS
ASCII comma count / positions	0 / []
Official-equivalent comma result	PASS
Sentence count / exception	4 / None
Official-equivalent sentence result	PASS
Comma-like non-ASCII marks	{"U+060C ARABIC COMMA": {"count": 0, "positions": []}, "U+3001 IDEOGRAPHIC COMMA": {"count": 0, "positions": []}, "U+FF0C FULLWIDTH COMMA": {"count": 0, "positions": []}}
Selects option A	PASS
Ecological explanation	PASS
Aesthetic value framed as secondary	PASS
Audit material in scored block	Absent

Manual task checks - separate from official score

All-capital English presentation	PASS
No ASCII or comma-like Unicode mark	PASS
Fewer than 16 sentences	PASS - 4 detected
Option A selected and explained	PASS
Ecological function is prioritized	PASS
No invented authority or current fact	PASS
No audit material in scored block	PASS

13. Scores

Instruction-level score	3 / 3 encoded constraints passed
Prompt-level strict score	PASS
Prompt-level loose score	PASS
Final score	1 / 1 prompt-level strict pass
Official vs local discrepancy	None
Manual checks	Reported separately; not added to official IFEval score

14. Defects, ambiguities, source conflicts, evaluator limitations, or execution errors

No source conflict or official evaluator execution error was observed. The initial candidate passed private preflight with zero substantive corrections. No content was changed after lock.

CapitalLettersEnglishChecker combines Python's isupper predicate with language detection and returns True on its LangDetectException path. CommaChecker detects only ASCII U+002C and does not reject comma-like Unicode punctuation. NumberOfSentences relies on the official NLTK Punkt tokenizer and does not judge sentence quality or meaning.

The official constraints do not assess whether the answer accurately reasons about ecological landscapes, whether the prose is useful, or whether unsupported claims are made. The local manual checks are separate and are not ecological or environmental certification.

ROUTE VERIFICATION NOT EXPOSED remains a material limitation. The requested Luna Max condition is recorded as supplied and is not proof of a hidden backend build or route.

15. Evidence status

Official item evidence	FROZEN - exact raw JSONL line preserved; raw-line SHA-256 recorded
Official checker evidence	FROZEN - instructions.py blob identity recorded
Sentence utility evidence	FROZEN - instructions_util.py blob identity recorded
Scored response evidence	LOCKED - SHA-256 recorded before official evaluation
Official evaluator evidence	EXECUTED - direct, strict, and loose results recorded
Local-check evidence	EXECUTED - transparent mirror and manual checks recorded
Runtime-control package	HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified
PDF visual QA	PASS - final rendered pages inspected; no clipping, overlap, missing glyphs, or footer collision observed

16. Record-integrity statement

This is an INTERNAL ONE-SHOT RECORD for the requested FASTPATH 5.6 / LUNA MAX condition. The frozen item, initial candidate, zero disclosed preflight corrections, locked response, evaluator outputs, local checks, and PDF were recorded in order. Locked response SHA-256: 25287b4bd0b229d28425f4cc6815e98a9c638c108b2e7cb3c81483b2dd6d7245. Official item-line SHA-256: f7d61dbdc55a72d2dd8c999d9d1e9dd85ff0c5f24f793983bd893dd02c952d51. The record preserves the observed result and does not claim certification, independent verification, leaderboard standing, guaranteed determinism, or proof of hidden backend identity.