

Grounded DI OS Benchmark Record

Prompt 42 of the Grounded DI cross-domain benchmark series - official Google Research IFEval item key 3276

INTERNAL ONE-SHOT RECORD - FASTPATH 5.6 / LUNA MAX - OFFICIAL EVALUATOR PASS

1. Document identity and condition

Primary document	Grounded DI OS Benchmark Record
Benchmark-series position	Cross-domain Prompt 42
Requested condition	Grounded DI OS / FastPath 5.6 / Luna Max / Protocol A
Run controls	Tier 18; DIP 0-166; Invocation 140A / 140; Delta H 0.0041; Benchmark Mode
Benchmark identifier	Google Research IFEval item key 3276
Encoded controls	change_case:english_capital; length_constraints:number_sentences (at least 5; less than 7)
Benchmark phase	Composite interaction test - three encoded controls including two same-family thresholds
Labels applied	INTERNAL ONE-SHOT RECORD; FASTPATH 5.6 / LUNA MAX; OFFICIAL EVALUATOR PASS; LOCAL CHECK PASS
Labels not applied	certified; independently verified; universal; representative; official leaderboard score; proof of hidden backend identity

2. Purpose and scope

This record documents one observed composite execution of official Google Research IFEval item key 3276 under the requested FastPath 5.6 / Luna Max / Protocol A condition. The item combines an English-all-caps control with two independent sentence-count thresholds, evaluated against one locked response.

The scored task asks whether the sentence "He hurried through the archaic rooms of the museum" has grammatical errors. Manual review checks the grammatical conclusion and the supporting identification of the subject, verb, adjective, and prepositional phrase. These semantic checks remain distinct from the official IFEval score.

The operator-supplied GDRCP v1.1.3 package is treated only as an evidence-harness specification: HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified.

3. Official benchmark URLs

Series architecture: <https://github.com/Grounded-DI/Benchmarks-Cross-Domain>

Audit reference: <https://github.com/Grounded-DI/briefwise-di-performance-on-harvey-lab-benchmark-audit>

Official evaluator tree: https://github.com/google-research/google-research/tree/master/instruction_following_eval

Official data: https://raw.githubusercontent.com/google-research/google-research/master/instruction_following_eval/data/input_data.jsonl

Official checker: https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions.py

Word and sentence utility: https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions_util.py

Paper: <https://arxiv.org/abs/2311.07911>

4. Exact official item (key 3276)

The exact official JSONL line was frozen before response generation. Dataset blob: cbe52f6eecf3986fdac745b4acba4da1408eb146. instructions.py blob: 0007359bd37c7b9c1d9557f8937cb6586d2788a7. instructions_util.py blob: d50702400738df4378fe286c0090a416f9a172fa. Raw-line SHA-256: f24128da29c73980dec862e344d6c5f3118094875624d1757a30d21b99b5753a.

```
{
  "key": 3276,
  "prompt": "Does the sentence \"He hurried through the archaic rooms of the museum\" have any grammatical errors? Answer in all capital letters, and organize your entire response in 5 or 6 sentences.",
  "instruction_id_list": [
    "change_case:english_capital",
    "length_constraints: number_sentences",
    "length_constraints:number_sentences"
  ],
  "kwargs": [
    {},
    {
      "relation": "at least",
      "num_sentences": 5
    },
    {
      "relation": "less than",
      "num_sentences": 7
    }
  ]
}
```

5. Exact benchmark prompt (scored task only)

The following exact prompt is the only scored task. Record material is outside the scored-response boundary.

```
Does the sentence "He hurried through the archaic rooms of the museum" have any grammatical errors? Answer in all capital letters, and organize your entire response in 5 or 6 sentences.
```

6. Observed runtime identity

Requested condition	Grounded DI OS / FastPath 5.6 / Luna Max / Protocol A
Exposed runtime identity	Codex agent environment; precise backend build or model snapshot not exposed
Route verification	ROUTE VERIFICATION NOT EXPOSED - requested condition is not independent proof of a hidden route
Runtime-control boundary	HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified

7. Timestamp and timezone

Source freeze	2026-09-03 02:14:13 UTC report time; source and item freeze occurred before candidate lock
Initial candidate	One complete candidate preserved before private preflight
Response lock	Locked before official evaluation; no post-lock response mutation
Official evaluation	Recorded after response lock; direct, strict, and loose paths executed
Timezone basis	UTC

8. Tool and web-search status

Official Google Research data and checker sources were retrieved and matched to the listed freeze identities before response generation. No web search was used for the scored answer. No current-fact research, proprietary material, credential, external service, or repository modification was used.

The prepared evaluator environment supplied NLTK Punkt and Punkt_tab resources for the official sentence-count path; no response bytes were changed after lock.

9. Draft and preflight status

INITIAL DRAFT STATUS	PASS
PRIVATE PREFLIGHT CORRECTIONS	0
Correction	None. The complete candidate already satisfied all three encoded controls and the manual content boundary.
LOCK STATUS	LOCKED BEFORE OFFICIAL EVALUATION
FINAL SUBMISSION STATUS	PASS

One complete candidate was generated. Private preflight found no substantive defect, so no correction was made. No retry, selection among multiple finals, or post-lock rewrite occurred.

10. Final submitted benchmark response

LOCKED SCORED RESPONSE - answer content only. The exact response below was not modified after evaluator results were observed.

THE SENTENCE HAS NO GRAMMATICAL ERRORS. HE IS THE SUBJECT AND HURRIED IS THE PAST-TENSE VERB. THE ADJECTIVE ARCHAIC CORRECTLY MODIFIES ROOMS. THE PREPOSITIONAL PHRASE OF THE MUSEUM SPECIFIES WHICH ROOMS ARE MEANT. THE SENTENCE IS COMPLETE AND NATURAL IN STANDARD ENGLISH.

Locked file: 272 bytes; 5 visible lines; terminal newline present. SHA-256: 0df3fa1ad04fe8607655330d9d34765ecbd4f5abd80ef84636823e6aebb6a549.

11. Official evaluator result

Official Google Research source code was executed against the locked response. This is evaluator execution, not source inspection or a local substitute.

Direct CapitalLettersEnglishChecker	True
Direct NumberOfSentences (at least 5)	True
Direct NumberOfSentences (less than 7)	True
follow_instruction_list (strict)	[True, True, True]
follow_all_instructions (strict)	True
follow_instruction_list (loose)	[True, True, True]
follow_all_instructions (loose)	True
Official evaluator label	OFFICIAL EVALUATOR PASS

CapitalLettersEnglishChecker requires str.isupper() and langdetect.detect(value) == 'en' and returns True on its LangDetectException path. The two NumberOfSentences checks use the official Punkt tokenizer and independently apply the at-least-five and less-than-seven thresholds. None checks grammatical correctness or usefulness.

12. Manual / local constraint check

LOCAL REIMPLEMENTATION - NOT OFFICIAL EVALUATOR. The transparent local check agrees with the official result.

Response byte / visible-line count	272 bytes / 5 lines; terminal newline present
Raw response SHA-256	0df3fa1ad04fe8607655330d9d34765ecbd4f5abd80ef84636823e6aebb6a549
Raw official item SHA-256	f24128da29c73980dec862e344d6c5f3118094875624d1757a30d21b99b5753a
Leading / trailing non-whitespace bytes	0 / 1
str.isupper()	True
Local language detector	en
Language-detector exception	None
Lowercase ASCII letters / non-ASCII cased	0 / []
Official-equivalent capital result	PASS
Official utility sentence count	5
Independent sentence count / exception	5 / None
Official-equivalent at-least-five result	PASS
Official-equivalent less-than-seven result	PASS
Exact 5/6 sentence window	PASS
Period byte positions	[38, 92, 140, 212, 270]
Audit material in scored block	Absent

Manual task checks - separate from official score

All-capital English presentation	PASS
Exactly five complete sentences	PASS
No sentence-count ambiguity in the locked block	PASS
Grammar conclusion: no errors	PASS
Subject identified	PASS
Past-tense verb identified	PASS
Adjective and prepositional phrase identified	PASS
No invented authority or current fact	PASS
No audit material in scored block	PASS

13. Scores

Instruction-level score	3 / 3 encoded constraints passed
Prompt-level strict score	PASS
Prompt-level loose score	PASS
Final score	1 / 1 prompt-level strict pass
Official vs local discrepancy	None
Manual checks	Reported separately; not added to official IFEval score

14. Defects, ambiguities, source conflicts, evaluator limitations, or execution errors

No source conflict or official evaluator execution error was observed. The initial candidate passed private preflight with zero substantive corrections. No content was changed after lock.

CapitalLettersEnglishChecker combines Python's isupper predicate with language detection and returns True on its LangDetectException path. The two sentence-count paths use the official tokenizer and do not establish that the prose is grammatical or meaningful.

The official constraints do not assess whether the grammar explanation is accurate, whether all phrase labels are useful, or whether unsupported claims are made. Local and manual checks are separate and are not linguistic certification.

ROUTE VERIFICATION NOT EXPOSED remains a material limitation. The requested Luna Max condition is recorded as supplied and is not proof of a hidden backend build or route.

15. Evidence status

Official item evidence	FROZEN - exact raw JSONL line preserved; raw-line SHA-256 recorded
Official checker evidence	FROZEN - instructions.py blob identity recorded
Sentence utility evidence	FROZEN - instructions_util.py blob identity recorded
Scored response evidence	LOCKED - SHA-256 recorded before official evaluation
Official evaluator evidence	EXECUTED - direct, strict, and loose results recorded
Local-check evidence	EXECUTED - transparent mirror and manual checks recorded
Runtime-control package	HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified
PDF visual QA	PASS - final rendered pages inspected; no clipping, overlap, missing glyphs, or footer collision observed

16. Record-integrity statement

This is an INTERNAL ONE-SHOT RECORD for the requested FASTPATH 5.6 / LUNA MAX condition. The frozen item, initial candidate, zero disclosed preflight corrections, locked response, evaluator outputs, local checks, and PDF were recorded in order. Locked response SHA-256: 0df3fa1ad04fe8607655330d9d34765ecbd4f5abd80ef84636823e6aebb6a549. Official item-line SHA-256: f24128da29c73980dec862e344d6c5f3118094875624d1757a30d21b99b5753a. The record preserves the observed result and does not claim certification, independent verification, leaderboard standing, guaranteed determinism, or proof of hidden backend identity.