

Grounded DI OS Benchmark Record

Prompt 43 of the Grounded DI cross-domain benchmark series - official Google Research IFEval item key 1021

INTERNAL ONE-SHOT RECORD - FASTPATH 5.6 / LUNA MAX - OFFICIAL EVALUATOR PASS

1. Document identity and condition

Primary document	Grounded DI OS Benchmark Record
Benchmark-series position	Cross-domain Prompt 43
Requested condition	Grounded DI OS / FastPath 5.6 / Luna Max / Protocol A
Run controls	Tier 18; DIP 0-166; Invocation 140A / 140; Delta H 0.0041; Benchmark Mode
Benchmark identifier	Google Research IFEval item key 1021
Encoded controls	change_case:english_capital; detectable_format:multiple_sections (PARAGRAPH; 2 sections)
Benchmark phase	Composite interaction test - all-capital English plus labeled two-section structure
Labels applied	INTERNAL ONE-SHOT RECORD; FASTPATH 5.6 / LUNA MAX; OFFICIAL EVALUATOR PASS; LOCAL CHECK PASS
Labels not applied	certified; independently verified; universal; representative; official leaderboard score; proof of hidden backend identity

2. Purpose and scope

This record documents one observed composite execution of official Google Research IFEval item key 1021 under the requested FastPath 5.6 / Luna Max / Protocol A condition. The item combines an English-all-caps control with a two-section marker control and is evaluated against one locked response.

The scored task asks for a critique of the sentence "If the law is bad, you should not follow it". Manual review checks that the response explains why the sentence is overbroad and distinguishes legal validity from moral duty. These conceptual checks remain distinct from the official IFEval score and are not legal advice.

The operator-supplied GDRCP v1.1.3 package is treated only as an evidence-harness specification: HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified.

3. Official benchmark URLs

Series architecture: <https://github.com/Grounded-DI/Benchmarks-Cross-Domain>
Audit reference: <https://github.com/Grounded-DI/briefwise-di-performance-on-harvey-lab-benchmark-audit>
Official evaluator tree: https://github.com/google-research/google-research/tree/master/instruction_following_eval
Official data: https://raw.githubusercontent.com/google-research/google-research/master/instruction_following_eval/data/input_data.jsonl
Official checker: https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions.py
Paper: <https://arxiv.org/abs/2311.07911>

4. Exact official item (key 1021)

The exact official JSONL line was frozen before response generation. Dataset blob: cbe52f6eecf3986fdac745b4acba4da1408eb146. instructions.py blob: 0007359bd37c7b9c1d9557f8937cb6586d2788a7. Raw-line SHA-256: b912ded8ced908de9ad46606a75a312695e0493adc6b568403bc3f2f17921a6a.

```
{
  "key": 1021,
  "prompt": "Write a 2 paragraph critique of the following sentence in all capital letters, no lowercase letters allowed: \\\"If the law is bad, you should not follow it\\\". Label each paragraph with PARAGRAPH X.",
  "instruction_id_list": [
    "change_case:english_capital",
    "detectable_format:multiple_sections"
  ],
  "kwargs": {
    {},
    {
      "section_splitter": "PARAGRAPH",
      "num_sections": 2
    }
  ]
}
```

5. Exact benchmark prompt (scored task only)

The following exact prompt is the only scored task. Record material is outside the scored-response boundary.

```
Write a 2 paragraph critique of the following sentence in all capital letters, no lowercase letters allowed: "If the law is bad, you should not follow it". Label each paragraph with PARAGRAPH X.
```

6. Observed runtime identity

Requested condition	Grounded DI OS / FastPath 5.6 / Luna Max / Protocol A
Exposed runtime identity	Codex agent environment; precise backend build or model snapshot not exposed
Route verification	ROUTE VERIFICATION NOT EXPOSED - requested condition is not independent proof of a hidden route
Runtime-control boundary	HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified

7. Timestamp and timezone

Source freeze	2026-09-03 02:50:51 UTC report time; source and item freeze occurred before candidate lock
Initial candidate	One complete candidate preserved before private preflight
Response lock	Locked before official evaluation; no post-lock response mutation
Official evaluation	Recorded after response lock; direct, strict, and loose paths executed
Timezone basis	UTC

8. Tool and web-search status

Official Google Research data and checker sources were retrieved and matched to the listed freeze identities before response generation. No web search was used for the scored answer. No current-fact research, proprietary material, credential, external service, or repository modification was used.

No dependency repair was needed for the executed checker paths. The locked response bytes were unchanged throughout evaluation.

9. Draft and preflight status

INITIAL DRAFT STATUS	PASS
PRIVATE PREFLIGHT CORRECTIONS	0
Correction	None. The complete candidate already satisfied both encoded controls and the manual content boundary.
LOCK STATUS	LOCKED BEFORE OFFICIAL EVALUATION
FINAL SUBMISSION STATUS	PASS

One complete candidate was generated. Private preflight found no substantive defect, so no correction was made. No retry, selection among multiple finals, or post-lock rewrite occurred.

10. Final submitted benchmark response

LOCKED SCORED RESPONSE - answer content only. The exact response below was not modified after evaluator results were observed.

```
PARAGRAPH 1
THE SENTENCE CAPTURES A MORAL INTUITION BUT IT IS TOO BROAD. A LAW CAN BE UNJUST YET STILL BE A LAW. CALLING A RULE BAD DOES NOT BY ITSELF
IDENTIFY THE HARM OR SHOW WHETHER DISOBEDIENCE IS JUSTIFIED.

PARAGRAPH 2
A STRONGER CRITIQUE WOULD SEPARATE LEGAL VALIDITY FROM MORAL DUTY. PEOPLE CAN SEEK REVIEW REPEAL OR OTHER LAWFUL REMEDIES AND MAY CONSIDER
CONSCIENTIOUS CIVIL DISOBEDIENCE WHEN A RULE GRAVELY VIOLATES BASIC RIGHTS. THE BEST VERSION THEREFORE INVITES REASONED JUDGMENT RATHER
THAN AUTOMATIC OBEDIENCE OR AUTOMATIC DEFIANCE.
```

Locked file: 548 bytes; 5 visible lines; terminal newline present. SHA-256: 416775968814b0402c7a576f71e539b9c2456b5c9957119e9216405020b33f58.

11. Official evaluator result

Official Google Research source code was executed against the locked response. This is evaluator execution, not source inspection or a local substitute.

Direct CapitalLettersEnglishChecker	True
Direct SectionChecker (PARAGRAPH; 2)	True
follow_instruction_list (strict)	[True, True]
follow_all_instructions (strict)	True
follow_instruction_list (loose)	[True, True]
follow_all_instructions (loose)	True
Official evaluator label	OFFICIAL EVALUATOR PASS

CapitalLettersEnglishChecker requires str.isupper() and langdetect.detect(value) == 'en' and returns True on its LangDetectException path. SectionChecker splits on its PARAGRAPH-plus-number pattern and counts markers; it does not independently prove exactly two prose paragraphs. Neither checker evaluates legal correctness or usefulness.

12. Manual / local constraint check

LOCAL REIMPLEMENTATION - NOT OFFICIAL EVALUATOR. The transparent local check agrees with the official result.

Response byte / visible-line count	548 bytes / 5 lines; terminal newline present
Raw response SHA-256	416775968814b0402c7a576f71e539b9c2456b5c9957119e9216405020b33f58
Raw official item SHA-256	b912ded8ced908de9ad46606a75a312695e0493adc6b568403bc3f2f17921a6a
Leading / trailing non-whitespace bytes	0 / 1
str.isupper()	True
Lowercase ASCII letters / non-ASCII cased	0 / []
Local language detector / exception	en / None
Section marker count / spans	2 / [[0, 12], [212, 225]]
Local non-empty section count	2
All local sections non-empty	True
Official-equivalent capital result	PASS
Official-equivalent section result	PASS
Exact two-section result	PASS
Audit material in scored block	Absent

Manual task checks - separate from official score

All-capital English presentation	PASS
Exactly two labeled non-empty sections	PASS
Each label uses PARAGRAPH plus a number	PASS
First section identifies the sentence as overbroad	PASS
Second section separates legal validity from moral duty	PASS
Lawful remedies and civil disobedience are qualified	PASS
No invented authority or current fact	PASS
No audit material in scored block	PASS

13. Scores

Instruction-level score	2 / 2 encoded constraints passed
Prompt-level strict score	PASS
Prompt-level loose score	PASS
Final score	1 / 1 prompt-level strict pass
Official vs local discrepancy	None
Manual checks	Reported separately; not added to official IFEval score

14. Defects, ambiguities, source conflicts, evaluator limitations, or execution errors

No source conflict or official evaluator execution error was observed. The initial candidate passed private preflight with zero substantive corrections. No content was changed after lock.

CapitalLettersEnglishChecker combines Python's isupper predicate with language detection and returns True on its LangDetectException path. SectionChecker recognizes marker patterns and counts them; it does not prove paragraph semantics or prohibit extra text.

The official constraints do not assess whether the critique is legally complete, whether a particular act of disobedience is lawful, or whether the conceptual analysis is persuasive. The local and manual checks are separate and are not legal advice or independent legal certification.

ROUTE VERIFICATION NOT EXPOSED remains a material limitation. The requested Luna Max condition is recorded as supplied and is not proof of a hidden backend build or route.

15. Evidence status

Official item evidence	FROZEN - exact raw JSONL line preserved; raw-line SHA-256 recorded
Official checker evidence	FROZEN - instructions.py blob identity recorded
Scored response evidence	LOCKED - SHA-256 recorded before official evaluation
Official evaluator evidence	EXECUTED - direct, strict, and loose results recorded
Local-check evidence	EXECUTED - transparent mirror and manual checks recorded
Runtime-control package	HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified
PDF visual QA	PASS - final rendered pages inspected; no clipping, overlap, missing glyphs, or footer collision observed

16. Record-integrity statement

This is an INTERNAL ONE-SHOT RECORD for the requested FASTPATH 5.6 / LUNA MAX condition. The frozen item, initial candidate, zero disclosed preflight corrections, locked response, evaluator outputs, local checks, and PDF were recorded in order. Locked response SHA-256: 416775968814b0402c7a576f71e539b9c2456b5c9957119e9216405020b33f58. Official item-line SHA-256: b912ded8ced908de9ad46606a75a312695e0493adc6b568403bc3f2f17921a6a. The record preserves the observed result and does not claim certification, independent verification, leaderboard standing, guaranteed determinism, legal advice, or proof of hidden backend identity.