

Grounded DI OS Benchmark Record

Prompt 45 of the Grounded DI cross-domain benchmark series - official Google Research IFEval item key 349

INTERNAL ONE-SHOT RECORD - FASTPATH 5.6 / LUNA MAX - OFFICIAL EVALUATOR PASS

1. Document identity and condition

Primary document	Grounded DI OS Benchmark Record
Benchmark-series position	Cross-domain Prompt 45
Requested condition	Grounded DI OS / FastPath 5.6 / Luna Max / Protocol A
Run controls	Tier 18; DIP 0-166; Invocation 140A / 140; Delta H 0.0041; Benchmark Mode
Benchmark identifier	Google Research IFEval item key 349
Encoded controls	detectable_format:number_highlighted_sections; length_constraints:number_words
Benchmark phase	Composite interaction test - practical garden watering advice with exactly three italic sections and fewer than 40 words
Labels applied	INTERNAL ONE-SHOT RECORD; FASTPATH 5.6 / LUNA MAX; OFFICIAL EVALUATOR PASS; LOCAL CHECK PASS
Labels not applied	certified; independently verified; universal; representative; official leaderboard score; proof of hidden backend identity

2. Purpose and scope

This record documents one observed composite execution of official Google Research IFEval item key 349 under the requested FastPath 5.6 / Luna Max / Protocol A condition. The item combines an italic-section control with a strict word budget and is evaluated against one locked response.

The scored task asks for concise advice on watering a small backyard garden for vegetables and flowers. Manual review checks that the advice is meaningful and practically coherent while keeping those quality judgments separate from the official IFEval score.

The operator-supplied GDRCP v1.1.3 package is treated only as an evidence-harness specification: HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified.

3. Official benchmark URLs

Series architecture: <https://github.com/Grounded-DI/Benchmarks-Cross-Domain>
Audit reference: <https://github.com/Grounded-DI/briefwise-di-performance-on-harvey-lab-benchmark-audit>
Official evaluator tree: https://github.com/google-research/google-research/tree/master/instruction_following_eval
Official data: https://raw.githubusercontent.com/google-research/google-research/master/instruction_following_eval/data/input_data.jsonl
Official checker: https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions.py
Word-count utility: https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions_util.py
Paper: <https://arxiv.org/abs/2311.07911>

4. Exact official item (key 349)

The exact official JSONL line was frozen before response generation. Dataset blob: cbe52f6eef3986fdac745b4acba4da1408eb146. instructions.py blob: 0007359bd37c7b9c1d9557f8937cb6586d2788a7. instructions_util.py blob: d50702400738df4378fe286c0090a416f9a172fa. Raw-line SHA-256: 0f1371fd570e070c904398254c1573142513b8970f8706a998d7f6cae8e635b9.

```
{"key": 349, "prompt": "I want to start a garden for veggies and flowers in my small backyard. Can you give me some advice on how to water my plants? Have at least 3 italic text sections, such as: *italic text 1*, *italic text 2*, etc.\n\nMake your reply short -- the whole reply should contain less than 40 words.", "instruction_id_list": ["detectable_format:number_highlighted_sections", "length_constraints:number_words"], "kwargs": [{"num_highlights": 3}, {"relation": "less than", "num_words": 40}]}
```

5. Exact benchmark prompt (scored task only)

The following exact prompt is the only scored task. Record material is outside the scored-response boundary.

I want to start a garden for veggies and flowers in my small backyard. Can you give me some advice on how to water my plants? Have at least 3 italic text sections, such as: *italic text 1*, *italic text 2*, etc.

Make your reply short -- the whole reply should contain less than 40 words.

6. Observed runtime identity

Requested condition	Grounded DI OS / FastPath 5.6 / Luna Max / Protocol A
Exposed runtime identity	Codex agent environment; precise backend build or model snapshot not exposed
Route verification	ROUTE VERIFICATION NOT EXPOSED - requested condition is not independent proof of a hidden route
Runtime-control boundary	HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified

7. Timestamp and timezone

Source freeze	2026-09-03 03:56:25 UTC report time; source and item freeze occurred before candidate lock
Initial candidate	One complete candidate preserved before private preflight
Response lock	Locked before official evaluation; no post-lock response mutation
Official evaluation	Recorded after response lock; direct, strict, and loose paths executed
Timezone basis	UTC

8. Tool and web-search status

Official Google Research data and checker sources were retrieved and matched to the listed freeze identities before response generation. No web search was used for the scored answer. No current-fact research, proprietary material, credential, or external service was needed.

No dependency repair was needed for the executed checker paths. The locked response bytes were unchanged throughout evaluation; repository publication is conditional on the full pass and artifact checks.

9. Draft and preflight status

INITIAL DRAFT STATUS	PASS
PRIVATE PREFLIGHT CORRECTIONS	0
Correction	None. The complete candidate already satisfied both encoded controls and the manual content boundary.
LOCK STATUS	LOCKED BEFORE OFFICIAL EVALUATION
FINAL SUBMISSION STATUS	PASS

One complete candidate was generated. Private preflight found no substantive defect, so no correction was made. No retry, selection among multiple finals, or post-lock rewrite occurred.

10. Final submitted benchmark response

LOCKED SCORED RESPONSE - answer content only. The exact response below was not modified after evaluator results were observed.

```
Water *deeply* at the roots when the soil feels dry. Check *soil moisture* before watering again. Use *morning watering* to reduce evaporation.
```

Locked file: 144 bytes; 1 visible lines; terminal newline present. SHA-256: 55a479e15b205b958b571a57780bc17b958200976760e8e1fdf647c120c13f9e.

11. Official evaluator result

Official Google Research source code was executed against the locked response. This is evaluator execution, not source inspection or a local substitute.

Direct HighlightSectionChecker	True
Direct NumberOfWords	True
follow_instruction_list (strict)	[True, True]
follow_all_instructions (strict)	True
follow_instruction_list (loose)	[True, True]
follow_all_instructions (loose)	True
Official evaluator label	OFFICIAL EVALUATOR PASS

HighlightSectionChecker searches for non-empty single-asterisk and double-asterisk spans. NumberOfWords counts tokens using the official RegexpTokenizer(r'\w+'). Neither checker evaluates the horticultural quality, local conditions, or usefulness of the advice.

12. Manual / local constraint check

LOCAL REIMPLEMENTATION - NOT OFFICIAL EVALUATOR. The transparent local check agrees with the official result.

Response byte / visible-line count	144 bytes / 1 line; terminal newline present
Raw response SHA-256	55a479e15b205b958b571a57780bc17b958200976760e8e1fdf647c120c13f9e
Raw official item SHA-256	0f1371fd570e070c904398254c1573142513b8970f8706a998d7f6cae8e635b9
Leading / trailing non-whitespace bytes	0 / 1
Official-equivalent highlighted sections	3
Single / double highlight counts	3 / 0
Highlight match values	['*deeply*', '*soil moisture*', '*morning watering*']
Official-equivalent word-token count	22
Strictly below 40 words	PASS
Non-whitespace material outside answer	Absent
Meaningful watering advice	PASS
Official-equivalent highlight result	PASS
Official-equivalent word result	PASS
Audit material in scored block	Absent

Manual task checks - separate from official score

Exactly three italic sections appear	PASS
Advice addresses roots and soil moisture	PASS
Morning watering and evaporation guidance	PASS
Concise non-empty practical explanation	PASS
No unsupported authority, policy, or current fact	PASS
No audit material in scored block	PASS

13. Scores

Instruction-level score	2 / 2 encoded constraints passed
Prompt-level strict score	PASS
Prompt-level loose score	PASS
Final score	1 / 1 prompt-level strict pass
Official vs local discrepancy	None
Manual checks	Reported separately; not added to official IFEval score

14. Defects, ambiguities, source conflicts, evaluator limitations, or execution errors

No source conflict or official evaluator execution error was observed. The initial candidate passed private preflight with zero substantive corrections. No content was changed after lock.

HighlightSectionChecker recognizes certain single- and double-asterisk spans and does not prove conventional Markdown rendering, exact semantic emphasis, or absence of extra prose. NumberOfWords counts official \w+ tokens rather than judging brevity or sentence quality.

The official constraints do not assess whether the watering advice fits a particular climate, soil, plant species, or season. Manual checks are general content review, not horticultural certification.

ROUTE VERIFICATION NOT EXPOSED remains a material limitation. The requested Luna Max condition is recorded as supplied and is not proof of a hidden backend build or route.

15. Evidence status

Official item evidence	FROZEN - exact raw JSONL line preserved; raw-line SHA-256 recorded
Official checker evidence	FROZEN - instructions.py and instructions_util.py blob identities recorded
Scored response evidence	LOCKED - SHA-256 recorded before official evaluation
Official evaluator evidence	EXECUTED - direct, strict, and loose results recorded
Local-check evidence	EXECUTED - transparent mirror and manual checks recorded
Runtime-control package	HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified
PDF visual QA	PASS - final rendered pages inspected; no clipping, overlap, missing glyphs, or footer collision observed

16. Record-integrity statement

This is an INTERNAL ONE-SHOT RECORD for the requested FASTPATH 5.6 / LUNA MAX condition. The frozen item, initial candidate, zero disclosed preflight corrections, locked response, evaluator outputs, local checks, and PDF were recorded in order. Locked response SHA-256: 55a479e15b205b958b571a57780bc17b958200976760e8e1fdf647c120c13f9e. Official item-line SHA-256: 0f1371fd570e070c904398254c1573142513b8970f8706a998d7f6cae8e635b9. The record preserves the observed result and does not claim certification, independent verification, leaderboard standing, guaranteed determinism, horticultural certification, or proof of hidden backend identity.