

Grounded DI OS Benchmark Record

Prompt 46 of the Grounded DI cross-domain benchmark series - official Google Research IFEval item key 371

INTERNAL ONE-SHOT RECORD - FASTPATH 5.6 / LUNA MAX - OFFICIAL EVALUATOR PASS

1. Document identity and condition

Primary document	Grounded DI OS Benchmark Record
Benchmark-series position	Cross-domain Prompt 46
Requested condition	Grounded DI OS / FastPath 5.6 / Luna Max / Protocol A
Run controls	Tier 18; DIP 0-166; Invocation 140A / 140; Delta H 0.0041; Benchmark Mode
Benchmark identifier	Google Research IFEval item key 371
Encoded controls	detectable_format:json_format; keywords:forbidden_words
Benchmark phase	Blindly selected composite interaction test - riddle content constrained by JSON structure and a forbidden target word
Labels applied	INTERNAL ONE-SHOT RECORD; FASTPATH 5.6 / LUNA MAX; OFFICIAL EVALUATOR PASS; LOCAL CHECK PASS
Labels not applied	certified; independently verified; universal; representative; official leaderboard score; proof of hidden backend identity

2. Purpose and blind selection scope

This record documents one observed composite execution of official Google Research IFEval item key 371 under the requested FastPath 5.6 / Luna Max / Protocol A condition. The item was selected before its natural-language prompt was retrieved.

The benchmark selector was deliberately blind to prompt text, topic, answer, and content-derived features. It saw only fresh-item keys, ordered instruction IDs, and ordered kwargs from the official dataset. It preferred two independent deterministic checker families and then applied the stated smallest-burden and lowest-key tie-breaks.

Blind selector result	selected_key = 371
Selector input	Metadata only; 541 official rows; 26 unique previously represented keys excluded
Selector shortlist	Seven eligible JSON-format plus forbidden-word metadata records; key 371 had one forbidden word and won the lowest-key tie-break
Selector prompt SHA-256	947aa34b9e787648a7892baf7738f22a8354c00cc9993360522be1189381cec4
Selection boundary	Exact official prompt fetched only after key 371 was frozen

The operator-supplied GDRCP v1.1.3 package is treated only as an evidence-harness specification: HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified.

3. Official benchmark URLs

Series architecture: <https://github.com/Grounded-DI/Benchmarks-Cross-Domain>

Audit reference: <https://github.com/Grounded-DI/briefwise-di-performance-on-harvey-lab-benchmark-audit>

Official evaluator tree: https://github.com/google-research/google-research/tree/master/instruction_following_eval

Official data: https://raw.githubusercontent.com/google-research/google-research/master/instruction_following_eval/data/input_data.jsonl

Official checker: https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions.py

Evaluator library: https://github.com/google-research/google-research/blob/master/instruction_following_eval/evaluation_lib.py

Word-count utility: https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions_util.py

Paper: <https://arxiv.org/abs/2311.07911>

4. Exact official item (key 371)

The exact official JSONL line was frozen after blind selection and before response generation. Dataset blob: cbe52f6eecf3986fdac745b4acba4da1408eb146. instructions.py blob: 0007359bd37c7b9c1d9557f8937cb6586d2788a7. instructions_util.py blob: d50702400738df4378fe286c0090a416f9a172fa. evaluation_lib.py blob: c87422822f9e1565ca7574ac0b3cee5389c9f44e. Raw-line SHA-256: 021cca041229fbcadb2cbc6c3bf224bf1f839da86486fd8ac9eea07eeeb314d0.

```
{
  "key": 371,
  "prompt": "Write a riddle that describes the word 'key' but doesn't use the word 'key'. Wrap all words into one JSON block. The word 'key' should not appear in your entire reply.",
  "instruction_id_list": [
    "detectable_format:json_format",
    "keywords:forbidden_words"
  ],
  "kwargs": [
    {},
    {"forbidden_words": ["key"]}
  ]
}
```

5. Exact benchmark prompt (scored task only)

The following exact prompt is the only scored task. Record material is outside the scored-response boundary.

Write a riddle that describes the word "key" but doesn't use the word "key". Wrap all words into one JSON block. The word "key" should not appear in your entire reply.

6. Observed runtime identity

Requested condition	Grounded DI OS / FastPath 5.6 / Luna Max / Protocol A
Exposed runtime identity	Codex agent environment; precise backend build or model snapshot not exposed
Route verification	ROUTE VERIFICATION NOT EXPOSED - requested condition is not independent proof of a hidden route
Runtime-control boundary	HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified

7. Timestamp and timezone

Source freeze	2026-09-03 15:02:33 UTC report time; blind metadata selection and item freeze occurred before candidate lock
Initial candidate	One complete candidate preserved before private preflight
Response lock	Locked before official evaluation; no post-lock response mutation
Official evaluation	Recorded after response lock; direct, strict, and loose paths executed
Timezone basis	UTC

8. Tool and web-search status

The official Google Research dataset and checker sources were retrieved through the selected GitHub source and matched to the listed freeze identities before response generation. No web search was used for the scored answer. No current-fact research, proprietary material, credential, or external service was needed.

The selector saw metadata only. The exact prompt and its content were not used for selection. No dependency repair was needed for the executed checker paths. The locked response bytes were unchanged throughout evaluation; repository publication remains conditional on the full pass and artifact checks.

9. Draft and preflight status

INITIAL DRAFT STATUS	PASS
PRIVATE PREFLIGHT CORRECTIONS	0
Correction	None. The complete candidate already satisfied valid JSON, a non-empty riddle, whole-word absence of the forbidden target, a clear target description, and the manual content boundary.
LOCK STATUS	LOCKED BEFORE OFFICIAL EVALUATION
FINAL SUBMISSION STATUS	PASS

One complete candidate was generated. Private preflight found no substantive defect, so no correction was made. No retry, selection among multiple scored finals, or post-lock rewrite occurred.

10. Final submitted benchmark response

LOCKED SCORED RESPONSE - answer content only. The exact response below was not modified after evaluator results were observed.

```
{"riddle":"I can open a lock and start an engine. I may be made of metal and carried in a pocket. What am I?"}
```

Locked file: 111 bytes; 1 visible line; terminal newline present. SHA-256: 081c53a5a079a8186f88b56b37387e078e95606ba672d3a85fa608080b0969.

11. Official evaluator result

Official Google Research source code was executed against the locked response. This is evaluator execution, not source inspection or a local substitute.

Direct JsonFormat checker	True
Direct ForbiddenWords checker	True
follow_instruction_list (strict)	[True, True]
follow_all_instructions (strict)	True
follow_instruction_list (loose)	[True, True]
follow_all_instructions (loose)	True
Official evaluator label	OFFICIAL EVALUATOR PASS

JsonFormat strips optional markdown fences and calls json.loads on the complete remaining value. ForbiddenWords searches case-insensitively for the configured word as a whole-word regex. Neither checker evaluates the riddle's semantic quality.

12. Manual / local constraint check

LOCAL REIMPLEMENTATION - NOT OFFICIAL EVALUATOR. The transparent local check agrees with the official result.

Response byte / visible-line count	111 bytes / 1 line; terminal newline present
Raw response SHA-256	081c53a5a079a8186f88b56b37387e078e95606ba672d3a85faff608080b0969
Raw official item SHA-256	021cca041229fbc dab2cbc6c3bf224bf1f839da86486fd8ac9eea07eeeb314d0
Leading / trailing non-whitespace bytes	0 / 1
JSON parse error	None
JSON top-level type	dict
Non-empty riddle field	PASS
Forbidden-word matches	[]
Official-equivalent forbidden-word result	PASS
Riddle describes target without spelling target	PASS
Non-whitespace material outside answer	Absent
Audit material in scored block	Absent

Manual task checks - separate from official score

One valid JSON block	PASS
Riddle describes a lock-opening and engine-starting object	PASS
Riddle includes a metal-and-pocket clue	PASS
Forbidden target word absent from the response	PASS
No unsupported authority, policy, or current fact	PASS
No audit material in scored block	PASS

13. Scores

Instruction-level score	2 / 2 encoded constraints passed
Prompt-level strict score	PASS
Prompt-level loose score	PASS
Final score	1 / 1 prompt-level strict pass
Official vs local discrepancy	None
Manual checks	Reported separately; not added to official IFEval score

14. Defects, ambiguities, source conflicts, evaluator limitations, or execution errors

No source conflict or official evaluator execution error was observed. The initial candidate passed private preflight with zero substantive corrections. No content was changed after lock.

JsonFormat accepts any value that parses as JSON after removing optional code-fence wrappers; it does not require a particular schema or prove that the riddle is a good description. ForbiddenWords uses a case-insensitive whole-word match and therefore does not test every possible spelling or semantic reference to the target.

The official constraints do not assess whether the riddle is elegant, uniquely solvable, or pedagogically useful. The local semantic checks are separate and are not part of the official IFEval score.

ROUTE VERIFICATION NOT EXPOSED remains a material limitation. The requested Luna Max condition is recorded as supplied and is not proof of a hidden backend build or route.

15. Evidence status

Blind selection evidence	FROZEN - selector prompt and metadata-only shortlist preserved; exact prompt withheld until key freeze
Official item evidence	FROZEN - exact raw JSONL line preserved; raw-line SHA-256 recorded
Official checker evidence	FROZEN - instructions.py and instructions_util.py blob identities recorded
Scored response evidence	LOCKED - SHA-256 recorded before official evaluation
Official evaluator evidence	EXECUTED - direct, strict, and loose results recorded
Local-check evidence	EXECUTED - transparent mirror and manual checks recorded
Runtime-control package	HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified
PDF visual QA	PASS - final rendered pages inspected; no clipping, overlap, missing glyphs, or footer collision observed

16. Record-integrity statement

This is an INTERNAL ONE-SHOT RECORD for the requested FASTPATH 5.6 / LUNA MAX condition. The metadata-only selector, frozen item, initial candidate, zero disclosed preflight corrections, locked response, evaluator outputs, local checks, and PDF were recorded in order. Locked response SHA-256: 081c53a5a079a8186f88b56b37387e078e95606ba672d3a85faff608080b0969. Official item-line SHA-256: 021cca041229fbc dab2cbc6c3bf224bf1f839da86486fd8ac9eea07eeeb314d0. Selector-prompt SHA-256: 947aa34b9e787648a7892baf7738f22a8354c00cc9993360522be1189381cec4. The record preserves the observed result and does not claim certification, independent verification, leaderboard standing, guaranteed determinism, semantic riddle certification, or proof of hidden backend identity.