

Grounded DI OS Benchmark Record

Prompt 47 of the Grounded DI cross-domain benchmark series - official Google Research IFEval item key 2497

INTERNAL ONE-SHOT RECORD - FASTPATH 5.6 / LUNA MAX - OFFICIAL EVALUATOR PASS

1. Document identity and condition

Primary document	Grounded DI OS Benchmark Record
Benchmark-series position	Cross-domain Prompt 47
Requested condition	Grounded DI OS / FastPath 5.6 / Luna Max / Protocol A
Run controls	Tier 18; DIP 0-166; Invocation 140A / 140; Delta H 0.0041; Benchmark Mode
Benchmark identifier	Google Research IFEval item key 2497
Encoded controls	detectable_format:number_highlighted_sections; startend:quotation
Benchmark phase	Blindly selected composite interaction test - one highlighted section and whole-response quotation wrapping
Labels applied	INTERNAL ONE-SHOT RECORD; FASTPATH 5.6 / LUNA MAX; OFFICIAL EVALUATOR PASS; LOCAL CHECK PASS
Labels not applied	certified; independently verified; universal; representative; official leaderboard score; proof of hidden backend identity

2. Purpose and blind selection scope

This record documents one observed composite execution of official Google Research IFEval item key 2497 under the requested FastPath 5.6 / Luna Max / Protocol A condition. The item was selected before its natural-language prompt was retrieved.

The benchmark selector was blind to prompt text, topic, answer, and content-derived features. It saw only fresh-item keys, ordered instruction IDs, and ordered kwargs from the official dataset. It preferred a fresh two-control checker tuple with minimal deterministic burden and then applied the lowest-key tie-break.

Blind selector result	selected_key = 2497
Selector input	Metadata only; 541 official rows; 43 unique previously represented keys excluded
Eligibility filter	498 fresh rows; 62 metadata-only candidates after the two-control and burden filters
Selector shortlist	Eight eligible fresh metadata records; key 2497 had one highlight and whole-response quotation wrapping
Selector prompt SHA-256	658368f9edb41c9105fa33943f69d2083de9f51ddd93ad75b574975ca47905e6
Selection boundary	Exact official prompt fetched only after key 2497 was frozen

The operator-supplied GDRCP v1.1.3 package is treated only as an evidence-harness specification: HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified.

3. Official benchmark URLs

Series architecture: <https://github.com/Grounded-DI/Benchmarks-Cross-Domain>

Audit reference: <https://github.com/Grounded-DI/briefwise-di-performance-on-harvey-lab-benchmark-audit>

Official evaluator tree: https://github.com/google-research/google-research/tree/master/instruction_following_eval

Official data: https://raw.githubusercontent.com/google-research/google-research/master/instruction_following_eval/data/input_data.jsonl

Official checker: https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions.py

Evaluator library: https://github.com/google-research/google-research/blob/master/instruction_following_eval/evaluation_lib.py

Word-count utility: https://github.com/google-research/google-research/blob/master/instruction_following_eval/instructions_util.py

Paper: <https://arxiv.org/abs/2311.07911>

4. Exact official item (key 2497)

The exact official JSONL line was frozen after blind selection and before response generation. Dataset blob: cbe52f6eefc3986fdac745b4acba4da1408eb146. instructions.py blob: 0007359bd37c7b9c1d9557f8937cb6586d2788a7. instructions_util.py blob: d50702400738df4378fe286c0090a416f9a172fa. evaluation_lib.py blob: c87422822f9e1565ca7574ac0b3cee5389c9f44e. Raw-line SHA-256: d3c575f2c64ca4212cf93615d3a28727966f8234edbe08ad50b1758d1509db55.

```
{
  "key": 2497,
  "prompt": "Give me a logical question about chatting and conquering. Make it interesting. Highlight at least one section of your answer with markdown, i.e. *highlighted section*. Put your entire response inside double quotation marks.",
  "instruction_id_list": [
    "detectable_format: number_highlighted_sections",
    "startend: quotation"
  ],
  "kwargs": {
    "num_highlights": 1
  }
}
```

5. Exact benchmark prompt (scored task only)

The following exact prompt is the only scored task. Record material is outside the scored-response boundary.

```
Give me a logical question about chatting and conquering. Make it interesting. Highlight at least one section of your answer with markdown, i.e. *highlighted section*. Put your entire response inside double quotation marks.
```

6. Observed runtime identity

Requested condition	Grounded DI OS / FastPath 5.6 / Luna Max / Protocol A
Exposed runtime identity	Codex agent environment; precise backend build or model snapshot not exposed
Route verification	ROUTE VERIFICATION NOT EXPOSED - requested condition is not independent proof of a hidden route
Runtime-control boundary	HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified

7. Timestamp and timezone

Source freeze	2026-09-03 15:44:29 UTC report time; blind metadata selection and item freeze occurred before candidate lock
Initial candidate	One complete candidate preserved before private preflight
Response lock	Locked before official evaluation; no post-lock response mutation
Official evaluation	Recorded after response lock; direct, strict, and loose paths executed
Timezone basis	UTC

8. Tool and web-search status

The official Google Research dataset and checker sources were retrieved through the selected GitHub source and matched to the listed freeze identities before response generation. No web search was used for the scored answer. No current-fact research, proprietary material, credential, or external service was needed.

The selector saw metadata only. The exact prompt and its content were not used for selection. No dependency repair was needed for the executed checker paths. The locked response bytes were unchanged throughout evaluation; repository publication remains conditional on the full pass and artifact checks.

9. Draft and preflight status

INITIAL DRAFT STATUS	PASS
PRIVATE PREFLIGHT RESPONSE CORRECTIONS	0
Correction	None. The complete candidate satisfied the official format controls and the manual content boundary. A local-only diagnostic predicate repair occurred after evaluation; it did not change the response or any official result.
LOCK STATUS	LOCKED BEFORE OFFICIAL EVALUATION
FINAL SUBMISSION STATUS	PASS

One complete candidate was generated and preserved as the single scored response. No retry, selection among multiple scored finals, or post-lock rewrite occurred. The local mirror's first question-mark predicate was corrected after evaluation because the mark appears immediately before the closing quote; the locked response remained byte-identical.

10. Final submitted benchmark response

LOCKED SCORED RESPONSE - answer content only. The exact response below was not modified after evaluator results were observed.

```
"If a clever commander can *win a kingdom by listening before speaking* how should they decide when a conversation has become a conquest?"
```

Locked file: 139 bytes; 1 visible line; terminal newline present. SHA-256: d77623efa0f02b2649ef4e0cb86497ee0c7a1d7008bc7fcf33fa499eccec724c.

11. Official evaluator result

Official Google Research source code was executed against the locked response. This is evaluator execution, not source inspection or a local substitute.

Direct HighlightSectionChecker	True
Direct QuotationChecker	True
follow_instruction_list (strict)	[True, True]
follow_all_instructions (strict)	True
follow_instruction_list (loose)	[True, True]
follow_all_instructions (loose)	True
Official evaluator label	OFFICIAL EVALUATOR PASS

HighlightSectionChecker recognizes non-empty asterisk spans and QuotationChecker checks the first and last non-whitespace characters after stripping. Neither checker evaluates whether the question is logically interesting or whether the highlighted phrase is semantically useful.

12. Manual / local constraint check

LOCAL REIMPLEMENTATION - NOT OFFICIAL EVALUATOR. The transparent local check agrees with the official result after one disclosed diagnostic-predicate repair.

Response byte / visible-line count	139 bytes / 1 line; terminal newline present
Raw response SHA-256	d77623efa0f02b2649ef4e0cb86497ee0c7a1d7008bc7fcf33fa499eccec724c
Raw official item SHA-256	d3c575f2c64ca4212cf93615d3a28727966f8234edbe08ad50b1758d1509db55
Leading / trailing non-whitespace bytes	0 / 1
Official-equivalent highlighted sections	1
Highlight match values	['*win a kingdom by listening before speaking*']
Whole-response double-quote boundary	PASS
Question mark inside quoted response	PASS
Logical chatting/conquering content	PASS
Non-whitespace material outside answer	Absent
Audit material in scored block	Absent
Local diagnostic repair	1 predicate repair after evaluator; locked bytes unchanged

Manual task checks - separate from official score

One non-empty highlighted section	PASS
Entire response enclosed in double quotes	PASS
Question concerns conversation and conquest	PASS
Question is phrased as a logical question	PASS
No unsupported authority, policy, or current fact	PASS
No audit material in scored block	PASS

13. Scores

Instruction-level score	2 / 2 encoded constraints passed
Prompt-level strict score	PASS
Prompt-level loose score	PASS
Final score	1 / 1 prompt-level strict pass
Official vs local discrepancy	None after disclosed local diagnostic repair
Manual checks	Reported separately; not added to official IFEval score

14. Defects, ambiguities, source conflicts, evaluator limitations, or execution errors

No source conflict or official evaluator execution error was observed. The initial candidate passed the substantive private preflight with zero response corrections. No content was changed after lock.

The first local mirror run used a predicate that looked for a question mark after the closing quote. That diagnostic was repaired after official evaluation to inspect the quoted interior correctly. The locked response bytes and official direct, strict, and loose results were unchanged.

HighlightSectionChecker recognizes certain asterisk spans and QuotationChecker checks boundary characters; neither establishes conventional Markdown rendering, logical novelty, or semantic quality. The local manual checks are separate and are not part of the official IFEval score.

ROUTE VERIFICATION NOT EXPOSED remains a material limitation. The requested Luna Max condition is recorded as supplied and is not proof of a hidden backend build or route.

15. Evidence status

Blind selection evidence	FROZEN - selector prompt and metadata-only shortlist preserved; exact prompt withheld until key freeze
Official item evidence	FROZEN - exact raw JSONL line preserved; raw-line SHA-256 recorded
Official checker evidence	FROZEN - instructions.py and instructions_util.py blob identities recorded
Scored response evidence	LOCKED - SHA-256 recorded before official evaluation
Official evaluator evidence	EXECUTED - direct, strict, and loose results recorded
Local-check evidence	EXECUTED - transparent mirror and manual checks recorded
Runtime-control package	HARNESS_SPECIFIED / NOT EXECUTED; runtime tests false; replay not executed; PackageReady false; not certified
PDF visual QA	PASS - final rendered pages inspected; no clipping, overlap, missing glyphs, or footer collision observed

16. Record-integrity statement

This is an INTERNAL ONE-SHOT RECORD for the requested FASTPATH 5.6 / LUNA MAX condition. The metadata-only selector, frozen item, initial candidate, zero response corrections, disclosed local diagnostic repair, locked response, evaluator outputs, local checks, and PDF were recorded in order. Locked response SHA-256: d77623efa0f02b2649ef4e0cb86497ee0c7a1d7008bc7fcf33fa499eccec724c. Official item-line SHA-256: d3c575f2c64ca4212cf93615d3a28727966f8234edbe08ad50b1758d1509db55. Selector-prompt SHA-256: 658368f9edb41c9105fa33943f69d2083de9f51ddd93ad75b574975ca47905e6. The record preserves the observed result and does not claim certification, independent verification, leaderboard standing, guaranteed determinism, semantic question certification, or proof of hidden backend identity.